

---

# Verification of the Blocking and Non-Blocking Michael-Scott Queue Algorithms

Mathias Pedersen, 201808137

---

Master's Thesis, Computer Science

May 2024

Advisor: Amin Timany



AARHUS  
UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE

# Abstract

►in English...◄

# Resumé

►in Danish...◄

# Acknowledgments



*Mathias Pedersen*  
*Aarhus, May 2024.*

# Contents

<b>Abstract</b>	<b>ii</b>
<b>Resumé</b>	<b>iii</b>
<b>Acknowledgments</b>	<b>iv</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Preliminaries</b>	<b>2</b>
2.1 HeapLang . . . . .	2
2.2 The Iris Program Logic Framework . . . . .	2
2.2.1 Fundamentals of Iris . . . . .	2
2.2.2 Hoare Triples and Weakest Pre-Condition . . . . .	2
2.2.3 Later Modality . . . . .	2
2.2.4 Invariants . . . . .	3
2.2.5 Resource Algebra . . . . .	3
2.2.6 Update modality and Viewshift . . . . .	3
2.2.7 Locks . . . . .	3
2.3 Formalisation in Coq . . . . .	3
2.3.1 Compiling the Project . . . . .	4
<b>3 Queue Specifications</b>	<b>5</b>
3.1 Specifications for Queues . . . . .	5
3.2 Defining a Sequential Specification . . . . .	5
3.3 Defining a Concurrent Specification . . . . .	6
3.4 Defining a Hocap-style Specification . . . . .	6
3.5 Deriving Sequential and Concurrent Specifications . . . . .	8
3.5.1 Deriving the Sequential Specification . . . . .	8
3.5.2 Deriving the Concurrent Specification . . . . .	10
<b>4 The Two-Lock Michael-Scott Queue</b>	<b>12</b>
4.1 Introduction . . . . .	12
4.2 Implementation . . . . .	12
4.2.1 Initialise . . . . .	12
4.2.2 Enqueue . . . . .	13
4.2.3 Dequeue . . . . .	13
4.2.4 Observations on the Two-Lock Michael-Scott Queue . . . . .	13
<b>5 Proving Specifications for the Two-Lock Michael-Scott Queue</b>	<b>18</b>
5.1 Introduction . . . . .	18
5.2 Sequential Specification . . . . .	18
5.2.1 Sequential Queue Predicate . . . . .	18
5.2.2 Proof Outline . . . . .	19
5.3 Concurrent Specification . . . . .	21
5.3.1 Concurrent Queue Predicate . . . . .	21

5.3.2	Linearisation Points . . . . .	24
5.3.3	Proof Outline . . . . .	24
5.3.4	Discussing the need for Old Nodes . . . . .	27
5.4	Hocap-Style Specification . . . . .	28
5.4.1	Introduction . . . . .	28
5.4.2	Hocap-Style Queue Predicate . . . . .	28
5.4.3	Proof Sketch . . . . .	29
<b>6</b>	<b>The Lock-Free Michael-Scott Queue</b>	<b>31</b>
6.1	Introduction . . . . .	31
6.2	Implementation . . . . .	31
6.2.1	Initialise . . . . .	31
6.2.2	Enqueue . . . . .	31
6.2.3	Dequeue . . . . .	32
6.2.4	Prophecies . . . . .	32
<b>7</b>	<b>Proving Hocap-style Specification for the Lock-Free Michael-Scott Queue</b>	<b>34</b>
7.1	Introduction . . . . .	34
7.2	Reachability . . . . .	34
7.2.1	Concrete Reachability . . . . .	34
7.2.2	Abstract Reachability . . . . .	35
7.3	Specifications for Lock-Free M&S Queue . . . . .	36
7.4	Hocap-style Queue Predicate . . . . .	36
7.5	Proof Outline . . . . .	37
7.6	Discussion . . . . .	42
<b>8</b>	<b>On the Resource Algebras Used</b>	<b>44</b>
8.1	Tokens . . . . .	44
8.2	Abstract State . . . . .	44
8.3	Abstract Points-to and Reachability . . . . .	45
<b>9</b>	<b>Conclusion and Future Work</b>	<b>47</b>
	<b>Bibliography</b>	<b>48</b>
<b>A</b>	<b>Common Definitions and Lemmas</b>	<b>49</b>
A.1	Lists and Nodes . . . . .	49
A.2	Results About the isLL Predicate . . . . .	50

# Chapter 1

## Introduction

►motivate and explain the problem to be addressed◄

►example of a citation: Michael and Scott [1996]◄ ►get your bibtex entries from <https://dblp.org/>◄ ►Emphasise that specs for two-lock and lock-free are equivalent, with only the queue functions being different. This is even shown in the coq proofs◄ ►Have a bullet point list of contributions:◄

- Giving direct specifications for TL and LF Queues
- Prove that M&S Queues satisfy those specifications (instead of refinement)
- Proven that TL and LF satisfy *same* specification
- Shown that Hocap implies sequential and concurrent
- Mechanised everything presented in the paper in coq

# Chapter 2

## Preliminaries

►Mention that the project uses heaplang and the program logic iris, and hence we need to know about them◄

### 2.1 HeapLang

►Write about heaplang language: ml like, with a heap, Cas◄ ►Refer to ILN for specification◄  
►Talk about Syntactic sugar: i.e.  $e1 \;; e2 = (\text{lam } v, e2) \; e1$  where  $v$  is fresh, and CAS ... as Snd (CMPXHG ...), and derived rules for them◄ ►also mention prophecies and refer to the section talking about them◄

### 2.2 The Iris Program Logic Framework

In this section, we give a brief introduction to Iris – the logic we use to reason about the M&S Queues. Iris is quite expressive and supports a myriad of features and derived rules, many of which have been utilised in this project. As such, it will be impossible to cover all facets of Iris in detail, so we limit ourselves to give an overview of the main aspects of Iris. If the reader wishes a more thorough introduction, or wants to see further details of the topics covered, please consult the Iris Lecture Notes [Birkedal and Bizjak, 2017].

►FOCUS IN THIS SECTION is on explaining the basic ideas of the concepts and showing the notation used. Assume the reader is somewhat familiar with program verification.  
◄

#### 2.2.1 Fundamentals of Iris

►program logic framework◄ ►instantiate with a language, here heaplang◄ ►higher order logic◄ ►separation logic◄ ►- exclusive ownership and duplicability◄ ►introduce persistent modality, duplicability◄ ►points-to predicates◄ ►resource oriented◄ ►derivation rules examples...◄

#### 2.2.2 Hoare Triples and Weakest Pre-Condition

►explain hoare triples◄ ►talk about Partial correctness and Adequacy◄ ►explain weakest pre-conditions◄ ►show how they are related◄

#### 2.2.3 Later Modality

►explain concept◄ ►ties propositions to program steps◄ ►Explain the löb induction rule, mention the intuition when  $P$  is hoare triple◄



The later modality adds a notion of step indexing to Iris propositions. This step indexing is technically parallel to the notion of taking steps in the programs, but the way we define hoare-triples and weakest-preconditions tie program steps together with steps in the logic. In other words, a single Later corresponds to a single step in the program.

## 2.2.4 Invariants

►explain the idea◄ ►open around atomic expression◄ ►Explain that we technically only get the resources Later, when opening invariant, but since we are also only required to prove the invariant again, later, and many of the structural rules work with a later, we shall not worry about getting the resources later in the proof outlines, as it usually isn't an issue. We mention the later in the few cases where only getting the resources Later is important.◄ ►only one opening at a time◄ ►namespaces◄ ►Show the basic rules. Inv-alloc (FUP), ...◄

## 2.2.5 Resource Algebra

►Introduce the idea◄ ►Give the formal definition and explain the parts◄ ►Cite first line of definition Definition 8.10 in ILN◄ ►Commutative semigroup  $\mathcal{M}$ : set with associative and commutative operation, denoted  $\cdot$ ◄ ►preorder, extension order:  $\preceq$ ◄ ►Validity◄ ►Core and persistency◄ ►Unital◄ ►tying them into the logic◄ ►also ghost names◄ ►ghost-allocate own-op, own-valid...◄ ►mention that points-to predicate is also a resource algebra◄ ►composing resource algebras◄ ►we introduce the resource algebras we use at the point of use◄

## 2.2.6 Update modality and Viewshift

►explain the ideas◄ ►frame preserving update◄

## 2.2.7 Locks

►show and explain the lock specification◄

## 2.3 Formalisation in Coq

Iris has been mechanised in the Coq proof assistant<sup>1</sup> – a tool to machine-check proofs of mathematical assertions. All results in this project have been completely machine-verified in the Iris mechanisation in Coq. Specifications in the Coq formalisation of Iris are usually written in terms of hoare-triples but proved by first converting the hoare-triples to equivalent weakest-preconditions, as this is usually easier to work with. The proofs presented in this report will follow suit and give specifications using hoare-triples, but prove them assuming they are weakest-preconditions. The proofs presented in the report thus follow the mechanised proofs very closely, making it possible to “step-through” the mechanised proofs in tandem with reading the paper-proofs presented in this report.

One caveat is that there is somewhat of a discrepancy between iris on paper, and the Iris formalisation in coq. Working with the latter requires a bit deeper understanding of the model of Iris. Jung et al. [2018] explains the underlying model of an older version of Iris, but many of the concepts discussed are still relevant.<sup>2</sup>

Table 2.1 gives an overview over the files developed in the project and how they relate to this report. All files related to the project are available at <https://github.com/MatteP1/thesis>.

<sup>1</sup>The mechanisation can be found at <https://gitlab.mpi-sws.org/iris/iris/>

<sup>2</sup>For an up-to-date presentation, please consult the Technical Reference at <https://iris-project.org/>

File Name	Relevant Sections	Description
queue_specs.v	Chapter 3	Implementation-independent specifications and derivations.
twoLockMSQ_derived.v	Section 3.5	
lockFreeMSQ_derived.v	Section 3.5	
MSQ_common.v	Chapters 4 - 7	Common definitions and lemmas.
twoLockMSQ_impl.v	Chapter 4	Two-Lock M&S Queue implementation and proofs of sequential, concurrent, and Hocap-style specifications.
twoLockMSQ_sequential_spec.v	Section 5.2	
twoLockMSQ_concurrent_spec.v	Section 5.3	
twoLockMSQ_hocap_spec.v	Section 5.4	
lockFreeMSQ_impl.v	Chapter 6	Lock-Free M&S Queue implementation and Hocap-style specification proof.
lockFreeMSQ_hocap_spec.v	Chapter 7	
lockAndCCFreeMSQ_impl.v	Section 7.6	Consistency-Check-Free version of Lock-Free M&S Queue.
lockAndCCFreeMSQ_hocap_spec.v	Section 7.6	

Table 2.1: Overview of Coq Files.

### 2.3.1 Compiling the Project

Compiling the project requires both Coq and Iris to be installed. Once installed, open a terminal and navigate to the project folder, `/thesis/`, which contains the `_CoqProject` file. Here, run `make`. This will both create a Coq Makefile and run it. The project is known to compile with Coq version 8.19.0 and Iris version 4.2.0.

## Chapter 3

# Queue Specifications

### 3.1 Specifications for Queues

In this chapter we discuss some possible specifications for queue data structures in general. As such, we need to make some basic assumption about what we expect from a queue. Firstly, we adopt the convention that a queue consists of three functions: initialize, enqueue, and dequeue. Exactly what these functions do will depend on the specific implementation of the queue, but we give some general pointers to what we expect of them.

The initialize function should create a queue which is initially empty. The functions enqueue and dequeue can then be invoked subsequently on said queue.

In addition to being parametrised on the queue, the enqueue function should also take a value as input. When invoking enqueue with such a value, the function should add this value to the end of the queue.

The dequeue function will attempt to dequeue an element from the queue. Since queues are allowed to be empty, the dequeue function is assumed to return an option value. If the queue is empty, then dequeue will return None, and otherwise it will remove an element from the front of the queue, and return this wrapped in a Some.

Working in a concurrent setting often gets quite complicated quite fast, and proving that ones queue satisfies the above desirables can become quite tricky. In fact, even defining those qualities formally can become non-trivial. As such, we give three different specifications for queues. In section 3.2 we give a specification that assumes the queue is run in a sequential setting. Next, in section 3.3, we give a specification that *does* allow for concurrency but gives up on some of the above qualities. Primarily, it does not track the contents of the queue; invoking enqueue on a value doesn't guarantee that the value is added to the queue. Finally, in section 3.4 we give a specification that allows for both concurrency *and* tracking of the contents of the queue.

### 3.2 Defining a Sequential Specification

Let us first consider a specification in the simple case where we do not allow for concurrency. In this case, we know that only a single thread will interact with the queue at any given point in a sequential manner. The specification we give will track the exact contents of the queue. To this end, we shall define the *abstract state* of the queue, denoted  $xs_v$  as a list of HeapLang values. I.e.  $xs_v : List\ Val$ . We adopt the convention that enqueueing an element is done by adding it to the front of the list, and dequeueing removes the last element of the list (if such an element exists). The reason for this choice is purely technical.

To allow queues to use whichever ghost names they like, we introduce the type “*SeqQghostnames*” whose purpose is to keep track of the ghost names used for a specific queue. One may think of *SeqQghostnames* as set of fixed-length tuples of ghost names.

With this, we give our definition of a sequential specification for a queue.

**Definition 3.2.1** (Sequential Specification).

$$\begin{aligned}
& \exists \text{isQueues} : \text{Val} \rightarrow \text{List Val} \rightarrow \text{SeqQgnames} \rightarrow \text{Prop}. \\
& \quad \{\text{True}\} \text{ initialize } () \{v_q. \exists G. \text{isQueues}(v_q, [], G)\} \\
& \wedge \quad \forall v_q, v, xs_v, G. \{\text{isQueues}(v_q, xs_v, G)\} \text{ enqueue } v_q \ v \{w. \text{isQueues}(v_q, (v :: xs_v), G)\} \\
& \wedge \quad \forall v_q, xs_v, G. \{\text{isQueues}(v_q, xs_v, G)\} \\
& \quad \text{ dequeue } v_q \\
& \quad \left\{ \begin{array}{l} (xs_v = [] * w = \text{None} * \text{isQueues}(v_q, xs_v, G)) \vee \\ w. (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * \text{isQueues}(v_q, xs'_v, G)) \end{array} \right\}
\end{aligned}$$

The proposition  $\text{isQueues}(v_q, xs_v, G)$  captures that the value  $v_q$  is a queue, whose contents matches that of our abstract representation  $xs_v$ , and the queue uses the ghost names described by  $G$ . Clients of a queue satisfying this specification can then apply the three Hoare-triples at invocations for initialize, enqueue, and dequeue, which tells them how the queue changes as a result of the invocation.

Note that the  $\text{isQueues}$  predicate is not required to be persistent, hence it cannot be duplicated and given to multiple threads. Since this predicate is required to apply the Hoare-triples, then only a single thread can use the specification at a time. This is the sense in which the specification is sequential.

### 3.3 Defining a Concurrent Specification

As discussed in the previous section, a concurrent specification will need the queue predicate to be duplicable to allow multiple threads to use the specification concurrently. For this specification, we denote the queue predicate by  $\text{isQueue}_C$ .

To achieve duplicability of  $\text{isQueue}_C$ , we shall give up on tracking the abstract state of the queue. The reason for doing so is the following. Consider the case where we own the queue predicate and it states that the abstract state of queue is  $xs_v$ . We spawn two threads and give them both the queue predicate stating that the abstract value is  $xs_v$ . Now they both invoke enqueue, one with value  $v$  and the other with  $v'$ . This make one thread conclude that the abstract state of the queue is  $v ++ xs_v$ , and the other that  $v' ++ xs_v$ , whereas in reality, the queue contains both  $v$  and  $v'$ . This is of course not desirable, and it is not immediately obvious how to solve this issue.

We *can* however make the queue guarantee that a given property holds of all the values in the queue. We do this by parametrising  $\text{isQueue}_C$  with a predicate,  $\Phi$ , which it will maintain holds for all elements in the queue. In this way, when dequeuing, we at least know that if we get some value, then  $\Phi$  holds of this value. The specification we wish to prove is as follows.

**Definition 3.3.1** (Concurrent Specification).

$$\begin{aligned}
& \exists \text{isQueue}_C : (\text{Val} \rightarrow \text{Prop}) \rightarrow \text{Val} \rightarrow \text{ConcQgnames} \rightarrow \text{Prop}. \\
& \forall \Phi : \text{Val} \rightarrow \text{Prop}. \\
& \quad \forall v_q, G. \text{isQueue}_C(\Phi, v_q, G) \implies \Box \text{isQueue}_C(\Phi, v_q, G) \\
& \wedge \quad \{\text{True}\} \text{ initialize } () \{v_q. \exists G. \text{isQueue}_C(\Phi, v_q, G)\} \\
& \wedge \quad \forall v_q, v, G. \{\text{isQueue}_C(\Phi, v_q, G) * \Phi(v)\} \text{ enqueue } v_q \ v \{w. \text{True}\} \\
& \wedge \quad \forall v_q, G. \{\text{isQueue}_C(\Phi, v_q, G)\} \text{ dequeue } v_q \{w. w = \text{None} \vee (\exists v. w = \text{Some } v * \Phi(v))\}
\end{aligned}$$

This specification additionally requires that  $\text{isQueue}_C$  is persistent, which means that it can be duplicated and given to multiple threads. We again use a collection of ghost names, which we here denote  $\text{ConcQgnames}$ .

### 3.4 Defining a Hocap-style Specification

In this section we will explore our most general specification: one that allows for both concurrency and tracking of the contents of the queue. We shall refer to this specification as a Hocap-style specification – Higher Order Concurrent Abstract Predicate – since it will be concurrent and parametrised by abstract

predicates. This specification is more general than both the sequential and concurrent specifications in the sense that they are derivable from the Hocap-style specification. We prove this in section 3.5.

As with concurrent specification, we cannot simply parametrise the queue predicate (now denoted `isQueue`) with the abstract state of the queue. So to allow clients to keep track of the contents of the queue, we will “split” the abstract state into two parts; the authoritative view and the fragmental view. The clients will then own the fragmental view, allowing them to keep track of the contents of the queue, whereas the `isQueue` predicate will own the authoritative view. We will in particular make sure that, if one has both the fragmental and authoritative views, then these agree on the abstract state of the queue. Further, it is only possible to update the abstract state of the queue if one possess both the authoritative and fragmental views. Hence, clients will have to supply the fragmental view to be able to apply the specifications for enqueue and dequeue.

For an abstract state  $xs_v$  and a ghost name  $\gamma$ , we shall use the notation  $\gamma \models_{\bullet} xs_v$  to mean that the authoritative view of the abstract state associated with  $\gamma$  is  $xs_v$ . Similarly we write  $\gamma \models_{\circ} xs_v$  to mean that the fragmental view associated with  $\gamma$  is  $xs_v$ .

We introduce three lemmas to help working with these predicates. These are proved in section 8.2. The first lemma shows that we can create fresh authoritative and fragmental views for any abstract state. The abstract state will usually be empty when allocating.

**Lemma 1** (Abstract State Alloc).  $\forall xs_v.$

$$\vdash \exists \gamma. \gamma \models_{\bullet} xs_v * \gamma \models_{\circ} xs_v$$

The second shows that the authoritative and fragmental views of the abstract state agree.

**Lemma 2** (Abstract State Agree).  $\forall \gamma, xs_v, xs'_v.$

$$\gamma \models_{\bullet} xs'_v * \gamma \models_{\circ} xs_v \vdash xs_v = xs'_v$$

The final lemma shows that, if we own both the authoritative and fragmental views, we are allowed to update the abstract state to whatever we like.

**Lemma 3** (Abstract State Update).  $\forall \gamma, xs_v, xs'_v, xs''_v.$

$$\gamma \models_{\bullet} xs'_v * \gamma \models_{\circ} xs_v \Rightarrow \gamma \models_{\bullet} xs''_v * \gamma \models_{\circ} xs''_v$$

The collection of ghost names is now denoted  $Qnames$ , but as all queues will have to deal with the ghost name keeping track of the abstract state as describe above, we require that any  $Qnames$  will contain a ghost name used for this purpose. We shall refer to this ghost name as  $\gamma_{Abst}$ .

With this, we now give the Hocap-style specification, and explain the intricacies of it afterwards.

**Definition 3.4.1** (Hocap Specification).

$$\exists \text{isQueue} : Val \rightarrow Qnames \rightarrow \text{Prop}.$$

$$\forall v_q, G. \text{isQueue}(v_q, G) \implies \Box \text{isQueue}(v_q, G)$$

$$\wedge \{ \text{True} \} \text{ initialize } () \{ v_q. \exists G. \text{isQueue}(v_q, G) * G.\gamma_{Abst} \models_{\circ} [] \}$$

$$\wedge \forall v_q, v, G, P, Q. (\forall xs_v. G.\gamma_{Abst} \models_{\bullet} xs_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i^{\uparrow}} \triangleright G.\gamma_{Abst} \models_{\bullet} (v :: xs_v) * Q) \multimap \{ \text{isQueue}(v_q, G) * P \} \text{ enqueue } v_q \ v \{ w.Q \}$$

$$\wedge \forall v_q, G, P, Q.$$

$$\left( \forall xs_v. G.\gamma_{Abst} \models_{\bullet} xs_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i^{\uparrow}} \triangleright \left( \begin{array}{l} (xs_v = [] * G.\gamma_{Abst} \models_{\bullet} xs_v * Q(\text{None})) \\ \vee \left( \begin{array}{l} \exists v, xs'_v. xs_v = xs'_v ++ [v] * \\ G.\gamma_{Abst} \models_{\bullet} xs'_v * Q(\text{Some } v) \end{array} \right) \end{array} \right) \right) \multimap \{ \text{isQueue}(v_q, G) * P \} \text{ dequeue } v_q \{ w.Q(w) \}$$

As for the concurrent specification, we require that the queue predicate, `isQueue`, is persistent, giving us support for concurrent clients.

Next, the specification for `initialize` gives clients an additional resource in the postcondition: the ownership of the fragmental view of the empty list,  $G.\gamma_{Abst} \models_{\circ} []$ . As discussed above, this allows them to keep track of the contents of the queue.

The specifications for `enqueue` and `dequeue` have now been parametrised by two predicates:  $P$  and  $Q$ . The clients get to pick  $P$  and  $Q$ , and the choice depends on what the client wishes to prove;  $P$  describes those resources that the client has before `enqueue` or `dequeue`, and  $Q$  the resources it will have after.

Hence  $P$  is in the precondition and  $Q$  in the postcondition of the associated Hoare-triples. However, before the client gets access to the Hoare-triple for enqueue or dequeue they must prove a viewshift. This viewshift states how the abstract state of the queue will change as a result of running enqueue or dequeue, and further shows that  $P$  can be updated to  $Q$ . Note that the consequent of the viewshift contains a  $\triangleright$ . This signifies that the update in the abstract state is tied to a step in the code. The mask on the viewshift further disallows opening of invariants in the namespace  $\mathcal{N}.i$ . This is to allow queues that use invariants to apply the viewshifts while their invariants are open (their invariants must of course be within the namespace  $\mathcal{N}.i$ ).

It might seem a bit strange that the client has to prove that the abstract state can be updated, but remember that the client owns the fragmental view, and that both this and the authoritative view, which is owned by the queue, is needed to update the abstract state. When proving the viewshift, clients are not updating the abstract state of the queue, they are merely showing that they can supply the fragmental view, allowing the abstract state to be updated. This then enables the queue to update the authoritative view of the abstract state (using the proved viewshift) in conjunction with updating the concrete view.

Exactly how client supply the fragmental view depends on what the client wants to achieve. We will see two options, when we derive the sequential and concurrent specifications from this Hocap-style specification in the next section.

## 3.5 Deriving Sequential and Concurrent Specifications

It is technically possible to derive the sequential and concurrent specifications from the Hocap-style specification without having proven the Hocap-style specification for a specific queue. However, it might be beneficial for the reader to first see how we can prove each specification *directly* for a specific queue, which we show in chapter 5, and then return to this section.

In this section we show that we can derive the sequential and concurrent specifications from sections 3.2 and 3.3 from the Hocap-style specification we saw in the previous chapter. These derivations are implementation independent, so we assume that we have some functions initialize, enqueue, and dequeue which satisfy the Hocap-style specification of definition 3.4.1, and we wish to prove that they also satisfy the sequential and concurrent specifications of definitions 3.2.1 and 3.3.1. That is, we assume that we have a persistent queue predicate, `isQueue`, and the three Hocap-style specifications for initialize, enqueue, and dequeue.

### 3.5.1 Deriving the Sequential Specification

Recall the sequential specification specified in definition 3.2.1. It demands a queue predicate, `isQueueS`. We here choose to define it as follows.

**Definition 3.5.1** (`isQueueS` Predicate (Derive)).

$$\begin{aligned} \text{isQueue}_S(v_q, xs_v, G_S) &\triangleq \exists G_H \in Qgnames. \\ &\quad \text{projQgnames}_S(G_H) = G_S * \\ &\quad \text{isQueue}(v_q, G_H) * \\ &\quad G_H.\gamma_{\text{Abst}} \mapsto_{\circ} xs_v \end{aligned}$$

Here, `projQgnamesS(GH)` simply creates an element of *SeqQgnames*, with ghost names matching those of  $G_H$ . We proceed to prove the sequential specifications for the three queue functions.

#### Sequential Initialise Specification

Recall the sequential specification for initialise:

$$\{\text{True}\} \text{ initialize } () \{v_q. \exists G_S. \text{isQueue}_S(v_q, [], G_S)\}$$

This Hoare-triple follows almost directly from the Hocap-style initialise specification – they only differ in the post-condition. Recall that the post-condition in the Hocap-style specification (lemma 3.4.1) states

$\exists G_H. \text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} []$ . By the rule of consequence, we are done if we can show that this post-condition implies the sequential post-condition. So we assume we have some  $G_H$  such that

$$\text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} [] \quad (3.1)$$

Note that from 3.1, we can easily establish  $\text{isQueues}(v_q, [], \text{projQnames}_S(G_H))$ . Hence, by choosing  $\text{projQnames}_S(G_H)$  for  $G_S$ , we are done.

### Sequential Enqueue Specification

This specification states:

$$\forall v_q, v, xs_v, G_S. \{ \text{isQueues}(v_q, xs_v, G_S) \} \text{ enqueue } v_q \ v \ \{ w. \text{isQueues}(v_q, (v :: xs_v), G_S) \}$$

So assume some  $v_q, v, xs_v$ , and  $G_S$ . Unfolding  $\text{isQueues}$ , we additionally assume some  $G_H$  which projects to  $G_S$ . Our goal thus becomes:

$$\{ \text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \} \text{ enqueue } v_q \ v \ \{ w. \text{isQueues}(v_q, (v :: xs_v), G_S) \}$$

To prove the Hoare-triple, we shall use the Hocap-style specification for enqueue. This however requires us to pick  $P$  and  $Q$ , and prove the resulting viewshift. We choose

$$P \triangleq G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \quad Q \triangleq G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} (v :: xs_v)$$

and assume some  $xs'_v$ . We must then prove the viewshift:

$$G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs'_v * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i \uparrow} \triangleright G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} (v :: xs'_v) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} (v :: xs_v)$$

Assume  $G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs'_v$  and  $G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v$ . By lemma 2,  $xs_v = xs'_v$ , hence, we can apply lemma 3 to update the authoritative and fragmental views to  $(v :: xs_v)$ , which exactly proves the consequent of the viewshift.<sup>1</sup>

With this, we now get access to the following Hoare-triple:

$$\{ \text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \} \text{ enqueue } v_q \ v \ \{ w. G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} (v :: xs_v) \} \quad (3.2)$$

The pre-condition already matches our goal, so we just have to get the post-condition to match. To do this, we must get  $\text{isQueue}(v_q, G_H)$  in the post-condition of 3.2, but since  $\text{isQueue}(v_q, G_H)$  is persistent and we have it in the pre-condition, we may also assume it in post-condition.

### Sequential Dequeue Specification

We use a similar approach to above to prove the sequential dequeue specification:

$$\begin{aligned} & \forall v_q, xs_v, G_C. \{ \text{isQueues}(v_q, xs_v, G_C) \} \\ & \text{ dequeue } v_q \\ & \left\{ w. \begin{aligned} & (xs_v = [] * w = \text{None} * \text{isQueues}(v_q, xs_v, G_C)) \vee \\ & (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * \text{isQueues}(v_q, xs'_v, G_C)) \end{aligned} \right\} \end{aligned}$$

So we assume some  $v_q, xs_v, G_S$ , and  $G_H$  such that  $G_H$  projects to  $G_S$ . Our goal is now:

$$\begin{aligned} & \{ \text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \} \\ & \text{ dequeue } v_q \\ & \left\{ w. \begin{aligned} & (xs_v = [] * w = \text{None} * \text{isQueues}(v_q, xs_v, G_S)) \vee \\ & (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * \text{isQueues}(v_q, xs'_v, G_S)) \end{aligned} \right\} \end{aligned}$$

This time we instantiate the Hocap-style dequeue specification with the following choices:

$$\begin{aligned} P & \triangleq G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v \\ Q(w) & \triangleq \begin{aligned} & (xs_v = [] * w = \text{None} * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v) \vee \\ & (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs'_v) \end{aligned} \end{aligned}$$

---

<sup>1</sup>The consequent technically has a  $\triangleright$ , but proving something *now* is stronger than proving it *later*.

We must now prove the resulting viewshift to get the Hoare-triple (note that we haven't substituted in  $P$  and  $Q$  for the sake of readability). So assume some  $xs'_v$ . We must show:

$$G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs'_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i^\dagger} \triangleright \left( \begin{array}{l} (xs'_v = [] * G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs'_v * Q(\text{None})) \\ \vee \left( \begin{array}{l} \exists v, xs''_v. xs'_v = xs''_v ++ [v] * \\ G_H.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs''_v * Q(\text{Some } v) \end{array} \right) \end{array} \right)$$

By lemma 2,  $xs'_v$  must be equal to  $xs_v$ . We do a case analysis on  $xs_v$ . If  $xs_v$  is empty, we prove the left disjunct in the consequent of the viewshift, *without* updating the authoritative and fragmental views. If  $xs_v$  is non-empty, i.e.  $xs_v = xs''_v ++ [v]$  for some  $xs''_v$  and  $v$ , then we prove the right-side of the consequent in the viewshift by using lemma 3 to update the authoritative and fragmental views to the new abstract state,  $xs''_v$ .

With this, we get access to the Hoare-triple (now with  $P$  and  $Q$  substituted in):

$$\begin{aligned} & \{\text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v\} \\ & \text{dequeue } v_q \\ & \left\{ w. \begin{array}{l} (xs_v = [] * w = \text{None} * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v) \vee \\ (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs'_v) \end{array} \right\} \end{aligned} \quad (3.3)$$

As before, this almost matches the Hoare-triple we have to prove. We use the rule of consequence to massage the post-condition into the correct shape, which concludes the derivation.

### 3.5.2 Deriving the Concurrent Specification

We prove the concurrent specification of definition 3.3.1. Remember that we need the  $\text{isQueue}_C$  predicate to be persistent, hence we cannot simply assert  $G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v$  as we did for  $\text{isQueues}$ .<sup>2</sup> Instead, we will put it into an invariant. The queue predicate we will use looks as follows.

**Definition 3.5.2** ( $\text{isQueue}_C$  Predicate (Derive)).

$$\begin{aligned} \text{isQueue}_C(\Psi, v_q, G_C) & \triangleq \exists G_H \in Q\text{gnames}. \\ & \text{projQgnames}_C(G_H) = G_C * \\ & \text{isQueue}(v_q, G_H) * \\ & \boxed{\exists xs_v. G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v * \text{All}(xs_v, \Psi)}^{\mathcal{N}.c} \end{aligned}$$

Persistency of  $\text{isQueue}_C$  follows by the persistency of  $\text{isQueue}$  and the fact that invariants and equalities are persistent.

#### Concurrent Initialise Specification

We have to derive the specification:

$$\{\text{True}\} \text{ initialize } () \{v_q. \exists G_C. \text{isQueue}_C(\Phi, v_q, G_C)\}$$

As before, this specification only differs from the Hocap-style specification for initialise in the post-condition. We use the *generalised* rule of consequence  $\blacktriangleright \text{ref} \blacktriangleleft$  and show that the post-condition of the Hocap-style specification, i.e.  $\exists G_H. \text{isQueue}(v_q, G_H) * G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} []$ , implies the post-condition above but with an update modality,  $\Rightarrow$ , in front. As before, we pick  $\text{projQgnames}_S(G_H)$  for  $G$ , allowing us to prove all but the invariant:

$$\Rightarrow \boxed{\exists xs_v. G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v * \text{All}(xs_v, \Psi)}^{\mathcal{N}.c}$$

We have  $G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} []$  from the Hocap-style post-condition, and  $\text{All}([], \Psi)$  is equivalent to  $\text{True}$ . Hence, we can deduce  $\exists xs_v. G_H.\gamma_{\text{Abst}} \Rightarrow_{\circ} xs_v * \text{All}(xs_v, \Psi)$ . We are done by  $\text{Inv-alloc} \blacktriangleright \text{ref} \blacktriangleleft$ , which turns this into an invariant.

<sup>2</sup>As we explain in section 8.2, no elements of the abstract state resource algebra are persistent.



## Concurrent Enqueue Specification

We must derive:

$$\forall v_q, v, G_C. \{ \text{isQueue}_C(\Phi, v_q, G_C) * \Phi(v) \} \text{ enqueue } v_q \ v \{ w. \text{True} \}$$

Assume some  $v_q, v, G_C, G_H$  such that  $G_H$  projects to  $G_S$ , and the invariant  $\boxed{\exists xs_v. G_H. \gamma_{\text{Abst}} \mapsto_{\circ} xs_v}^{\mathcal{N}.c}$ . Our goal becomes the following Hoare-triple:

$$\{ \text{isQueue}(v_q, G_H) * \Phi(v) \} \text{ enqueue } v_q \ v \{ w. \text{True} \}$$

We specialise the Hocap-style enqueue specification with  $P \triangleq \Psi(v)$  and  $Q \triangleq \text{True}$ . With this choice, the Hoare-triple we get after proving the viewshift exactly matches our goal. Hence, we are done if we can prove the viewshift.

We assume  $xs'_v$ , and prove the viewshift:

$$G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs'_v * \Psi(v) \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i \uparrow} \triangleright G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} (v :: xs'_v) * \text{True}$$

So assume  $G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs'_v$  and  $\Psi(v)$ . We open the invariant giving us  $G_H. \gamma_{\text{Abst}} \mapsto_{\circ} xs_v$  and  $\text{All}(xs_v, \Psi)$  for some  $xs_v$ . By lemma 2 we know that  $xs_v = xs'_v$ . We now update the abstract state to  $(v :: xs_v)$  using lemma 3, obtaining  $G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} (v :: xs_v)$  and  $G_H. \gamma_{\text{Abst}} \mapsto_{\circ} (v :: xs_v)$ . We use the first to prove the consequent of the viewshift. Before we are done, we must close the invariant. We use the fragmental part together with  $\text{All}(xs_v, \Psi)$  and  $\Psi(v)$  to do this.

## Concurrent Dequeue Specification

Finally, we derive the concurrent specification for dequeue.

$$\forall v_q, G_C. \{ \text{isQueue}_C(\Phi, v_q, G_C) \} \text{ dequeue } v_q \{ w. w = \text{None} \vee (\exists v. w = \text{Some } v * \Phi(v)) \}$$

So we assume some  $v_q, G_S, G_H$  such that  $G_H$  projects to  $G_S$ , and the invariant. Our goal is now:

$$\{ \text{isQueue}(v_q, G_H) \} \text{ dequeue } v_q \{ w. w = \text{None} \vee (\exists v. w = \text{Some } v * \Phi(v)) \}$$

We make the following choices for  $P$  and  $Q$ .

$$P \triangleq \text{True} \qquad Q(w) \triangleq w = \text{None} \vee (\exists v. w = \text{Some } v * \Psi(v))$$

Again, with this choice, the Hoare-triple we get after proving the viewshift exactly matches our goal.

We assume some  $xs'_v$ , and prove the viewshift (again, without substituting in  $P$  and  $Q$ ):

$$G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs'_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i \uparrow} \triangleright \left( \begin{array}{l} (xs'_v = [] * G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs'_v * Q(\text{None})) \\ \vee \left( \begin{array}{l} \exists v, xs''_v. xs'_v = xs''_v ++ [v] * \\ G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs''_v * Q(\text{Some } v) \end{array} \right) \end{array} \right)$$

Assume  $G_H. \gamma_{\text{Abst}} \mapsto_{\bullet} xs'_v$ . Opening the invariant, we get  $\text{All}(xs_v, \Psi)$  and the fragmental part,  $G_H. \gamma_{\text{Abst}} \mapsto_{\circ} xs_v$ , for some  $xs_v$ , which by lemma 2 we know is equal to  $xs'_v$ . We proceed by case analysis on  $xs_v$ . If it is empty, we simply close the invariant again, and proceed to prove the first disjunct of the consequent. If it is not empty, then we have  $xs_v = xs''_v ++ [v]$  for some  $xs''_v$  and  $v$ . We use lemma 3 to update the abstract state to  $xs''_v$ , and split  $\text{All}(xs_v, \Psi)$  into  $\Psi(v)$  and  $\text{All}(xs''_v, \Psi)$ . Using this and the fragmental part, we close the invariant again. To finish, we prove the right disjunct of the consequent using the authoritative part and  $\Psi(v)$ .<sup>3</sup>

<sup>3</sup>Note here the importance of the  $\triangleright$  in the consequent. From the invariant, we technically only have  $\triangleright \Psi(v)$ . If we didn't have the later in the consequent, we would have to prove  $\Psi(v)$  from  $\triangleright \Psi(v)$  which isn't possible.

## Chapter 4

# The Two-Lock Michael-Scott Queue

In this chapter, we give an implementation of the blocking version of the M&S Queue, the Two-Lock M&S Queue, in HeapLang. This implementation differs slightly from the original, presented in Michael and Scott [1996], but most changes simply reflect the differences in the two languages.

### 4.1 Introduction

The underlying data structure making up the queue is a singly-linked list. The linked-list will always contain at least one element, called the *head* node, marking the beginning of the queue. Note that the head node is itself not part of the queue, but all nodes following it are. The queue keeps a head pointer ( $\ell_{\text{head}}$ ) which always points to the head node, and a tail pointer ( $\ell_{\text{tail}}$ ) which points to some node in the linked list, denoted the tail node.

In my implementation, a node is a triple  $(\ell_{i\_in}, w_i, \ell_{i\_out})$  satisfying that location  $\ell_{i\_in}$  points to the pair  $(w_i, \ell_{i\_out})$ . Here,  $w_i$  either contains the value of the node  $v_i$  wrapped in a *Some* (i.e.  $w_i = \text{Some } v_i$ ) or it is *None*.  $\ell_{i\_out}$  either points to *None* which represents the null pointer, or to the next node in the linked list. When we say that a location  $\ell$  points to a node  $(\ell_{i\_in}, w_i, \ell_{i\_out})$ , we mean that  $\ell \mapsto \ell_{i\_in}$ . Hence, if we have two adjacent nodes  $(\ell_{i\_in}, w_i, \ell_{i\_out})$ ,  $(\ell_{i+1\_in}, w_{i+1}, \ell_{i+1\_out})$  in the linked list, then we have the following structure:  $\ell_{i\_in} \mapsto (w_i, \ell_{i\_out})$ ,  $\ell_{i\_out} \mapsto \ell_{i+1\_in}$ , and  $\ell_{i+1\_in} \mapsto w_{i+1}, \ell_{i+1\_out}$ . For a given triple  $x = (\ell_{in}, w, \ell_{out})$ , we introduce the following notation:

$$\text{in}(x) = \ell_{in} \qquad \text{val}(x) = w \qquad \text{out}(x) = \ell_{out}$$

The reader may wonder why there is an extra, intermediary “in” pointer, between the pairs of the linked list, and why the “out” pointer couldn’t point directly to the next pair. This comes down to difference between HeapLang and the C-like language used in the original implementation [Michael and Scott, 1996]. Variables in the C-like language are technically just locations, and the assignment operator for variables simply corresponds to a store operation. In HeapLang, variables are modelled directly as locations which gives us an apparent extra pointer indirection compared to the original implementation.

### 4.2 Implementation

The queue consists of 3 functions: initialize, enqueue, and dequeue. Their implementation is shown in figure 4.4. As the name of the data structure suggests, the functions rely on two locks. To this end, we shall assume that we have some lock implementation given. In the accompanying Coq mechanisation, a “spin-lock” is used, but the only part we really care about is its specification which we discussed in section 2.2.7.

#### 4.2.1 Initialise

initialize will first create a single node – the head node – marking the start of the linked list. It then creates two locks,  $h_{\text{lock}}$  and  $t_{\text{lock}}$ , protecting the head and tail pointers, respectively. Finally, it creates




Figure 4.1: Queue after initialisation.

the head and tail pointers, both pointing to the head node. The queue is then a pointer to a structure containing the head and tail pointers, and the two locks.

Figure 4.1 illustrates the structure of the queue after initialisation. Note that one of the pointers is decorated with a square. This represents a *persistent* pointer; a pointer that will never be updated again. All “in” pointers  $\ell_{i\_in}$ , are persistent, meaning that, once created, they will only ever point to  $(w_i, \ell_{i\_out})$ . We shall use the notation  $\ell \mapsto^\square v$  (introduced in Vindum and Birkedal [2021]) to mean that  $\ell$  points to  $v$  persistently.

Note that in the original specification, a queue is a pointer to a 4-tuple  $(\ell_{\text{head}}, \ell_{\text{tail}}, h_{\text{lock}}, t_{\text{lock}})$ . Since HeapLang doesn’t support 4-tuples, we instead represent the queue as a pointer to a pair of pairs:  $((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}}))$ .

### 4.2.2 Enqueue

To enqueue a value, we must create a new node, append it to the underlying linked-list, and swing the tail pointer to this new node. These three operations are depicted in figure 4.2.

enqueue takes as argument the value to be enqueued and creates a new node containing this value (corresponding to figure 4.2a). This creation doesn’t interact with the underlying queue data-structure, hence why we don’t acquire  $t_{\text{lock}}$  first. After creating the new node, we must make the last node in the linked list point to it. Since this operation interacts with the queue, we first acquire  $t_{\text{lock}}$ . Once we obtain the lock, we make the last node in the linked list point to our new node (figure 4.2b). Following this, we swing  $\ell_{\text{tail}}$  to the newly inserted node (figure 4.2c).

Figure 4.2 also illustrates when pointers become persistent; once the previous last node is updated to point to the newly inserted node, that pointer will never be updated again, hence becoming persistent.

### 4.2.3 Dequeue

It is of course only possible to dequeue an element from the queue if the queue contains at least one element. Hence, the first thing dequeue does is check if the queue is empty. We can detect an empty queue by checking if the head node is the last node in the linked list. Being the last node in the linked list corresponds to having the “out” node be None. If this is the case, then the queue is empty and the function returns None. Otherwise, there is a node just after the head node, which is the first node of the queue. To dequeue it, we first read the associated value, and next we swing the head pointer to it, making it the new head node. Finally, we return the value we read.

Since all of these operations interact with the queue, we shall only perform them after having acquired  $h_{\text{lock}}$ .

Figure 4.3 illustrates running dequeue on a non-empty queue. Note that the only change is that the head pointer is swung to the next node in the linked list; the old head node is not deleted, it just becomes unreachable from the head pointer. In this way, the linked list only ever grows.

### 4.2.4 Observations on the Two-Lock Michael-Scott Queue

Now that we have seen the implementation, we point out the following noteworthy observations about the behaviour of the queue.

1. The tail node is always either the last or second last node in the linked list.
2. All but the last pointer in the linked list (the pointer to None) never change.
3. Nodes in the linked list are never deleted. Hence, the linked list only ever grows.

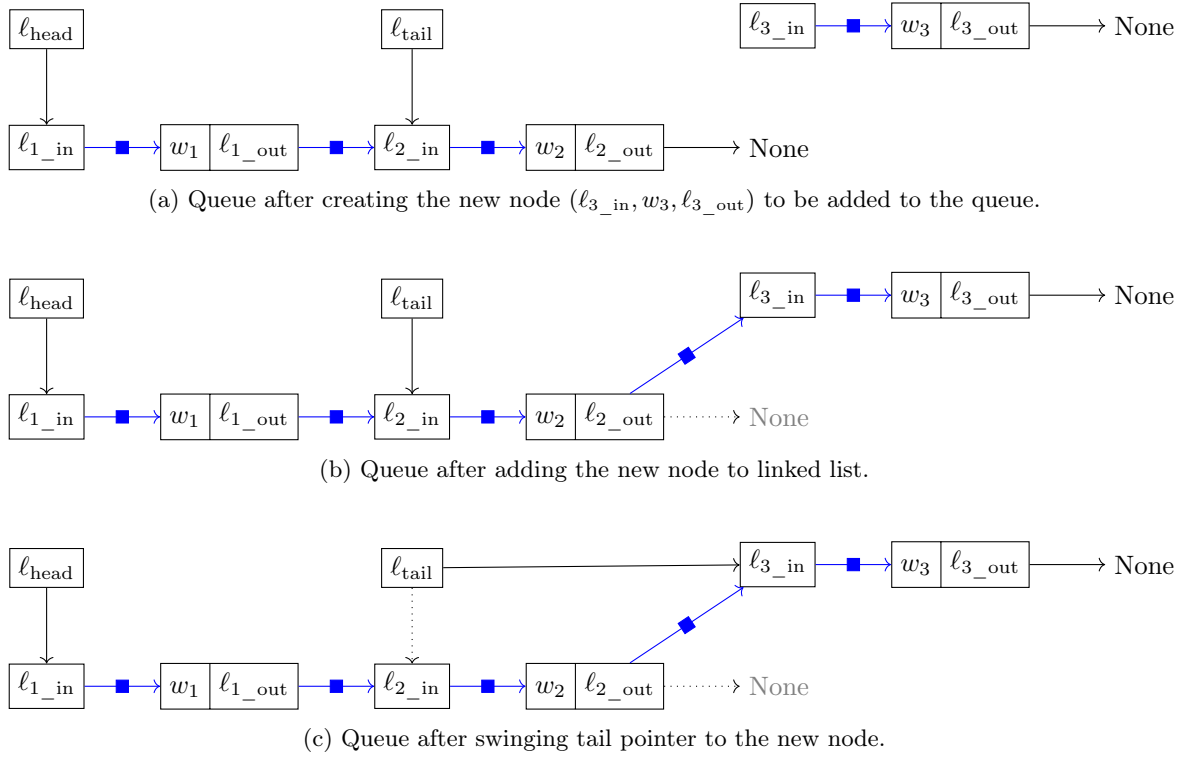


Figure 4.2: Enqueueing an element to a queue with one element. The illustrations assume that no dequeue is happening, hence  $\ell_{head}$  stays the same.




Figure 4.3: Dequeueing an element ( $w_2$ ) from a queue with two elements ( $w_2, w_3$ ). The dotted line represents the state before the dequeue, and the dashed line is the state after dequeueing.

```

1 initialize  $\triangleq$ 
2   let node = ref (None, ref (None)) in
3   let H_lock = newLock() in
4   let T_lock = newLock() in
5   ref ((ref (node), ref (node)), (H_lock, T_lock))
6
7 enqueue Q value  $\triangleq$ 
8   let node = ref (Some value, ref (None)) in
9   acquire(snd(snd(!Q)));
10  snd(! (snd(fst(!Q))))  $\leftarrow$  node;
11  snd(fst(!Q))  $\leftarrow$  node;
12  release(snd(snd(!Q)))
13
14 dequeue Q  $\triangleq$ 
15  acquire(fst(snd(!Q)));
16  let node = !(fst(fst(!Q))) in
17  let new_head = !(snd(!node)) in
18  if new_head = None then
19    release(fst(snd(!Q)));
20    None
21  else
22    let value = fst(!new_head) in
23    fst(fst(!Q))  $\leftarrow$  new_head;
24    release(fst(snd(!Q)));
25    value

```

Figure 4.4: Implementation of Two-Lock M&S Queue in HeapLang.

4. The tail can lag one node behind the head.
5. At any given time, the queue is in one of four states:
  - (a) No threads are interacting with the queue (**Static**)
  - (b) A thread is enqueueing (**Enqueue**)
  - (c) A thread is dequeuing (**Dequeue**)
  - (d) A thread is enqueueing and a thread is dequeuing (**Both**)

Observation 1 captures the fact that, while enqueueing, a new node is first added to the linked list, and then later the tail pointer is updated to point to the newly added node. Since only one thread can enqueue a node at a time (due to the lock), then the tail pointer will only ever point to the last or second last node. In a sequential setting, the tail will always appear to point to the last node, as no one can inspect the queue while the tail points to the second last.

Insight 2 means that we can mark all pointers in the queue, except the pointer to None, as persistent.


In the original implementation [Michael and Scott, 1996], nodes are freed after being dequeued. However, since our language, HeapLang, is a garbage-collected language, we do not explicitly free the nodes. This then leads to observation 3. From our perspective, the linked list only grows.

Observation 4 might seem a little surprising, and indeed it stands in contrast to property 5 in Michael and Scott [1996], which states “Tail always points to a node in the linked list, because it never lags behind Head, so it can never point to a deleted node.”. I also didn’t realise this possibility until a proof attempt using a model that “forgot” old nodes lead to an unprovable case (see section 5.3.4). The situation can occur when the queue is empty, and a thread performs an incomplete enqueue; it attaches the new node to the end, but before it can swing the tail to this new node, another thread performs a dequeue, which dequeues this new node, swinging the head to it. Now the tail is lagging one node behind the head. Figure 4.5 illustrates such a program trace.


It is not possible for the tail to point more than one node behind the head, as in order for this to happen, more nodes must be enqueued, but this can’t happen before the current enqueue finishes, which will update the tail and bring it up to speed with the head.

Fortunately, this isn’t an issue for safety, but a consequence of this possibility is that when modelling the queue, we must remember at least one “old” node (i.e. a dequeued node), as the tail might be pointing to this node. For the sake of simplicity in the model, the choice is made to remember an arbitrary amount of old nodes, which is represented by the list  $xs_{old}$ .


Observation 5 is a simple consequence of the implementation using two locks.



(a) The queue is initially empty, and an enqueueing thread has created a node  $(\ell_{2\_in}, w_2, \ell_{2\_out})$  that it wishes to append to the linked list.



(b) The thread executes line 10 which adds the created node to the linked list.



(c) Before the thread executes line 11 to swing the tail pointer  $\ell_{\text{tail}}$ , another thread dequeues the node that was just enqueued, which swings the head pointer,  $\ell_{\text{head}}$  to it. The tail is now lagging behind.

Figure 4.5: Illustrations of a scenario that makes the head lag behind the tail.

## Chapter 5

# Proving Specifications for the Two-Lock Michael-Scott Queue

### 5.1 Introduction

In this chapter we show that the Two-Lock M&S Queue satisfies the three queue specifications introduced in chapter 3. Section 3.5 of said chapter showed us that we get the sequential and concurrent specifications essentially for free, granted we can prove that Two-Lock M&S Queue satisfies the Hocap-style specification. However, for instructive purposes, we will in this chapter give the sequential and concurrent specifications directly, as the proof of the Hocap-style specification can get somewhat overwhelming.

In section 5.2 we prove that the Two-Lock M&S Queue satisfies the sequential specification. This also introduces some of the basic predicates we use for the queue. Next, in section 5.3, we prove the concurrent specification, which introduces some more complexities, such as invariant, and shows us how we can work with concurrent programs. Finally, section 5.4 shows how to modify the proof of the concurrent specification to prove the Hocap style specification.

### 5.2 Sequential Specification

#### 5.2.1 Sequential Queue Predicate

Since the queue uses two locks, we will need two ghost names; one for each lock (cf. section 2.2.7). We thus let  $SeqQgnames$  be the set of pairs containing ghost names. For an element  $G \in SeqQgnames$ , the first element of the pair, written  $G.\gamma_{Hlock}$ , will contain the ghost name for the head lock, and the second element,  $G.\gamma_{Tlock}$ , the ghost name for the tail lock.

To prove the specification, we must give a specific  $isQueue_S$  predicate. With the points we discussed in section 4.2.4 in mind, we give our definition of the queue predicate for the sequential specification and explain it afterwards.

**Definition 5.2.1** (Two-Lock M&S Queue -  $isQueue_S$  Predicate).

$$\begin{aligned} isQueue_S(v_q, xs_v, G) &\triangleq \exists \ell_{queue}, \ell_{head}, \ell_{tail} \in Loc. \exists h_{lock}, t_{lock} \in Val. \\ &v_q = \ell_{queue} * \ell_{queue} \mapsto^\square ((\ell_{head}, \ell_{tail}), (h_{lock}, t_{lock})) * \\ &\exists xs_{queue} \in List (Loc \times Val \times Loc). \exists x_{head}, x_{tail} \in (Loc \times Val \times Loc). \\ &projVal(xs_{queue}) = wrapSome(xs_v) * \\ &isLL(xs_{queue} ++ [x_{head}]) * \\ &\ell_{head} \mapsto in(x_{head}) * \\ &\ell_{tail} \mapsto in(x_{tail}) * isLast(x_{tail}, (xs_{queue} ++ [x_{head}])) * \\ &isLock(G.\gamma_{Hlock}, h_{lock}, True) * \\ &isLock(G.\gamma_{Tlock}, t_{lock}, True) \end{aligned}$$



This `isQueues` predicate states that the value  $v_q$  is a location, which persistently points to the structure containing the head and tail pointers, and the two locks. It also connects the abstract state  $xs_v$  with the concrete state by stating that if you strip away the locations in  $xs_{\text{queue}}$  (achieved by `projVal`; see appendix A.1.4) and wrap the values in the abstract state  $xs_v$  in `Some` (achieved by `wrapSome`; see appendix A.1.5), then the lists become equal.

Next, the predicate specifies the concrete state. There is some head node  $x_{\text{head}}$ , which the head points to. This head node and the nodes in  $xs_{\text{queue}}$  form the underlying linked list (specified using the `isLL` predicate below). There is also a tail node, which is the last node in the linked list, and the tail points to this node. The proposition `isLast( $x, xs$ )` simply asserts the existence of some  $xs'$ , so that  $xs = x :: xs'$  (defined formally in Appendix A.1.2).

Next, we have the `isLock` predicate for our two locks. Since we are in a sequential setting, the locks are superfluous, hence they simply protect `True`.

The `isLL` predicate essentially creates the structure seen in the examples of section 4.2. It is defined in two steps. Firstly, we create all the persistent pointers in the linked list using the `isLL_chain` predicate. Note that this in effect makes `isLL_chain( $xs$ )` persistent for all  $xs$ .

**Definition 5.2.2** (Linked List Chain Predicate).

$$\begin{aligned} \text{isLL\_chain}([]) &\equiv \text{True} \\ \text{isLL\_chain}([x]) &\equiv \text{in}(x) \mapsto^\square (\text{val}(x), \text{out}(x)) \\ \text{isLL\_chain}(x :: x' :: xs) &\equiv \text{in}(x) \mapsto^\square (\text{val}(x), \text{out}(x)) * \text{out}(x') \mapsto^\square \text{in}(x') * \text{isLL\_chain}(x' :: xs) \end{aligned}$$

Then, to define `isLL`, we add that the last node in the linked list points to `None`.

**Definition 5.2.3** (Linked List Predicate).

$$\begin{aligned} \text{isLL}([]) &\equiv \text{True} \\ \text{isLL}(x :: xs) &\equiv \text{out}(x) \mapsto \text{None} * \text{isLL\_chain}(x :: xs) \end{aligned}$$

For instance, if we wanted to capture the linked list in figure 4.2c, we would use the list  $xs = [(\ell_{3\_in}, w_3, \ell_{3\_out}); (\ell_{2\_in}, w_2, \ell_{2\_out}); (\ell_{1\_in}, w_1, \ell_{1\_out})]$ . `isLL( $xs$ )` will expand to  $\ell_{3\_out} \mapsto \text{None} * \text{isLL\_chain}(xs)$ , and `isLL_chain( $xs$ )` expands to

$$\begin{aligned} \ell_{3\_in} &\mapsto^\square (w_3, \ell_{3\_out}) * \ell_{2\_out} \mapsto^\square \ell_{3\_in} * \\ \ell_{2\_in} &\mapsto^\square (w_2, \ell_{2\_out}) * \ell_{1\_out} \mapsto^\square \ell_{2\_in} * \\ \ell_{1\_in} &\mapsto^\square (w_1, \ell_{1\_out}) \end{aligned}$$

Note how this matches the structure of the linked list in figure 4.2c.

The proofs require us manipulate specific `isLL` propositions quite a bit – appendix A.2 shows the specific lemmas we will use, but the proof outlines will generally not mention the lemmas explicitly.

## 5.2.2 Proof Outline

### Initialise

**Lemma 4** (Two-Lock M&S Queue Sequential Specification - Initialise).

$$\{\text{True}\} \text{initialise } () \{v_q, \exists G. \text{isQueues}(v_q, [], G)\}$$

*Proof.* Proving the initialise specification amounts to stepping through the code, giving us the required resources, and then using these to create an instance of `isQueues` with the obtained resources. To begin with, we step through line 2 which creates the first node  $x_1 = (\ell_{1\_in}, \text{None}, \ell_{1\_out})$  with  $\ell_{1\_out} \mapsto \text{None}$  and  $\ell_{1\_in} \mapsto (\text{None}, \ell_{1\_out})$ . We can then update the latter points-to proposition to become persistent, giving us  $\ell_{1\_in} \mapsto^\square (\text{None}, \ell_{1\_out})$ . Next, on lines 3 and 4, we create the two locks. We use the specification for `newLock` for each of its invocations. Both times, we specify that the lock should protect `True`. This gives us two ghost names,  $\gamma_{\text{Hlock}}, \gamma_{\text{Tlock}}$ , which we will collect in a *SeqQnames* pair,  $G$ . Finally, we step through the allocations of the head, tail, and queue pointers on line 5. This gives us locations  $\ell_{\text{head}}, \ell_{\text{tail}}, \ell_{\text{queue}}$ , such that both  $\ell_{\text{head}}$  and  $\ell_{\text{tail}}$  point to node  $x_1$ , and such that  $\ell_{\text{queue}}$  points to the structure

containing the head and tail pointers, and the two locks. This last points to predicate we update to become persistent. With this, we now have all the resources needed to prove the post-condition. Proving this follows by a sequence of framing away the resources we obtained and instantiating existentials with the values we got above. Most noteworthy, we pick the empty list for  $xs_{\text{queue}}$ , and node  $x_1$  for  $x_{\text{head}}$  and  $x_{\text{tail}}$ .  $\square$

## Enqueue

**Lemma 5** (Two-Lock M&S Queue Sequential Specification - Enqueue).

$$\forall v_q, v, xs_v, G. \{ \text{isQueues}(v_q, xs_v, G) \} \text{ enqueue } v_q \ v \ \{ w. \text{isQueues}(v_q, (v :: xs_v), G) \}$$

*Proof.* We assume the pre-condition  $\text{isQueues}(v_q, xs_v, G)$  which tells us that  $v_q$  is some location  $\ell_{\text{queue}}$ , and in particular, we have following:

$$\ell_{\text{queue}} \mapsto^\square ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) \quad (5.1)$$

$$\text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) \quad (5.2)$$

$$x_{\text{tail}} \mapsto \text{None} * \text{isLL\_chain}(xs_{\text{queue}} ++ [x_{\text{head}}]) \quad (5.3)$$

$$\ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, (xs_{\text{queue}} ++ [x_{\text{head}}])) \quad (5.4)$$

$$\text{isLock}(G.\gamma_{\text{Tlock}}, t_{\text{lock}}, \text{True}) \quad (5.5)$$

We proceed to step into the enqueue function. Firstly, on line 8, we create a new node  $x_{\text{new}}$ , with  $\text{val}(x_{\text{new}}) = \text{Some } v$ . Next, on line 9 we acquire the tail lock. We step through the dereference and the projections using 5.1. To acquire the lock we use the acquire specification with 5.5. This gives us  $\text{locked}(Qg.\gamma_{\text{Tlock}})$  and  $\text{True}$ .

Line 10 adds node  $x_{\text{new}}$  to the linked list. We first use 5.1 to step to the dereference of  $\ell_{\text{tail}}$ . From 5.4, we know the dereference results in  $\text{in}(x_{\text{tail}})$ . Because  $x_{\text{tail}}$  is in the linked list, we can use 5.3 to conclude that  $x_{\text{tail}}$  is a node, and hence we can perform the load and the projection to get to the final store operation:  $\text{out}(x_{\text{tail}}) \leftarrow \text{in}(x_{\text{new}})$ . Using the points-to proposition from 5.3, we perform the store, which in turn gives us  $\text{out}(x_{\text{tail}}) \mapsto \text{in}(x_{\text{new}})$ . We make this persistent and combine it with the  $\text{isLL\_chain}$  proposition from 5.3 to conclude  $\text{isLL}(x_{\text{new}} :: xs_{\text{queue}} ++ [x_{\text{head}}])$ .

The next line (line 11) swings the tail pointer to  $x_{\text{new}}$ . 5.4 and 5.4 gives us all the required resources to step through the code and perform the store. Afterwards, we get  $\ell_{\text{tail}} \mapsto \text{in}(x_{\text{new}})$ .

Finally, we get to line 12 which releases the lock. We use the specification for release, giving up  $\text{locked}(Qg.\gamma_{\text{Tlock}})$ . The only thing left is to prove the postcondition:  $\text{isQueues}(v_q, (v :: xs_v), G)$ . For the existentials, we pick  $x_{\text{new}} :: xs_{\text{queue}}$  as the queue, and as the tail node, we chose  $x_{\text{new}}$ . The remaining choices are similar to those we got from the pre-condition. With these choices, the remaining proof obligations become straightforward; we already have the required pointers and the  $\text{isLL}$  proposition. The relationship between the abstract and concrete states follows from  $\text{val}(x_{\text{new}}) = \text{Some } v$  and 5.2.  $\square$

## Deque

**Lemma 6** (Two-Lock M&S Queue Sequential Specification - Dequeue).

$$\begin{aligned} & \forall v_q, xs_v, G. \{ \text{isQueues}(v_q, xs_v, G) \} \\ & \text{ dequeue } v_q \\ & \left\{ w. \begin{aligned} & (xs_v = [] * w = \text{None} * \text{isQueues}(v_q, xs_v, G)) \vee \\ & (\exists v, xs'_v. xs_v = xs'_v ++ [v] * w = \text{Some } v * \text{isQueues}(v_q, xs'_v, G)) \end{aligned} \right\} \end{aligned}$$

*Proof.* We assume the pre-condition  $\text{isQueues}(v_q, xs_v, G)$  which gives us that  $v_q$  is some location  $\ell_{\text{queue}}$ , and the following propositions.

$$\ell_{\text{queue}} \mapsto^\square ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) \quad (5.6)$$

$$\text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) \quad (5.7)$$

$$x_{\text{tail}} \mapsto \text{None} * \text{isLL\_chain}(xs_{\text{queue}} ++ [x_{\text{head}}]) \quad (5.8)$$

$$\ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \quad (5.9)$$

$$\text{isLock}(G.\gamma_{\text{Hlock}}, h_{\text{lock}}, \text{True}) \quad (5.10)$$

We perform the function application and step into the function. We first do the superfluous acquire on line 15 using the acquire specification and 5.10. This gives us  $\text{locked}(G.\gamma_{\text{Hlock}})$ .

Next we step to line 16 which dereferences the head node. We perform the loads and projections using 5.6 and 5.9, which tells us that the *node* variable becomes  $\text{in}(x_{\text{head}})$ . From 5.8, we know that  $x_{\text{head}}$  is in the linked list, hence it must be a node.

We step to line 17 which finds out what  $x_{\text{head}}$  points to. As  $x_{\text{head}}$  is a node, we can step to the load:  $!(\text{out}(x_{\text{head}}))$ . The result of this load is either *None* or another node  $x_{\text{head\_next}}$ , depending on whether or not  $xs_{\text{queue}}$  is empty. So we consider both cases.

**Case  $xs_{\text{queue}}$  is empty:** In this case, 5.8 simply asserts  $\text{isLL}[x_{\text{head}}]$ , which by definition tells us that  $x_{\text{head}} \mapsto \text{None}$ . Hence, the “if” on line 18 will take the “then” branch, so we step to line 19. Here we release the lock, giving up  $\text{locked}(G.\gamma_{\text{Hlock}})$ , and return *None* on the next line. What remains is to prove the post-condition with  $w = \text{None}$ . We can easily do this by proving the first disjunction. From 5.7 with the fact that  $xs_{\text{queue}} = []$  we can conclude that  $xs_v$  is empty, and since we haven’t modified the queue, we can prove  $\text{isQueues}(v_q, xs_v, G)$  using the same resources we got from the pre-condition.

**Case  $xs_{\text{queue}}$  is not empty:** In this case, we can conclude that there must be some node  $x_{\text{head\_next}}$ , which is the first node in  $xs_{\text{queue}}$ . I.e.

$$xs_{\text{queue}} = xs'_{\text{queue}} ++ [x_{\text{head\_next}}] \quad (5.11)$$

We can thus use the  $\text{isLL}$  predicate to conclude that  $x_{\text{head}}$  must point to  $x_{\text{head\_next}}$ . Hence the “else” branch will be taken, so we step to line 25. Since  $x_{\text{head\_next}}$  is part of the linked list, it must be a node, allowing us to extract its value  $\text{val}(x_{\text{head\_next}})$  in the first line of the else branch.

We step to line 23 which swings the head pointer,  $\ell_{\text{head}}$  to  $x_{\text{head\_next}}$ . We perform these steps using 5.6 and 5.9, which then gives us  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head\_next}})$ .

Finally, we release the lock on line 24, giving up  $\text{locked}(G.\gamma_{\text{Hlock}})$  and return the value:  $\text{val}(x_{\text{head\_next}})$ . We must now prove the post-condition with  $w = \text{val}(x_{\text{head\_next}})$ , and this time we prove the second disjunct. From 5.7 and 5.11 we can deduce

$$xs_v = xs'_v ++ [v] \quad (5.12)$$

$$\text{projVal}(xs'_{\text{queue}}) = \text{wrapSome}(xs'_v) \quad (5.13)$$

$$\text{val}(x_{\text{head\_next}}) = \text{Some } v \quad (5.14)$$

5.12 and 5.14 proves the first part of the post-condition. What remains is to show  $\text{isQueues}(v_q, xs'_v, G)$ . For the existentials, we pick  $xs'_{\text{queue}}$  as the queue and  $x_{\text{head\_next}}$  as the head node. The rest are similar to the variables we got from the pre-condition. With these choices, we can prove the predicate using the resources we have established.

□

## 5.3 Concurrent Specification

### 5.3.1 Concurrent Queue Predicate

For the concurrent specification, we need to allow for multiple threads to access the queue resources to perform the queue operations, as showcased in the sequential case. The concurrent specification enforces this possibility by asserting that the queue predicate,  $\text{isQueue}_C$ , is persistent, hence duplicable. The resources needed by the queue, however, are not persistent. The solution is to collect the required resources in an *invariant* which *is* persistent and can hence be given to multiple threads. The invariant we define here additionally describes the concrete state of the queue. In the proofs of the queue functions, we shall then access the required resources through the invariant. We now present the invariant and explain it afterwards.

**Definition 5.3.1** (Two-Lock M&S Queue Concurrent Invariant).

$$\begin{aligned}
& I_{\text{TLC}}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G) \triangleq \\
& \exists xs_v. \text{All}(xs_v, \Phi) * \quad \text{(abstract state)} \\
& \exists xs, xs_{\text{queue}}, xs_{\text{old}}, x_{\text{head}}, x_{\text{tail}}. \quad \text{(concrete state)} \\
& xs = xs_{\text{queue}} ++ [x_{\text{head}}] ++ xs_{\text{old}} * \\
& \text{isLL}(xs) * \\
& \text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) * \\
& ( \\
& \quad \ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, xs) * \quad \text{(Static)} \\
& \quad \text{TokNE } G * \text{TokND } G * \text{TokUpdated } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \frac{1}{2} \text{in}(x_{\text{tail}}) * \quad \text{(Enqueue)} \\
& \quad (\text{isLast}(x_{\text{tail}}, xs) * \text{TokBefore } G \vee \text{isSndLast}(x_{\text{tail}}, xs) * \text{TokAfter } G) * \\
& \quad \text{TokE } G * \text{TokND } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto \frac{1}{2} \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, xs) * \quad \text{(Dequeue)} \\
& \quad \text{TokNE } G * \text{TokD } G * \text{TokUpdated } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto \frac{1}{2} \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \frac{1}{2} \text{in}(x_{\text{tail}}) * \quad \text{(Both)} \\
& \quad (\text{isLast}(x_{\text{tail}}, xs) * \text{TokBefore } G \vee \text{isSndLast}(x_{\text{tail}}, xs) * \text{TokAfter } G) * \\
& \quad \text{TokE } G * \text{TokD } G \\
& )
\end{aligned}$$

In contrast to the sequential specification, the abstract state is now existentially quantified, hence the exact contents of the queue are not tracked. Instead, we have added the proposition  $\text{All}(xs_v, \Phi)$ , which states that all values in  $xs_v$  (i.e. the values currently in the queue) satisfy the predicate  $\Phi$  (see appendix A.1.6 for formal definition). This will allow us to conclude that dequeued values satisfy  $\Phi$ .

The concrete state of the queue is still reflected in the abstract state through projecting out the values of the nodes (via  $\text{projVal}$ ), and wrapping the values in the queue in  $\text{Some}$  (via  $\text{wrapSome}$ ). Another difference is that we now also keep track of an arbitrary number of “old” nodes; nodes that are behind the head node,  $x_{\text{head}}$ . As discussed above, this inclusion is due to observation 4.

As before, we also assert that the concrete state forms a linked list, as described by the  $\text{isLL}$  predicate. The final part of the invariant describes the four possible states of the queue, as described in 5. Since the resources used by the queue are inside an invariant, and enqueueing/dequeueing threads need to access the resources of the queue multiple times (as shown in the sequential specification), then we will have to open and close the invariant multiple times. Each time we open the invariant, the existentially quantified variables will not be the same as those from early accesses of the invariant (as they are existentially quantified). Thus, the threads must be able to “match up” variables from previous accesses to later accesses. The way we shall achieve this is by allowing threads to keep a *fraction* of the points-to predicate that it is using. For instance, an enqueueing thread will have to access the points-to predicate concerning  $\ell_{\text{tail}}$  multiple times, and in between accesses of the invariant, it can get to keep half of the points-to predicate. Thus, when it opens the invariant later, it will have  $\ell_{\text{tail}} \mapsto \frac{1}{2} \text{in}(x_{\text{tail}})$  from an earlier access, and it will obtain the existence of some new  $x'_{\text{tail}}$ , such that  $\ell_{\text{tail}} \mapsto \frac{1}{2} \text{in}(x'_{\text{tail}})$ . Combining the two points-to predicates allows us to conclude that  $\text{in}(x_{\text{tail}}) = \text{in}(x'_{\text{tail}})$  and using lemma 33 we can further conclude  $x_{\text{tail}} = x'_{\text{tail}}$ . In this way, we can match up variables from earlier accesses to variables in later accesses. In the **Static** state where no thread is interacting with the queue, the queue owns all of the points-to predicates concerning the head and tail.

In the **Enqueue** state, the enqueueing thread owns half of the tail pointer, and we distinguish between two cases, as discussed in 1: either the enqueueing thread has yet to add the new node to the linked list and  $x_{\text{tail}}$  is still the last node, or the new node has been added, but the tail pointer hasn’t been updated,

meaning that  $x_{\text{tail}}$  is the second last node ( $\text{isSndLast}$  is defined similarly to  $\text{isLast}$ ; see appendix A.1.3). In the **Dequeue** state, the dequeueing thread owns half of the head pointer, and the tail is as in the **Static** state.

Finally, the **Both** state is essentially a combination of the **Enqueue** and **Dequeue** states.

To track which state the queue is in, we use *tokens*. Importantly, tokens are non-duplicable and owning two of the same token allows us to deduce *false*. Thus, if we own a particular token, then, upon opening the invariant, we can rule out certain states simply because they mention the token we own. We refer to section 8.1 for a presentation of the resource algebra used to define tokens.

We will use several tokens, each associated with their own ghost name. We collect all these new ghost names in *ConcQghostnames*, which then contains the ghost names for the two locks and the ghost names for the tokens.

To refer to a specific token associated with a tuple  $G$ , we write e.g.  $\text{TokE } G$  which projects out the appropriate ghost name and asserts ownership of the token associated with it. We proceed to explain the meaning of each of the tokens used in the invariant.

- $\text{TokNE } G$  represents that no threads are enqueueing.
- $\text{TokE } G$  represents that a thread is enqueueing.
- $\text{TokND } G$  represents that no threads are dequeueing.
- $\text{TokD } G$  represents that a thread is dequeueing.
- $\text{TokBefore } G$  represents that an enqueueing thread has not yet added the new node to the linked list.
- $\text{TokAfter } G$  represents that an enqueueing thread has added the new node to the linked list, but not yet swung the tail.
- $\text{TokUpdated } G$  is defined as  $\text{TokBefore } G * \text{TokAfter } G$ , and represents that the queue is up to date.

We note that the concurrent specification for the Two-Lock M&S Queue *can* be proven using queue invariant 5.3.1 directly, and the proof outline below will also be using this. However, a simpler (but arguably less intuitive) queue invariant was discovered. This simpler invariant is equivalent to 5.3.1 and has the benefit of being easier to work with in the mechanised proofs. Thus, in the mechanised proofs, we always rewrite to the simpler invariant when opening the invariant. We define it as follows.

**Definition 5.3.2** (Simplified Two-Lock M&S Queue Invariant).

$$\begin{aligned}
& I'_{\text{TLC}}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G) \triangleq \\
& \exists x_{s_v}. \text{All}(x_{s_v}, \Phi) * \quad \quad \quad \text{(abstract state)} \\
& \exists x_s, x_{s_{\text{queue}}}, x_{s_{\text{old}}}, x_{\text{head}}, x_{\text{tail}}. \quad \quad \quad \text{(concrete state)} \\
& x_s = x_{s_{\text{queue}}} ++ [x_{\text{head}}] ++ x_{s_{\text{old}}} * \\
& \text{isLL}(x_s) * \\
& \text{projVal}(x_{s_{\text{queue}}}) = \text{wrapSome}(x_{s_v}) * \\
& ((\ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \text{TokND } G) \vee (\ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{head}}) * \text{TokD } G)) * \\
& ( \\
& \quad (\ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, x_s) * \text{TokNE } G * \text{TokUpdated } G) \vee \\
& \quad ( \\
& \quad \quad \ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}}) * \text{TokE } G * \\
& \quad \quad ((\text{isLast}(x_{\text{tail}}, x_s) * \text{TokBefore } G) \vee (\text{isSndLast}(x_{\text{tail}}, x_s) * \text{TokAfter } G)) \\
& \quad ) \\
& )
\end{aligned}$$

As is evident, it contains the same propositions as the original invariant, but simply in a different structure. This structure allows for more linear reasoning; the common propositions between the different queue states are now “grouped” together.

With this, we now give our definition of the queue predicate,  $\text{isQueue}_C$ . In the below, we let  $\mathcal{N}$  be some namespace.

**Definition 5.3.3** (Two-Lock M&S Queue -  $\text{isQueue}_C$  Predicate).

$$\begin{aligned} \text{isQueue}_C(\Psi, v_q, G) \triangleq & \exists \ell_{\text{queue}}, \ell_{\text{head}}, \ell_{\text{tail}} \in \text{Loc}. \exists h_{\text{lock}}, t_{\text{lock}} \in \text{Val}. \\ & v_q = \ell_{\text{queue}} * \ell_{\text{queue}} \mapsto^\square ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) * \\ & \boxed{\text{ITLC}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.\text{queue}} * \\ & \text{isLock}(G.\gamma_{\text{Hlock}}, h_{\text{lock}}, \text{TokD } G) * \\ & \text{isLock}(G.\gamma_{\text{Tlock}}, t_{\text{lock}}, \text{TokE } G). \end{aligned}$$

In contrast to our definition of  $\text{isQueues}$ , the locks now protect  $\text{TokE } G$  and  $\text{TokD } G$ . The idea is that, when an enqueueing thread obtains  $t_{\text{lock}}$ , they will obtain the  $\text{TokE } G$  token, which allows them to conclude that the queue state is either **Static** or **Dequeue**. Similarly for a dequeuing thread.

### 5.3.2 Linearisation Points

An important notion for concurrent algorithms is *linearisability* [Herlihy and Wing, 1990]. Linearisability is a non-blocking property<sup>1</sup> that helps reason about which behaviours are possible. Both versions of the M&S Queue are linearisable, which rules out undesired behaviours.

One way to characterise linearisability is through the concept of *linearisation points*. We say that a function (such as enqueue or dequeue) is linearisable, if for all invocations of the function, there is a specific point in time between the invocation and the response where the effect of the function appears to takes place; before that point, the effect hasn’t taken place, and nothing needs to be done afterwards to complete the effect. We call such a point a linearisation point.

For example, enqueue has a single linearisation point at the instruction that appends the newly created node to the linked list (the store on line 10). At exactly that point, the effect of the enqueue takes place. dequeue is slightly more complicated as it has multiple linearisation points. If the queue is empty when we read the head node’s out pointer on line 17, then at that read, the dequeue function is guaranteed to return `None`, interpreting the queue as empty. In this case the dequeue doesn’t change the queue but merely observes it, and the observation happens exactly at the read of the head node’s out pointer, making it the linearisation point.

On the hand, if the queue was not empty at that point, then the linearisation point is at line 23, when we swing the head pointer. At that very store operation, the effect of dequeue occurs.

Linearisation points are closely tied to updates of the abstract state of the queue. The abstract state of the queue changes only at linearisation points. This is consistent with the notion that the effects of the function takes effect at the linearisation points – updates to the abstract state happen atomically as we would intuitively want it to. This link to the abstract state becomes even more prevalent when we prove Hocap-style specifications in sections 5.4 and section 7.5.

### 5.3.3 Proof Outline

Firstly, we must show that  $\text{isQueue}_C$  is persistent. This however follows from the fact that invariants are persistent, the  $\text{isLock}$  predicates are persistent, persistent points-to predicates are persistent, and persistency is preserved by  $*$  and quantifications (rules: persistently-sep, persistently- $\wedge$ , persistently- $\exists$ ).

The proofs of the three specifications largely have the same structure as the sequential counterparts. The major difference is that we don’t have access to the resources all the time; we must get them from the invariant. Further we also have to keep track of which state we are in. For the proof outlines below, these points will be the main focus.

<sup>1</sup>Other properties, such as serialisability, inherently require implementations to be blocking. Linearisability does not require this, but linearisable algorithms can of course still be blocking.

## Initialise

**Lemma 7** (Two-Lock M&S Queue Concurrent Specification - Initialise).

$$\{\text{True}\} \text{ initialize } () \{v_q. \exists G. \text{isQueue}_C(\Phi, v_q, G)\}$$

*Proof.* We first step through line 2 which gives us the head node of the linked list. Following this, we must create the two locks. As the locks had to protect tokens, we must first create these. To create the two tokens, we use the ghost-alloc $\blacktriangleright\text{ref}\blacktriangleleft$  rule twice, which gives us two ghost names, one for each of the tokens. We put the ghost names into a tuple  $G$ , and write TokE  $G$  and TokD  $G$  for the two ghost resources created by the ghost-alloc $\blacktriangleright\text{ref}\blacktriangleleft$  rule. We then use the newLock specification to step through the code and create the locks, giving up the two tokens. Next, we step to line 5 and create the  $\ell_{\text{queue}}$ ,  $\ell_{\text{head}}$ , and  $\ell_{\text{tail}}$  pointers with  $\ell_{\text{queue}} \mapsto ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}}))$ . We make this persistent.

All that remains then is to prove the postcondition; the  $\text{isQueue}_C$  predicate. We have both the points-to predicate of  $\ell_{\text{queue}}$  and the two  $\text{isLock}$  predicates (from the newLock specification). So all that remains is the invariant.

We create  $\text{I}_{\text{TLC}}$  in the **Static** state, most of which is analogous to the sequential specification. However, we will also need to supply the tokens required by the **Static** state. Thus, we allocate the four tokens TokNE  $G$ , TokND  $G$ , TokBefore  $G$ , and TokAfter  $G$  in the same way we allocated TokE  $G$  and TokD  $G$ . We combine TokBefore  $G$  and TokAfter  $G$  to get TokUpdated  $G$ , and we now have all the tokens we need to create  $\text{I}_{\text{TLC}}$  in the **Static** state. To create the invariant from  $\text{I}_{\text{TLC}}$ , we use the Inv-alloc rule $\blacktriangleright\text{ref}\blacktriangleleft$ .  $\square$

## Enqueue

**Lemma 8** (Two-Lock M&S Queue Concurrent Specification - Enqueue).

$$\forall v_q, v, G. \{\text{isQueue}_C(\Phi, v_q, G) * \Phi(v)\} \text{ enqueue } v_q \ v \ \{w. \text{True}\}$$

*Proof.* We assume the pre-condition, which tells us that  $v_q$  is a location  $\ell_{\text{queue}}$ , and we have:

$$\Phi(v) \tag{5.15}$$

$$\ell_{\text{queue}} \mapsto^\square ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) \tag{5.16}$$

$$\boxed{\text{I}_{\text{TLC}}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} \tag{5.17}$$

$$\text{isLock}(G.\gamma_{\text{Tlock}}, t_{\text{lock}}, \text{TokE } G) \tag{5.18}$$

We perform the function application and step into enqueue to line 8. We create the new node as before, giving us  $x_{\text{new}}$ , with  $\text{val}(x_{\text{new}}) = \text{Some } v$ .

On line 9 we then use the acquire specification with 5.18 to acquire  $\text{locked}(G.\gamma_{\text{Tlock}})$  and TokE  $G$ .

Following this, we step to line 10. Using 5.16 we can step to the load of  $\ell_{\text{tail}}$ . This is where we meet our first challenge; in order to perform the store, we must open the invariant to access the points-to predicate regarding  $\ell_{\text{tail}}$ .

As invariants can only be opened if the expression being considered is atomic, we use a bind rule to “focus” on the load of  $\ell_{\text{tail}}$ . We proceed to open the invariant, and since we have TokE  $G$ , we know that the queue is in state **Static** or **Dequeue**. In any case, we get that  $\ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}})$ , for some  $x_{\text{tail}}$  which is the last node in the linked list. As  $x_{\text{tail}}$  is in the linked list, we further deduce that it is a node.

We can now perform the load of  $\ell_{\text{tail}}$  which results in  $\text{in}(x_{\text{tail}})$ . We must now close the invariant. We split up the points-to predicate  $\ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}})$  in two, which leaves us with two of  $\ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}})$ . We keep one of them, and use the other to close the invariant in the before case of state **Enqueue** or **Both**, depending on which state we opened the invariant into. By doing this, we give up TokE  $G$ , but we gain TokNE  $G$  and TokAfter  $G$ .

As  $x_{\text{tail}}$  was a node we can step to  $\text{out}(x_{\text{tail}}) \leftarrow \text{in}(x_{\text{new}})$ . However, the points-to predicate concerning  $\text{out}(x_{\text{tail}})$  isn't persistent, and is hence inside the invariant. We thus have to open the invariant again. Since we have TokNE  $G$  and TokAfter  $G$ , we know that we are in the before case of either state **Enqueue** or **Both**.

The invariant hence gives us  $\ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x'_{\text{tail}})$  for some  $x'_{\text{tail}}$ . However, since we kept  $\ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}})$ , we can combine these, allowing us to conclude that  $\text{in}(x_{\text{tail}}) = \text{in}(x'_{\text{tail}})$ . As both  $x_{\text{tail}}$  and  $x'_{\text{tail}}$  are nodes, we apply 33 to conclude  $x_{\text{tail}} = x'_{\text{tail}}$ . This now gives us that  $\text{out}(x_{\text{tail}}) \mapsto \text{None}$ , allowing us to perform the

store. Thus,  $x_{\text{new}}$  is added to the linked list, and this is a linearisation point. As such we must update the abstract state to reflect the change. We do this by closing the invariant with  $(v :: xs_v)$  as the abstract state, where  $xs_v$  is the abstract state we got when we opened the invariant. Note that we have  $\Phi(v)$  (from 5.15), hence we will be able to conclude  $\text{All}((v :: xs_v), \Phi)$ . For the concrete state, we pick  $x_{\text{new}} :: xs$ , where  $xs$  is the concrete state we got when we opened the invariant. We close the invariant in the after case of either state **Enqueue** or **Both**, giving up TokAfter  $G$ , and obtaining TokBefore  $G$ .

We step to line 11, which swings the tail pointer to  $x_{\text{new}}$ . Using 5.16 we step to  $\ell_{\text{tail}} \leftarrow \text{in}(x_{\text{new}})$ . However, to perform this store we must first know that  $\ell_{\text{tail}}$  points to something. This resource is inside the invariant, so we open the invariant one last time. Due to having TokNE  $G$  and TokBefore  $G$ , we know that we are in the after case of state **Enqueue** or **Both**. This time we get  $\ell_{\text{tail}} \mapsto^{\frac{1}{2}} x''_{\text{tail}}$  for some  $x''_{\text{tail}}$ , where  $x''_{\text{tail}}$  is the *second* last node in the linked list. Hence there is some other node  $x'_{\text{new}}$ , which is the last node, with  $x''_{\text{tail}}$  pointing to it. As before, we use our half of the points-to predicate of  $\ell_{\text{tail}}$  (i.e.  $\ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}})$ ) to get that  $x''_{\text{tail}} = x_{\text{tail}}$ . Since  $x_{\text{tail}}$  points to  $x_{\text{new}}$ , and  $x''_{\text{tail}}$  points to  $x'_{\text{new}}$ , we can further conclude that  $x_{\text{new}} = x'_{\text{new}}$ . Thus, we can perform the store, which now gives us that  $\ell_{\text{tail}}$  points to  $x_{\text{new}}$ ; the last node in the linked list. With this, we can close the invariant in state **Static** or **Dequeue**, giving up TokNE  $G$  and TokUpdated  $G$ , and getting TokE  $G$ . Finally, on line 12, we release the lock which we can do since we have TokE  $G$  and locked( $G.\gamma_{\text{Tlock}}$ ). The postcondition merely asserts True, so there is nothing left to prove.  $\square$

## Dequeue

**Lemma 9** (Two-Lock M&S Queue Concurrent Specification - Dequeue).

$$\forall v_q, G. \{\text{isQueue}_C(\Phi, v_q, G)\} \text{ dequeue } v_q \{w.w = \text{None} \vee (\exists v. w = \text{Some } v * \Phi(v))\}$$

*Proof.* As usual, we assume the pre-condition giving us that  $v_q$  is a location  $\ell_{\text{queue}}$  with

$$\ell_{\text{queue}} \mapsto^{\square} ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) \quad (5.19)$$

$$\boxed{\text{ITLC}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} \quad (5.20)$$

$$\text{isLock}(G.\gamma_{\text{Hlock}}, h_{\text{lock}}, \text{True}) \quad (5.21)$$

We do the function application and step into dequeue. First, on line 15, we acquire the lock, which gives us locked( $G.\gamma_{\text{Hlock}}$ ) and TokD  $G$ .

Next, we step to line 16 and using 5.19 we get to  $!(\ell_{\text{head}})$ . To perform this load, we must open the invariant. We open it in state **Static** or **Enqueue** (as we have TokD  $G$ ), which gives us  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head}})$  for some  $x_{\text{head}}$ . The isLL predicate further tells us that  $x_{\text{head}}$  is a node:  $\text{in}(x_{\text{head}}) \mapsto^{\square} (\text{val}(x_{\text{head}}), \text{out}(x_{\text{head}}))$ . We perform the load, and take half of the points-to predicate:  $\ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{head}})$ . We use the other half to close the invariant in state **Dequeue** or **Both**, giving up TokD  $G$ , but obtaining TokND  $G$ .

We proceed to line 17, which finds out what  $x_{\text{head}}$  is pointing to. As  $x_{\text{head}}$  is a node we can step to  $!(\text{out}(x_{\text{head}}))$ . To perform this dereference, we must open the invariant. As we own TokND  $G$ , the queue must be in state **Dequeue** or **Both**. In any case, we get that there is some  $x'_{\text{head}}$  with  $\ell_{\text{head}} \mapsto^{\frac{1}{2}} x'_{\text{head}}$ . Using the fractional points-to predicate we kept from earlier and lemma 33, we can conclude that  $x'_{\text{head}} = x_{\text{head}}$ . We now perform a case analysis on the contents of the queue:  $xs_{\text{queue}}$ , similarly to the sequential proof.

**Case  $xs_{\text{queue}}$  is empty:** In this case, we use the isLL predicate to conclude  $\text{out}(x_{\text{head}}) \mapsto \text{None}$ . The expression  $!(\text{out}(x_{\text{head}}))$  hence resolves to None. At this point, we know dequeue will decide that the queue is empty, so this is a linearisation point. We close the invariant in state **Static** or **Enqueue**, giving up TokND  $G$  and obtaining TokD  $G$ .

As *new\_head* was set to None, the “if” on line 18 takes the “then” branch, so we step to line 19. We release the lock, giving up TokD  $G$  and locked( $G.\gamma_{\text{Hlock}}$ ), and return None. We must now prove the post-condition with  $w = \text{None}$ . We easily prove the first disjunct.

**Case  $xs_{\text{queue}}$  is not empty:** As in the sequential proof, we conclude that  $x_{\text{head}}$  points to  $x_{\text{head\_next}}$  for



some  $x_{\text{head\_next}}$ . That is, we have the following:

$$\text{out}(x_{\text{head}}) \mapsto \text{in}(x_{\text{head\_next}}) \quad (5.22)$$

$$\text{in}(x_{\text{head\_next}}) \mapsto^\square (\text{val}(x_{\text{head\_next}}), \text{out}(x_{\text{head\_next}})) \quad (5.23)$$

We perform the dereference, which resolves to  $\text{in}(x_{\text{head\_next}})$ . We close the invariant in **Dequeue** or **Both**, meaning we still have TokND  $G$  and  $\ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{head}})$ .

This time, the “if” will take the else branch, so we step to line 25. Using 5.23, we conclude that the return value will be  $\text{val}(x_{\text{head\_next}})$ .

Next, we step to line 23, which swings the head pointer to  $x_{\text{head\_next}}$ . Using 5.19, we step to  $\ell_{\text{head}} \leftarrow \text{in}(x_{\text{head\_next}})$ . Performing this store requires a points-to proposition for  $\ell_{\text{head}}$ . Hence, we open the invariant in state **Dequeue** or **Both** (since we have TokND  $G$ ), which gives us the following:

$$\text{All}(xs_v, \Phi) \quad (5.24)$$

$$xs = xs_{\text{queue}} ++ [x''_{\text{head}}] ++ xs_{\text{old}} \quad (5.25)$$

$$\text{isLL}(xs) \quad (5.26)$$

$$\text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) \quad (5.27)$$

$$\ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x''_{\text{head}}) \quad (5.28)$$

for some  $xs$ ,  $xs_{\text{queue}}$ ,  $x''_{\text{head}}$ ,  $xs_{\text{old}}$ , and  $xs_v$ . We combine 5.28 with our half of the points-to predicate to conclude  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head}})$  and use lemma 33 to conclude  $x''_{\text{head}} = x_{\text{head}}$ .

We can now perform the store, swinging the head pointer to  $x_{\text{head\_next}}$ , which gives us  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head\_next}})$ . This is a linearisation point, so we must update the abstract state  $xs_v$ , which we got from the invariant opening.

From 5.22 and 5.26, we can deduce that  $xs_{\text{queue}}$  isn't empty (as otherwise, we would have  $x_{\text{head}} \mapsto \text{None}$  which contradicts with 5.22), and in fact, its first element must be  $x_{\text{head\_next}}$ . That is,

$$xs_{\text{queue}} = xs'_{\text{queue}} ++ [x_{\text{head\_next}}] \quad (5.29)$$

Combining this with 5.27 allows us to give a similar conclusion to the sequential proof: there exists  $xs'_v$  and  $v$  such that

$$xs_v = xs'_v ++ [v] \quad (5.30)$$

$$\text{projVal}(xs'_{\text{queue}}) = \text{wrapSome}(xs'_v) \quad (5.31)$$

$$\text{val}(x_{\text{head\_next}}) = \text{Some } v \quad (5.32)$$

By 5.24 and 5.30, we deduce  $\text{All}(xs'_v, \Phi)$  and  $\Phi(v)$ . We are now finally ready to close the invariant. We use  $xs'_v$  for the abstract state, giving up  $\text{All}(xs'_v, \Phi)$ , but allowing us to keep  $\Phi(v)$ . For the concrete state, we use the same  $xs$ . Note that by combining 5.25 and 5.29, we have that  $xs = xs'_{\text{queue}} ++ [x_{\text{head\_next}}] ++ (x_{\text{head}} :: xs_{\text{old}})$ , which allows us to pick  $xs'_{\text{queue}}$  as the queue,  $x_{\text{head\_next}}$  as the head node, and  $(x_{\text{head}} :: xs_{\text{old}})$  as the old nodes. Since  $xs$  hasn't changed, we prove the isLL predicate using 5.26, and the relationship between  $xs'_v$  and  $xs'_{\text{queue}}$  follows from 5.31. Finally, we set the state of the queue to **Static** or **Enqueue**, giving up TokND  $G$  and  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head\_next}})$ , and obtaining TokD  $G$ .

We step to line 24 and release the lock by giving up TokD  $G$  and  $\text{locked}(G.\gamma_{\text{Hlock}})$ . Lastly, we return the dequeued value:  $\text{val}(x_{\text{head\_next}})$ , meaning we must prove the post-condition with  $w = \text{val}(x_{\text{head\_next}})$ . Recall that we got to keep  $\Phi(v)$ , so using 5.32, we can prove the right disjunct.  $\square$

### 5.3.4 Discussing the need for Old Nodes

As mentioned in the observations, it is possible for the tail to lag one node behind the head. This insight lead to including the old nodes of the queue in the queue invariant:  $xs_{\text{old}}$ . This addition manifests in

the end of the proof of dequeue. When we open the invariant to swing  $\ell_{\text{head}}$  to  $x_{\text{head\_next}}$ , we get that the entire linked list is  $xs$ . After performing the store, we can then close the invariant with the same  $xs$  that we opened the queue to, just written differently to signify that  $x_{\text{head}}$  is now “old”, and  $x_{\text{head\_next}}$  is the new head node. Because of this, we can supply the same predicate concerning the location of the  $x_{\text{tail}}$  in the linked list that we got when we opened the invariant to prove either the **Static** state or the **Enqueue** state.

Had we not used an  $xs_{\text{old}}$  and essentially just “forgotten” old nodes, we couldn’t have done this. Say we defined  $xs$  as  $xs = xs_{\text{queue}} ++ [x_{\text{head}}]$  instead. Then, once we have to close the invariant, we cannot supply  $xs$ , which we got when we opened the invariant. Our only choice (due to the fact that  $loc_{\text{head}}$  must point to  $x_{\text{head\_next}}$ ) is to close the invariant with  $xs' = xs_{\text{queue}} = xs'_{\text{queue}} ++ [x_{\text{head\_next}}]$ . However, clearly  $xs' \neq xs$ , so we cannot supply the same predicate concerning the location of  $x_{\text{tail}}$  that we got when opening the invariant, since this predicate talks about  $xs$ , not  $xs'$ . Now, if we opened the invariant in the state **Dequeue**, then we could conclude  $\text{isLast}(x_{\text{tail}}, xs')$  from  $\text{isLast}(x_{\text{tail}}, xs)$ , due to the relationship between  $xs$  and  $xs'$ , and still be able close the invariant. However, if we opened the invariant in state **Both**, then we would need to assert  $\text{isSndLast}(x_{\text{tail}}, xs')$  from  $\text{isSndLast}(x_{\text{tail}}, xs)$ . This is however not provable, since  $\text{isSndLast}(x_{\text{tail}}, xs)$  allows for the case where  $xs'_{\text{queue}}$  is empty. In this case  $xs' = [x_{\text{head\_next}}]$  which makes it impossible to prove  $\text{isSndLast}(x_{\text{tail}}, xs')$ , as it is impossible to be the second last element in a list of size one.

## 5.4 Hocap-Style Specification

### 5.4.1 Introduction

When proving the concurrent specification in the previous section, we were quite careful with tracking the state of the queue, and to some extent, even its contents. The contents may have been existentially quantified, but through saving half a pointer, we could match up the contents of the queue between invariant openings. Given this precision in the proof, it should be no surprise that proving the Hocap-style specification will be very similar. This section will hence focus on the parts where the proof of the Hocap-style specification differs from that of the concurrent specification.

### 5.4.2 Hocap-Style Queue Predicate

Our definition of the queue predicate,  $\text{isQueue}$ , is *almost* the same as  $\text{isQueue}_C$ . First and foremost, the collection of ghost names,  $Q\text{ghostnames}$ , now contains the additional ghost name:  $\gamma_{\text{Abst}}$ , as required. For the queue invariant, the differences are that we no longer take the predicate  $\Psi$ , and we assert  $G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs_v$  instead of  $\text{All}(xs_v, \Psi)$ . That is, our queue invariant for the Hocap-style specification now looks as follows.

**Definition 5.4.1** (Two-Lock M&S Queue Hocap Invariant).

$$\begin{aligned}
& \text{ITLH}(\ell_{\text{head}}, \ell_{\text{tail}}, G) \triangleq \\
& \exists xs_v. G.\gamma_{\text{Abst}} \mapsto_{\bullet} xs_v * \quad (\text{abstract state}) \\
& \exists xs, xs_{\text{queue}}, xs_{\text{old}}, x_{\text{head}}, x_{\text{tail}}. \quad (\text{concrete state}) \\
& xs = xs_{\text{queue}} ++ [x_{\text{head}}] ++ xs_{\text{old}} * \\
& \text{isLL}(xs) * \\
& \text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) * \\
& ( \\
& \quad \ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, xs) * \quad (\text{Static}) \\
& \quad \text{TokNE } G * \text{TokND } G * \text{TokUpdated } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}}) * \quad (\text{Enqueue}) \\
& \quad (\text{isLast}(x_{\text{tail}}, xs) * \text{TokBefore } G \vee \text{isSndLast}(x_{\text{tail}}, xs) * \text{TokAfter } G) * \\
& \quad \text{TokE } G * \text{TokND } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \text{isLast}(x_{\text{tail}}, xs) * \quad (\text{Dequeue}) \\
& \quad \text{TokNE } G * \text{TokD } G * \text{TokUpdated } G \\
& \vee \\
& \quad \ell_{\text{head}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{head}}) * \ell_{\text{tail}} \mapsto^{\frac{1}{2}} \text{in}(x_{\text{tail}}) * \quad (\text{Both}) \\
& \quad (\text{isLast}(x_{\text{tail}}, xs) * \text{TokBefore } G \vee \text{isSndLast}(x_{\text{tail}}, xs) * \text{TokAfter } G) * \\
& \quad \text{TokE } G * \text{TokD } G \\
& )
\end{aligned}$$

For the queue predicate, the only real difference is that we no longer take the predicate  $\Psi$ . That is, our queue predicate is the following.

**Definition 5.4.2** (Two-Lock M&S Queue - isQueue Predicate (Hocap)).

$$\begin{aligned}
\text{isQueue}_C(v_q, G) & \triangleq \exists \ell_{\text{queue}}, \ell_{\text{head}}, \ell_{\text{tail}} \in \text{Loc}. \exists h_{\text{lock}}, t_{\text{lock}} \in \text{Val}. \\
& v_q = \ell_{\text{queue}} * \ell_{\text{queue}} \mapsto^{\square} ((\ell_{\text{head}}, \ell_{\text{tail}}), (h_{\text{lock}}, t_{\text{lock}})) * \\
& \boxed{\text{ITLH}(\Phi, \ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} * \\
& \text{isLock}(G.\gamma_{\text{Hlock}}, h_{\text{lock}}, \text{TokD } G) * \\
& \text{isLock}(G.\gamma_{\text{Tlock}}, t_{\text{lock}}, \text{TokE } G).
\end{aligned}$$

### 5.4.3 Proof Sketch

The Hocap-style proofs are largely similar their concurrent counterparts. However, instead of having to handle the  $\Psi$  predicate, we must now work with the authoritative and fragmental views of the abstract state. For initialise, we must additionally get ownership of the authoritative and fragmental view of the empty abstract state, and for enqueue and dequeue, the only real changes happen at the linearisation points. We sketch these challenges below.

#### Initialise

**Lemma 10** (Two-Lock M&S Queue Hocap Specification - Initialise).

$$\{\text{True}\} \text{ initialize } () \{v_q. \exists G. \text{isQueue}(v_q, G) * G.\gamma_{\text{Abst}} \mapsto_{\circ} []\}$$

As discussed, we must obtain  $G.\gamma_{\text{Abst}} \mapsto_{\bullet} [] * G.\gamma_{\text{Abst}} \mapsto_{\circ} []$ , for some ghost name  $G.\gamma_{\text{Abst}}$ . We achieve this by applying lemma 1 instantiated with the empty state:  $[]$ . We use  $G.\gamma_{\text{Abst}} \mapsto_{\bullet} []$  to establish the queue invariant, and  $G.\gamma_{\text{Abst}} \mapsto_{\circ} []$  to prove the post-condition.

## Enqueue

**Lemma 11** (Two-Lock M&S Queue Hocap Specification - Enqueue).

$$\forall v_q, v, G, P, Q. \quad (\forall x_{s_v}. G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v} * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i \uparrow} \triangleright G.\gamma_{\text{Abst}} \models_{\bullet} (v :: x_{s_v}) * Q) \multimap \\ \{\text{isQueue}(v_q, G) * P\} \text{ enqueue } v_q \ v \ \{w.Q\}$$

We start by assuming the viewshift which allows us to update  $P$  to  $Q$  and  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$  to  $G.\gamma_{\text{Abst}} \models_{\bullet} (v :: x_{s_v})$ , for any  $x_{s_v}$ . The only real change from the previous proof happens the second time we open the invariant; the first and third times, the abstract state doesn't change, hence we simply frame away the newly added authoritative fragment, and continue as we did before. The second time we open the invariant is on line 10, around the expression that adds the newly created node,  $x_{\text{new}}$ , to the linked list. When opening the invariant, we now get  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$ . And as before, we also get all the resources to match up variables and step through the code, updating the concrete state. To close the invariant, we must make the same choice of abstract state as we did previously:  $(v :: x_{s_v})$ . This, however, requires us to obtain  $G.\gamma_{\text{Abst}} \models_{\bullet} (v :: x_{s_v})$ . However, since we have  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$  and  $P$  (from the pre-condition), we can apply the viewshift to obtain it, along with  $Q$ . This then allows us to close the invariant, and the proof proceeds as previously. At the end, we must now additionally prove the post-condition  $Q$ , but this is no issue as we obtained that from the viewshift.

## Deque

**Lemma 12** (Two-Lock M&S Queue Hocap Specification - Dequeue).

$$\forall v_q, G, P, Q. \\ \left( \forall x_{s_v}. G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v} * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i \uparrow} \triangleright \left( \begin{array}{l} (x_{s_v} = [] * G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v} * Q(\text{None})) \\ \vee \left( \begin{array}{l} \exists v, x'_{s_v}. x_{s_v} = x'_{s_v} ++ [v] * \\ G.\gamma_{\text{Abst}} \models_{\bullet} x'_{s_v} * Q(\text{Some } v) \end{array} \right) \end{array} \right) \right) \multimap \\ \{\text{isQueue}(v_q, G) * P\} \text{ dequeue } v_q \ \{w.Q(w)\}$$

We assume the viewshift and proceed as in the concurrent proof until we get to the second time we open the invariant. This invariant opening is on line 17 around the expression that reads the head node's out pointer. It is here that we figure out whether or not the queue is empty by doing case analysis on  $x_{s_{\text{queue}}}$ . We open the invariant and do the same case analysis here.

**Case  $x_{s_{\text{queue}}}$  is empty:** In the case that the queue is empty, the abstract state of the queue will not change. We thus apply the viewshift (we have  $P$  from the precondition and  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$  from the invariant), which gives us the consequent of the viewshift. The right disjunct states that the abstract state,  $x_{s_v}$ , is non-empty, but since the abstract state is reflected in  $x_{s_{\text{queue}}}$ , which *is* empty, then we know that the right disjunct is impossible. Hence we may assume the left disjunct. I.e.  $x_{s_v} = [] * G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v} * Q(\text{None})$ . We now proceed as before, this time giving up  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$  to close the invariant. After stepping through the code, we are left with proving the postcondition:  $Q(\text{None})$ , which we got from the viewshift.

**Case  $x_{s_{\text{queue}}}$  is not empty:** If the queue is not empty, then we do not apply the viewshift (as the abstract state doesn't change within this invariant opening), and simply continue as we did previously. The next time we open the invariant is on line 23, around the expression that swings  $\ell_{\text{head}}$  to  $x_{\text{head\_next}}$ . It is this store operation that updates the abstract state of the queue, so it is within this invariant opening that we apply the viewshift (again, we have  $P$  from the pre-condition and  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$  from the invariant). As we did previously, we deduce that the current  $x_{s_{\text{queue}}}$  is non-empty, and since  $x_{s_v}$  is reflected in  $x_{s_{\text{queue}}}$ , then we can conclude that the first disjunct is impossible, so the viewshift gives us  $\exists v, x'_{s_v}. x_{s_v} = x'_{s_v} ++ [v] * G.\gamma_{\text{Abst}} \models_{\bullet} x'_{s_v} * Q(\text{Some } v)$ . As before, we conclude that *Some*  $v$  is the return value (through the reflection between  $x_{s_{\text{queue}}}$  and  $x_{s_v}$ ), and proceed to close the invariant, this time giving up  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{s_v}$ . Stepping through the code, we end up having to prove the post-condition  $Q(\text{Some } v)$ , which we got from the viewshift.

## Chapter 6

# The Lock-Free Michael-Scott Queue

### 6.1 Introduction

In this chapter we will study the non-blocking version of the M&S Queue, the Lock-Free M&S Queue. As with the two-lock version, the original implementation can be found in Michael and Scott [1996]. As the name “Lock-Free” suggests, the implementation doesn’t rely on locks to achieve correct behaviour. Instead, it uses the atomic operation: [CAS](#) which we discussed in section 2.1.

### 6.2 Implementation

The implementation shares many commonalities with the two-lock variant, most importantly, the underlying queue is still a linked list. The major differences are how we manipulate the linked list and the head and tail pointers.

#### 6.2.1 Initialise

As the implementation doesn’t use locks, the queue data structure is now just a pointer to a pair consisting of the head and tail pointers.

#### 6.2.2 Enqueue

Enqueueing a node consists of the same two steps as before: add the newly created node to the linked list, and swing the tail pointer. In this way, figures 4.2a, 4.2b, and 4.2c still accurately reflect the state changes that the queue goes through during an enqueue. However, one big difference is that swinging the tail pointer to the newly inserted node is not necessarily done by the thread that enqueued it. We discuss this below.

Since other threads can work on the linked list at the same time, we must do some extra checks when enqueueing a new node.

Firstly, we must ensure that the tail pointer is actually pointing to the last node. We ensure this on line 11. If it isn’t, that means that another enqueueing thread has added a new node to the linked list, but hasn’t yet swung the tail pointer. So instead of waiting for it, we *help* it by trying to swing the tail pointer for it. Afterwards, we try to enqueue our node again.

Secondly, before we can add the node, we must ensure that no thread has performed an enqueue while we have been working, so that we don’t overwrite another threads enqueue. This is ensured with the [CAS](#) instruction on line 12 – it adds the new node to the linked list only if our tail node is still the last, and hence points to None. After adding the node, we attempt to swing the tail pointer to it. This may fail, but that simply means that another thread has already swung it for us. Finally, we have an additional *consistency check* on line 10. We discuss the purpose of this extra check in section 7.6.

### 6.2.3 Dequeue

Fundamentally, a dequeue operation still consists of swinging the head pointer to the next node in the linked list. The original authors decided that the tail pointer shouldn't lag behind the head pointer, hence dequeue also accesses the tail node and ensures it won't lag behind as result of the dequeue. On line 26, we check whether the head and tail nodes are the same. If this is the case, then the queue is either empty or the tail pointer is lagging behind. We distinguish between these two cases on line 27. If the tail node is indeed lagging behind, then some thread has enqueued a node, but not yet swung the tail pointer, so we help it out by swinging the pointer.

If the head and tail nodes are not the same, then it must be safe to dequeue, which we attempt in the else block on line 31.

### 6.2.4 Prophecies

Notice that the load on line 24 is a possible linearisation point – if the load resolves to None (i.e. the queue is empty) and the consistency check on the next line passes, the dequeue will conclude the queue is empty and return None. That is, the point where we “read” the state of the queue is the load on line 24, but we only conclude that the queue is empty if the consistency check on the next line succeeds. So when we are at the load, we need to know whether or not the consistency check passes; if it does, we should apply the viewshift, and if it doesn't, we shouldn't apply it. This is somewhat of a conundrum, as this requires us to reason about a future computation.

To solve this issue, we use a *prophecy variable*. Prophecies are a part of Iris and allow us to reason about the result of future expressions. They are only used in the logic and do not alter the semantics of the code. On line 23 in dequeue, we create one such prophecy variable, which is then “resolved” on line 25. The prophecy variable allows us to reason about the result of the expression inside the [Resolve](#) statement. In other words, we can reason about consistency check already at the load on line 24, allowing us to make the correct choice. We will see exactly how this works when we prove the specification for the Lock-Free M&S Queue in chapter 7.

```

1  initialize  $\triangleq$ 
2    let node = ref (None, ref (None)) in
3    ref (ref (node), ref (node))
4
5  enqueue Q value  $\triangleq$ 
6    let node = ref (Some value, ref (None)) in
7    (rec loop_ =
8      let tail = !(snd(!Q)) in
9      let next = !(snd(!tail)) in
10     if tail = !(snd(!Q)) then
11       if next = None then
12         if CAS (snd(!tail)) next node then
13           CAS (snd(!Q)) tail node
14         else loop ()
15       else CAS (snd(!Q)) tail next; loop ()
16     else loop ()
17   ) ()
18
19  dequeue Q  $\triangleq$ 
20    (rec loop_ =
21      let head = !(fst(!Q)) in
22      let tail = !(snd(!Q)) in
23      let p = newproph in
24      let next = !(snd(!head)) in
25      if head = Resolve(!(fst(!Q)), p, ()) then
26        if head = tail then
27          if next = None then
28            None
29          else
30            CAS(snd(!Q)) tail next; loop ()
31        else
32          let value = fst(!next) in
33          if CAS (fst(!Q)) head next then
34            value
35          else loop ()
36      else loop ()
37    )()

```

Figure 6.1: Implementation of Lock-Free M&S Queue in HeapLang.

## Chapter 7

# Proving Hocap-style Specification for the Lock-Free Michael-Scott Queue

### 7.1 Introduction

In this chapter we prove that the Lock-Free M&S Queue satisfies the Hocap-style specification given in section 3.4. Section 7.2 introduces the notion of *Reachability* which we will need in the proofs. In section 7.3 we discuss the specification we will be proving, and in sections 7.4 and 7.5 we give the proofs. Finally, in section 7.6, we discuss a simplification to the Lock-Free M&S Queue algorithm and its implications.

### 7.2 Reachability

An important aspect in the correctness of the Lock-Free M&S Queue is which nodes a particular node is able to *reach* through the linked list (i.e. by following the chain of pointers), and how the head and tail pointers change during the lifetime of the queue.

Firstly, the underlying linked list still only ever grows, and it does so only at the end. Hence, the set of nodes that a given node can reach only ever grows. Further, all nodes can always reach the last node in the linked list.

Secondly, similarly to the two-lock variant, the correctness of the queue relies on the fact that the head and tail pointers are only ever swung towards the end of the linked list. That is, if a node can reach, say, the tail node at one point during the program, then it can reach any future tail nodes.

Thirdly, whereas it was possible for the tail node to lag behind the head node in the two-lock version, it is not possible in the lock-free version. Indeed, if such a scenario could happen, dequeue could crash! Consider the scenario where the head node is the last node in the linked list (hence the queue is empty), and the tail is lagging behind the head. If someone invokes dequeue, the check on line 26, which is supposed to detect an empty queue or a lagging tail will result to false, and hence, incorrectly, take the “else” branch, which assumes that there is something to dequeue. But since the queue is empty, then *next* – the node after head – is None, and when we try to dereference *next* on line 32, we will crash. Therefore, our invariant must ensure that the tail never lags behind the head.

To capture these properties, we introduce two notions of reachability: concrete reachability and abstract reachability, which we introduce in the following sections. This way of modelling the queue was originally introduced in Vindum and Birkedal [2021]. The presentation here borrows the same ideas, but differs in the sense that it is node-oriented instead of location-oriented. Moreover, we prove some additional properties of reachability which allows us to simplify the queue invariant slightly.

#### 7.2.1 Concrete Reachability

We say that a node  $x_n$  in a linked list can concretely reach a node  $x_m$  when, if we start traversing succeeding nodes (by following the out and in pointers starting from  $x_n$ ), we will eventually get to  $x_m$ . If this is the case, we write  $x_n \rightsquigarrow x_m$ . We allow for traversing zero nodes to reach  $x_m$ , which essentially



means that all nodes can reach themselves. Formally, we define concrete reachability as an inductive predicate.

**Definition 7.2.1** (Concrete Reachability). Given two nodes,  $x_n$  and  $x_m$ , we say that  $x_n$  concretely reaches  $x_m$ , if they satisfy the following inductive predicate.

$$x_n \rightsquigarrow x_m \triangleq \text{in}(x_n) \mapsto^\square (\text{val}(x_n), \text{out}(x_n)) * (x_n = x_m \vee \exists x_p. \text{out}(x_n) \mapsto^\square \text{in}(x_p) * x_p \rightsquigarrow x_m)$$

This definition firstly states that  $x_n$  is a node:  $\text{in}(x_n) \mapsto^\square (\text{val}(x_n), \text{out}(x_n))$ . Secondly,  $x_n$  is either the node to be reached,  $x_m$ , or it has a succeeding node  $x_p$ , which can reach  $x_m$ . Note that the points-to propositions are all persistent, which mimics the fact that the linked list is only ever changed by appending new nodes to the end. This in turn makes concrete reachability a persistent predicate.

We proceed to prove some useful lemmas about concrete reachability.

**Lemma 13** (reach-reflexive).  $x_n \rightsquigarrow x_n ** \text{in}(x_n) \mapsto^\square (\text{val}(x_n), \text{out}(x_n))$

*Proof.* The  $*$  direction follows directly by the definition. To prove the  $**$  direction, it suffices to show  $(x_n = x_n \vee \exists x_p. \text{out}(x_n) \mapsto^\square \text{in}(x_p) * x_p \rightsquigarrow x_n)$ . Clearly, this follows as the left disjunction holds.  $\square$

**Lemma 14** (reach-transitive).  $x_n \rightsquigarrow x_m * x_m \rightsquigarrow x_o * x_n \rightsquigarrow x_o$

*Proof.* We proceed by induction in  $x_n \rightsquigarrow x_m$ .

B.C. In the base case,  $x_n = x_m$ . We get to assume that  $x_m \rightsquigarrow x_o$ , and must prove  $x_n \rightsquigarrow x_o$ . Since  $x_n = x_m$ , we are done.

I.C. In the inductive case, we assume that  $x_n$  is a node that points to some  $x_p$ , which satisfies  $x_m \rightsquigarrow x_o * x_p \rightsquigarrow x_o$ . Assuming  $x_m \rightsquigarrow x_o$ , we must prove  $x_n \rightsquigarrow x_o$ .

To prove  $x_n \rightsquigarrow x_o$  we must first show that  $x_n$  is a node, which we have already established. Next, we must show that either  $x_n = x_o$ , or  $x_n$  steps to some  $x'_p$  which can reach  $x_o$ . We prove the second case by choosing our  $x_p$  for  $x'_p$ . Thus, we have to show  $x_p \rightsquigarrow x_o$ . This then follow by the induction hypothesis together with our assumption that  $x_m \rightsquigarrow x_o$ .  $\square$

**Lemma 15** (reach-from-is-node).  $x_n \rightsquigarrow x_m * \text{in}(x_n) \mapsto^\square (\text{val}(x_n), \text{out}(x_n))$

*Proof.* This follow immediately from the definition of concrete reachability.  $\square$

**Lemma 16** (reach-to-is-node).  $x_n \rightsquigarrow x_m * \text{in}(x_m) \mapsto^\square (\text{val}(x_m), \text{out}(x_m))$

*Proof.* We proceed by induction in  $x_n \rightsquigarrow x_m$ . The base case follows by lemma 15 above. In the inductive case, we assume that  $x_n$  points to some  $x_p$ , which reaches  $x_m$ . Our induction hypothesis is  $\text{in}(x_m) \mapsto^\square (\text{val}(x_m), \text{out}(x_m))$ , which is also our proof obligation, so we are done.  $\square$

**Lemma 17** (reach-last).  $x_n \rightsquigarrow x_m * \text{out}(x_n) \mapsto \text{None} * x_n = x_m * \text{out}(x_n) \mapsto \text{None}$

*Proof.* Assuming  $x_n \rightsquigarrow x_m$  and  $\text{out}(x_n) \mapsto \text{None}$ , we must prove that  $x_n = x_m$  and  $\text{out}(x_n) \mapsto \text{None}$ . By  $x_n \rightsquigarrow x_m$ , we know that either  $x_n = x_m$ , or  $x_n$  points to some  $x_p$  and  $x_p \rightsquigarrow x_m$ . The first case immediately gives us everything we need to prove both goals. If we are in the second case, then we know that  $\text{out}(x_n) \mapsto^\square \text{in}(x_p)$ . But by our initial assumption,  $\text{out}(x_n) \mapsto \text{None}$ . Since  $\text{in}(x_p)$  is a location, then this is clearly a contradiction.  $\square$

## 7.2.2 Abstract Reachability

As discussed, we wish to capture that if a node can reach either the head or tail node at one point during the program, then it can reach any future head or tail nodes. To do this we introduce the notion of abstract reachability. The idea is to introduce ghost variables that can “point” to nodes in the linked list, just as the head and tail pointers do. We shall write  $\gamma \rightsquigarrow x$  to mean that the ghost variable  $\gamma$  *abstractly points* to the node  $x$ . We shall construct the abstract points-to predicate so that we can update  $\gamma \rightsquigarrow x$  to  $\gamma \rightsquigarrow y$  only if  $x$  can concretely reach  $y$ , i.e.  $x \rightsquigarrow y$ . This additional restriction compared to the normal

points-to predicate is what allows us to capture the property described above. We write  $x \dashrightarrow \gamma$  to mean that the node  $x$  can *abstractly reach* the ghost variable  $\gamma$ . The idea is that, if we have established  $x \dashrightarrow \gamma$ , then no matter what node  $\gamma$  abstractly points to, for instance  $\gamma \mapsto y$ , we can conclude  $x \leadsto y$ . This means that if we update  $\gamma \mapsto y$  in the future to, say  $\gamma \mapsto z$ , then we can conclude that  $x \leadsto z$ .

We refer to section 8.3 for the definition of the resource algebra we use to formally define abstract points-to and abstract reachability. For our purposes, we only need to know the following four lemmas to work with the predicates in the proofs. These lemmas are also proved in section 8.3.

Firstly, if we have a node, we may allocate some ghost variable  $\gamma$  which points to it and assert that the node can reach  $\gamma$ .

**Lemma 18** (Abs-Reach-Alloc).  $x \leadsto x \Rightarrow (\exists \gamma. \gamma \mapsto x * x \dashrightarrow \gamma)$

The second lemma allows us to get a *concrete* reachability predicate out of an abstract one. If a ghost name  $\gamma_m$  currently points abstractly to some node  $x_m$ , then any node that can abstractly reach  $\gamma$ , can also *concretely* reach  $x_m$ .

**Lemma 19** (Abs-Reach-Concr).  $x_n \dashrightarrow \gamma_m * \gamma_m \mapsto x_m \multimap x_n \leadsto x_m * \gamma_m \mapsto x_m$

We can also go the other way, and get an abstract reachability predicate out of a concrete one. If a ghost variable  $\gamma_m$  points abstractly to some node  $x_m$ , and a node  $x_n$  can *concretely* reach  $x_m$ , then we may deduce that  $x_n$  can *abstractly* reach  $\gamma_m$ , meaning that  $x_n$  can reach any node that  $\gamma_m$  will ever point to in the future.

**Lemma 20** (Abs-Reach-Abs).  $x_n \leadsto x_m * \gamma_m \mapsto x_m \Rightarrow x_n \dashrightarrow \gamma_m * \gamma_m \mapsto x_m$

The final lemma allows us update abstract pointers. As discussed above, we will require that whatever node we update the pointer to is a successor of the current node. That is, if a ghost variable  $\gamma_m$  currently points to  $x_m$ , then we must show that  $x_m$  can reach  $x_o$ , before we can update  $\gamma$  to point abstractly to  $x_o$ . After the update we additionally get that  $x_o$  can abstractly reach  $\gamma$ .

**Lemma 21** (Abs-Reach-Advance).  $\gamma_m \mapsto x_m * x_m \leadsto x_o \Rightarrow \gamma_m \mapsto x_o * x_o \dashrightarrow \gamma_m$

## 7.3 Specifications for Lock-Free M&S Queue

As we showed in section 3.5, the sequential and concurrent specifications can be derived from the Hocap-style specification *without* referring to the actual implementation. Thus, in this chapter we focus only on proving the Hocap-style specification, and rely on the derivations to prove the concurrent and sequential specifications.

We prove the Hocap-style specification for the initialize, enqueue, and dequeue functions introduced in chapter 6. We make the collection of ghost names,  $Q_{\text{ghostnames}}$ , be the set of four-tuples. Firstly, it contains the mandatory  $\gamma_{\text{Abst}}$  whose purpose is the same as before; to keep track of the abstract state of the queue. Additionally, it contains  $\gamma_{\text{Head}}$ ,  $\gamma_{\text{Tail}}$ , and  $\gamma_{\text{Last}}$ , which will abstractly point to the head, tail, and last node, respectively.

## 7.4 Hocap-style Queue Predicate

Since the queue predicate is required to be persistent and multiple threads need to access shared resources of the queue, we will be needing an invariant. The invariant we define has some commonalities with the invariant we used for the two-lock variant, but it incorporates the differences we discussed earlier in the chapter. In particular, it is important for the correctness of the queue that the tail does not lag behind the head. As such, our invariant will not allow for this behaviour. This has the extra implication that the head node is always the oldest node, meaning that we do not need to keep track of older nodes,  $x_{\text{old}}$ .

Unlike the two-lock variant, we assert the existence of an additional node  $x_{\text{last}}$ , which invariantly is the last (newest added) node in the linked list. This helps us reason about where the head and tail nodes are located; enqueue distinguishes between the cases where the tail is last and not last, and similarly for dequeue and head.

In this way,  $x_{\text{head}}$  is the first node,  $x_{\text{last}}$  is the last node, and  $x_{\text{tail}}$  either lies somewhere in between, is

one of them, or, in the case where the queue is empty, is both of them. To force this structure, we use our abstract reachability predicate from the previous section. We proceed to define the invariant.

**Definition 7.4.1** (Lock-Free M&S Queue Invariant).

$$\begin{aligned}
& I_{\text{LFH}}(\ell_{\text{head}}, \ell_{\text{tail}}, G) \triangleq \\
& \exists xs_v. G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs_v * & (\text{abstract state}) \\
& \exists xs, xs_{\text{queue}}, x_{\text{head}}, x_{\text{tail}}, x_{\text{last}}. & (\text{concrete state}) \\
& xs = xs_{\text{queue}} ++ [x_{\text{head}}] * \\
& \text{isLL}(xs) * \\
& \text{isLast}(x_{\text{last}}, xs) \\
& \text{projVal}(xs_{\text{queue}}) = \text{wrapSome}(xs_v) * \\
& \ell_{\text{head}} \mapsto \text{in}(x_{\text{head}}) * \\
& \ell_{\text{tail}} \mapsto \text{in}(x_{\text{tail}}) * \\
& G.\gamma_{\text{Head}} \mapsto x_{\text{head}} * x_{\text{head}} \dashrightarrow G.\gamma_{\text{Tail}} * \\
& G.\gamma_{\text{Tail}} \mapsto x_{\text{tail}} * x_{\text{tail}} \dashrightarrow G.\gamma_{\text{Last}} * \\
& G.\gamma_{\text{Last}} \mapsto x_{\text{last}}
\end{aligned}$$

The isQueue predicate is now quite simple: it states that the value representing the queue is a location which points persistently to a pair of locations, the head and tail pointers, which satisfy the invariant we defined above.

**Definition 7.4.2** (Lock-Free M&S Queue - isQueue Predicate).

$$\begin{aligned}
& \text{isQueue}(v_q, G) \triangleq \exists \ell_{\text{queue}}, \ell_{\text{head}}, \ell_{\text{tail}} \in \text{Loc}. \\
& v_q = \ell_{\text{queue}} * \ell_{\text{queue}} \mapsto^{\square} (\ell_{\text{head}}, \ell_{\text{tail}}) * \\
& \boxed{I_{\text{LFH}}(\ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.\text{queue}}
\end{aligned}$$

## 7.5 Proof Outline

We instantiate the specification with our definition of isQueue (7.4.2). By the definition of isQueue we easily show that isQueue is persistent. What remains to be shown is the specifications for initialize, enqueue, and dequeue. Both enqueue and dequeue has code that attempt to swing the tail pointer forward (for enqueue, lines 13 and 15, and for dequeue, line 30). These all behave similarly, so we additionally prove a specification for swinging the tail.

### Initialise

**Lemma 22** (Lock-Free M&S Queue Specification - Initialise).

$$\{\text{True}\} \text{ initialize } () \{v_q, \exists G. \text{isQueue}(v_q, G) * G.\gamma_{\text{Abst}} \Rightarrow_{\circ} []\}$$

*Proof.* We first step through line 2 which creates a new node:  $x_1 = (\ell_{1\_in}, \text{None}, \ell_{1\_out})$ , with  $\ell_{1\_out} \mapsto \text{None}$  and  $\ell_{1\_in} \mapsto (\text{None}, \ell_{1\_out})$ , the latter of which we make persistent. Next, we step through line 3 which gives us some locations  $\ell_{\text{head}}$ ,  $\ell_{\text{tail}}$ , and  $\ell_{\text{queue}}$ , with  $\ell_{\text{head}} \mapsto \text{in}(x_1)$  and  $\ell_{\text{tail}} \mapsto \text{in}(x_1)$ , and finally  $\ell_{\text{queue}} \mapsto (\ell_{\text{head}}, \ell_{\text{tail}})$ .

As we did for the two-lock version, we apply lemma 1 to allocate an empty abstract queue, giving us some ghost name  $\gamma_{\text{Abst}}$  that we put into  $G$ , and the resources  $G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} [] * G.\gamma_{\text{Abst}} \Rightarrow_{\circ} []$ . To allocate the invariant, we must additionally obtain abstract reach propositions. Since  $x_1$  is a node, we may use lemma 13 to conclude  $x_1 \rightsquigarrow x_1$ . We can now use lemma 18 three times, giving us ghost names  $\gamma_{\text{Head}}$ ,  $\gamma_{\text{Tail}}$ ,  $\gamma_{\text{Last}}$  which we again put into  $G$ , and the resources

$$G.\gamma_{\text{Head}} \mapsto x_1 * G.\gamma_{\text{Tail}} \mapsto x_1 * x_1 \dashrightarrow G.\gamma_{\text{Tail}} * G.\gamma_{\text{Last}} \mapsto x_1 * x_1 \dashrightarrow G.\gamma_{\text{Last}}$$

We now have all the resources we need to allocate the invariant with the head, tail, and last node being  $x_1$ .

With the invariant allocated, proving the post-condition becomes straightforward.  $\square$

### Swing Tail

The specification we wish to prove is the following.

**Lemma 23** (Swing Tail).  $\forall \ell_{\text{head}}, \ell_{\text{tail}}, x_{\text{tail}}, x_{\text{newtail}}, G.$   
 $\left\{ \boxed{\text{LFH}(\ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} * x_{\text{tail}} \rightsquigarrow x_{\text{newtail}} * x_{\text{newtail}} \dashrightarrow G.\gamma_{\text{Last}} \right\}$   
**CAS**  $\ell_{\text{tail}} \text{ in}(x_{\text{tail}}) \text{ in}(x_{\text{newtail}})$   
 $\{w.w = \text{true} \vee w = \text{false}\}$

*Proof.* The rule for **CAS** demands that we have a points-to predicate for  $\ell_{\text{tail}}$ . This is available inside the invariant, so we proceed to open it. This tells us that there is some  $x'_{\text{tail}}$  so that  $\ell_{\text{tail}} \mapsto \text{in}(x'_{\text{tail}})$ . We consider both cases of the **CAS**:

**Case CAS** succeeds. It must then have been the case that  $\text{in}(x'_{\text{tail}}) = \text{in}(x_{\text{tail}})$ . Since we have  $x_{\text{tail}} \rightsquigarrow x_{\text{newtail}}$ , we know that  $x_{\text{tail}}$  is a node (lemma 15). From the invariant, we additionally got that  $G.\gamma_{\text{Tail}} \rightsquigarrow x'_{\text{tail}}$  and  $x_{\text{head}} \dashrightarrow G.\gamma_{\text{Tail}}$ , which by lemma 19 means that  $x_{\text{head}} \rightsquigarrow x'_{\text{tail}}$ . We can thus also conclude that  $x'_{\text{tail}}$  is a node (lemma 16). So since both  $x_{\text{tail}}$  and  $x'_{\text{tail}}$  are nodes, and  $\text{in}(x'_{\text{tail}}) = \text{in}(x_{\text{tail}})$ , lemma 33 tells us that  $x_{\text{tail}} = x'_{\text{tail}}$ . In other words, we have that  $G.\gamma_{\text{Tail}} \rightsquigarrow x_{\text{tail}}$ . Since the **CAS** succeeded, we now have that  $\ell_{\text{tail}} \mapsto \text{in}(x_{\text{newtail}})$ . Since the invariant demands that  $G.\gamma_{\text{Tail}}$  and  $\ell_{\text{tail}}$  agree on the node they point to, we must update  $G.\gamma_{\text{Tail}} \rightsquigarrow x_{\text{tail}}$  to  $G.\gamma_{\text{Tail}} \rightsquigarrow x_{\text{newtail}}$ . We can do this using lemma 21 as we assumed  $x_{\text{tail}} \rightsquigarrow x_{\text{newtail}}$ . With this, we can close the invariant again, using  $x_{\text{newtail}}$  as the tail node.

The **CAS** evaluates to **true** which we use to prove the first disjunct of the post-condition.

**Case CAS** Fails. Since the **CAS** failed, nothing was updated, and we can close the invariant again with the same resources we got out of it. The **CAS** evaluates to **false**, hence we can prove the second disjunct in the post-condition.  $\square$

### Enqueue

**Lemma 24** (Lock-Free M&S Queue Specification - Enqueue).

$$\forall v_q, v, G, P, Q. \left( \forall x s_v. G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} x s_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i^{\uparrow}} \triangleright G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} (v :: x s_v) * Q \right) \multimap \{ \text{isQueue}(v_q, G) * P \} \text{enqueue } v_q \ v \ \{w.Q\}$$

*Proof.* We assume the viewshift and proceed to prove the hoare triple. By definition of `isQueue`, we know that  $v_q$  is a location  $\ell_{\text{queue}}$  and there are locations  $\ell_{\text{head}}$  and  $\ell_{\text{tail}}$  so that

$$\ell_{\text{queue}} \mapsto^{\square} (\ell_{\text{head}}, \ell_{\text{tail}}) * \boxed{\text{LFH}(\ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} \quad (7.1)$$

We first step through line 6 which creates a new node  $x_{\text{new}}$ , so that

$$\text{in}(x_{\text{new}}) \mapsto^{\square} (\text{Some } v, \text{out}(x_{\text{new}})) \quad (7.2)$$

$$\text{out}(x_{\text{new}}) \mapsto \text{None} \quad (7.3)$$

The next line is the beginning of the looping function. We proceed by *l*öb induction, which allows us to assume the hoare triple we wish to prove *later*. This means that, if we reach a recursive call, we will have the hoare triple that we must prove – the later will be immediately stripped when we apply  $()$  and “step into” the recursive function.

Line 8 first dereferences to the tail pointer  $\ell_{\text{tail}}$  using the resources in 7.1. We open the invariant to obtain

the points-to predicate concerning  $\ell_{\text{tail}}$ . We get that  $\ell_{\text{tail}}$  points to some  $x_{\text{tail}}$ . Using the resources from the invariant, we may conclude the following persistent information:

$$x_{\text{tail}} \dashrightarrow G.\gamma_{\text{Last}} \quad (7.4)$$

$$\text{in}(x_{\text{tail}}) \mapsto^{\square} (\text{val}(x_{\text{tail}}), \text{out}(x_{\text{tail}})) \quad (7.5)$$

The first part is directly from the invariant, and the second we may derive using lemmas 19 and 16. We perform the load, and close the invariant.

The next line (line 9) finds out what  $x_{\text{tail}}$  points to. Using 7.5 we step to  $!(\text{out}(x_{\text{tail}}))$ . The points-to predicate required to perform this dereference is owned by the invariant (as it might be non-persistent), so we open the invariant again. We get that there is some  $x_{\text{last}}$ , with  $G.\gamma_{\text{Last}} \mapsto x_{\text{last}}$ . From this, 7.4, and lemma 19 we conclude  $x_{\text{tail}} \leadsto x_{\text{last}}$ . This gives us two cases to consider: either  $x_{\text{tail}}$  is  $x_{\text{last}}$  (meaning that  $x_{\text{tail}}$  is not lagging behind), or it points to some node  $x_{\text{tail\_next}}$  which can reach  $x_{\text{last}}$  (meaning that  $x_{\text{tail}}$  is lagging behind).

**Case**  $x_{\text{tail}} = x_{\text{last}}$ . Since we had  $\text{isLast}(x_{\text{last}}, xs)$ , we know that  $x_{\text{tail}}$  is the last node in the linked list, hence it points to None. We perform the load which sets  $\text{next}$  to None, and close the invariant.

We proceed to the consistency check on line 10. As before, the points-to predicate for  $\ell_{\text{tail}}$  is in the invariant, so we open it. We get  $\ell_{\text{tail}} \mapsto \text{in}(x'_{\text{tail}})$ , for some  $x'_{\text{tail}}$ . Using this, we perform the dereference and close the invariant. The branch taken now depends on whether or not  $x'_{\text{tail}}$  is consistent with  $x_{\text{tail}}$ . In case they aren't, we take the "else" branch on line 16, which simply consists of a recursive call to the looping function. We are done by the induction hypothesis (from the l b induction).

If they are consistent, we take the "then" branch and step to line 11. Here we check whether or not  $\text{next}$  is None. We already know this is the case, so we proceed to line 12. This consists of a CAS instruction which attempts to add  $x_{\text{new}}$  to the linked list. The CAS will succeed if and only if  $\ell_{\text{tail}}$  still points to None. We open the invariant to gain access to the relevant points-to predicate. Similarly to what we did earlier, we apply lemma 19 to conclude  $x_{\text{tail}} \leadsto x'_{\text{last}}$ , where  $x'_{\text{last}}$  is the current last node of the linked list (according to the invariant). As before, we perform case analysis on  $x_{\text{tail}} \leadsto x'_{\text{last}}$ .

**Case**  $x_{\text{tail}} = x'_{\text{last}}$ . We now know that  $x_{\text{tail}}$  is still the last node in the linked list, hence  $\text{out}(x_{\text{tail}}) \mapsto \text{None}$ , and the CAS will succeed. This instruction makes  $x_{\text{tail}}$  point to  $x_{\text{new}}$ , which essentially adds it to the linked list. Thus, the value in  $x_{\text{new}}$  becomes enqueued. In other words, this is a linearisation point, so we must apply the viewshift. We instantiate the viewshift with the abstract state of the queue  $xs_v$  from the invariant opening, and supply  $G.\gamma_{\text{Abst}} \mapsto_{\bullet} xs_v$  from the invariant and the  $P$  from the pre-condition. We hence get  $G.\gamma_{\text{Abst}} \mapsto_{\bullet} (v :: xs_v)$  and  $Q$ . When closing the invariant, we use  $(v :: xs_v)$  for the abstract state,  $(x_{\text{new}} :: xs)$  for the concrete state,  $x_{\text{new}} :: xs_{\text{queue}}$  for the queue, and we take  $x_{\text{new}}$  to be the last node. The head and tail nodes remain the same. This means we give up  $G.\gamma_{\text{Abst}} \mapsto_{\bullet} (v :: xs_v)$  that we got from the viewshift, and the points-to predicate 7.3 (used to assert  $\text{isLL}(x_{\text{new}} :: xs)$ ). The only thing left to prove is  $G.\gamma_{\text{Last}} \mapsto x_{\text{new}}$ . From the invariant opening, we have  $G.\gamma_{\text{Last}} \mapsto x_{\text{tail}}$ . Since  $x_{\text{tail}} \leadsto x_{\text{new}}$ , we may apply lemma 21 to update the abstract points-to resource to  $G.\gamma_{\text{Last}} \mapsto x_{\text{new}}$  and additionally obtain  $x_{\text{new}} \dashrightarrow G.\gamma_{\text{Last}}$ . With this, we can close the invariant, and step to line 13. This line attempts to swing the tail, so we apply our swing-tail lemma (lemma 23) by supplying our invariant,  $x_{\text{tail}} \leadsto x_{\text{new}}$ , and  $x_{\text{new}} \dashrightarrow G.\gamma_{\text{Last}}$ . This tells us that the CAS is safe, and it either succeeds or fails. The resulting value is the returned value of the enqueue function, but since the post-condition is simply  $Q$ , which we own, we are done.

**Case**  $\text{out}(x_{\text{tail}}) \mapsto^{\square} \text{in}(x_{\text{tail\_next}}) * x_{\text{tail\_next}} \leadsto x_{\text{last}}$ . Since  $x_{\text{tail}}$  doesn't point to None, the CAS will fail. We close the invariant and step to line 14. We finish by applying the induction hypothesis.

**Case**  $\text{out}(x_{\text{tail}}) \mapsto^{\square} \text{in}(x_{\text{tail\_next}}) * x_{\text{tail\_next}} \leadsto x_{\text{last}}$ . Using this we perform the load, which sets  $\text{next}$  to  $\text{in}(x_{\text{tail\_next}})$ . Before closing the invariant, we apply lemma 20 with  $x_{\text{tail\_next}} \leadsto x_{\text{last}}$  and  $G.\gamma_{\text{Last}} \mapsto x_{\text{last}}$  to obtain  $x_{\text{tail\_next}} \dashrightarrow G.\gamma_{\text{Last}}$ . We now proceed to close the invariant. Next, we reach the consistency check. We handle it similarly to the previous case: open the invariant, get some  $x'_{\text{tail}}$ , close the invariant, and in case of inconsistency, apply induction hypothesis. If the nodes are consistent, we step to line 11. This time, the check will fail as  $\text{next}$  is  $\text{in}(x_{\text{tail\_next}})$  which is not None. Hence we step to line 15 which attempts to swing the tail pointer. We here apply lemma

23 which we can do as we own the invariant,  $x_{\text{tail}} \rightsquigarrow x_{\text{tail\_next}}$ , and  $x_{\text{tail\_next}} \dashrightarrow G.\gamma_{\text{Last}}$ . We step through to the recursive call and finish by applying the induction hypothesis.  $\square$

## Deque

**Lemma 25** (Lock-Free M&S Queue Specification - Dequeue).

$$\forall v_q, G, P, Q. \left( \begin{array}{l} \left( \forall x s_v. G.\gamma_{\text{Abst}} \Rightarrow \bullet x s_v * P \Rightarrow_{\mathcal{E} \setminus \mathcal{N}.i^\dagger} \triangleright \left( \begin{array}{l} (x s_v = [] * G.\gamma_{\text{Abst}} \Rightarrow \bullet x s_v * Q(\text{None})) \\ \vee \left( \begin{array}{l} \exists v, x s'_v. x s_v = x s'_v ++ [v] * \\ G.\gamma_{\text{Abst}} \Rightarrow \bullet x s'_v * Q(\text{Some } v) \end{array} \right) \end{array} \right) \right) * \\ \{\text{isQueue}(v_q, G) * P\} \text{ dequeue } v_q \{w.Q(w)\} \end{array} \right)$$

*Proof.* We assume the viewshift and must prove the hoare triple. As before, we know from `isQueue` that the queue,  $v_q$ , is a location,  $\ell_{\text{queue}}$ , and there are locations  $\ell_{\text{head}}$  and  $\ell_{\text{tail}}$  so that

$$\ell_{\text{queue}} \mapsto^\square (\ell_{\text{head}}, \ell_{\text{tail}}) * \boxed{\text{LFH}(\ell_{\text{head}}, \ell_{\text{tail}}, G)}^{\mathcal{N}.queue} \quad (7.6)$$

The body of `dequeue` is the loop, so we immediately apply `l b` induction. We step through the function application, and into the looping function to line 21. This line dereferences  $\ell_{\text{head}}$ , so we open the invariant to access the associated points-to predicate. We obtain that  $\ell_{\text{head}}$  points to some  $x_{\text{head}}$ , meaning the load results to `in`( $x_{\text{head}}$ ). We also derive the following information

$$\text{in}(x_{\text{head}}) \mapsto^\square (\text{val}(x_{\text{head}}), \text{out}(x_{\text{head}})) \quad (7.7)$$

$$x_{\text{head}} \dashrightarrow G.\gamma_{\text{Head}} \quad (7.8)$$

$$x_{\text{head}} \dashrightarrow G.\gamma_{\text{Tail}} \quad (7.9)$$

$$x_{\text{head}} \dashrightarrow G.\gamma_{\text{Last}} \quad (7.10)$$

From the abstract points-to predicates from the invariant and lemma 19 we get that  $x_{\text{head}} \rightsquigarrow x_{\text{tail}}$ , so by lemma 15, we know that  $x_{\text{head}}$  is a node. This shows 7.7. By reflexivity of `reach` (lemma 13) we additionally know that  $x_{\text{head}} \rightsquigarrow x_{\text{head}}$ . Lemma 20 then gives us 7.8 and 7.9.

Lastly, we use lemma 19 to deduce that  $x_{\text{tail}} \rightsquigarrow x_{\text{last}}$ . By transitivity of `reach` (lemma 14) we have that  $x_{\text{head}} \rightsquigarrow x_{\text{last}}$ , which we use with lemma 20 to get 7.10.

We now close the invariant and step to line 22 which attempts to read  $\ell_{\text{tail}}$ . We open the invariant, which tells us that  $\ell_{\text{tail}}$  points to some  $x_{\text{tail}}$  (not necessarily the same as the previous invariant opening), and we perform the load. From the abstract points-to and `reach` predicates from the invariant together with 7.9 and lemmas 19, 16, and 20 we get the following

$$\text{in}(x_{\text{tail}}) \mapsto^\square (\text{val}(x_{\text{tail}}), \text{out}(x_{\text{tail}})) \quad (7.11)$$

$$x_{\text{head}} \rightsquigarrow x_{\text{tail}} \quad (7.12)$$

$$x_{\text{tail}} \dashrightarrow G.\gamma_{\text{Tail}} \quad (7.13)$$

$$(7.14)$$

We close the invariant and step to line 23. This line creates our *prophecy variable*,  $p$ , which will be resolved on line 25. This allows us to reason about the result of the consistency check: we will later show that the expression associated with  $p$  (i.e. `!(fst(!Q))`) evaluates to some value  $v_p$ , but we can already now case on whether  $v_p$  will be equal to `in`( $x_{\text{head}}$ ) – the left hand side of the equality check on line 25.

**Case** `in`( $x_{\text{head}}$ ) =  $v_p$ . We continue to line 24, which finds out what  $x_{\text{head}}$  points to. As  $x_{\text{head}}$  could be the last node in the linked list, we don't have the relevant points-to predicate. We therefore open the invariant. We get the three nodes  $x'_{\text{head}}$ ,  $x'_{\text{tail}}$ , and  $x_{\text{last}}$ . Specifically,  $x_{\text{last}}$  is the last node in the linked list, and

$$G.\gamma_{\text{Last}} \rightsquigarrow x_{\text{last}} \quad (7.15)$$

Combining this with 7.10 and lemma 19 we conclude  $x_{\text{head}} \rightsquigarrow x_{\text{last}}$ . This gives us two cases to consider: either  $x_{\text{head}} = x_{\text{last}}$  or  $x_{\text{head}}$  points to some  $x_{\text{head\_next}}$ , which reaches  $x_{\text{last}}$ .

**Case**  $x_{\text{head}} = x_{\text{last}}$ . This corresponds to the queue being empty, which we derive below.

As  $x_{\text{last}}$  is the last node, we have that  $\text{out}(x_{\text{head}}) \mapsto \text{None}$ , hence the load resolves to  $\text{None}$ . Using the abstract points-to predicates from the invariant together with 7.8, 7.9, and lemma 19 we get  $x_{\text{head}} \rightsquigarrow x'_{\text{head}}$  and  $x_{\text{head}} \rightsquigarrow x'_{\text{tail}}$ . We can now apply lemma 17 three times to conclude  $x_{\text{head}} = x'_{\text{head}} = x'_{\text{tail}} = x_{\text{tail}}$ . Since  $x'_{\text{head}}$  points to  $\text{None}$ , then  $x_{\text{queue}}$  has to be empty (if it wasn't we could deduce that  $x'_{\text{head}}$  pointed to a node). This also implies that abstract state of the queue  $x_{\text{sv}}$ , is empty,  $x_{\text{sv}} = []$ .

Because the load resolved to  $\text{None}$ , then the variable  $\text{next}$  in the code will be  $\text{None}$ , and since we are in the case where the consistency check passes, and since we have derived that  $x_{\text{head}} = x_{\text{tail}}$ , we know that  $\text{dequeue}$  will return  $\text{None}$ . In other words, this is a linearisation point.

Since  $x_{\text{sv}} = []$ , then our abstract state predicate from the invariant states  $G.\gamma_{\text{Abst}} \models_{\bullet} []$ . We thus instantiate the viewshift with  $[]$ , and supply the  $P$  from the pre-condition. As  $x_{\text{sv}} = []$  we can conclude that the first disjunct must be true (the second contains a contradiction), so we get  $Q(\text{None})$  and  $G.\gamma_{\text{Abst}} \models_{\bullet} []$ . As we haven't changed any resources, we can close the invariant again.

We reach the consistency check on line 25. By 7.6, we know that  $!(\text{fst}(!Q))$  steps to  $!(\ell_{\text{head}})$ , but to resolve the prophecy, we must first show what  $!(\ell_{\text{head}})$  evaluates to. This resource is inside the invariant so we open it. We get that  $\ell_{\text{head}} \mapsto \text{in}(x''_{\text{head}})$  for some node  $x''_{\text{head}}$ . We close the invariant and resolve the prophecy:  $!(\text{fst}(!Q))$  evaluated to  $\text{in}(x''_{\text{head}})$ . In other words,  $v_p = \text{in}(x''_{\text{head}})$ , and therefore  $\text{in}(x_{\text{head}}) = \text{in}(x''_{\text{head}})$ . Since the remaining if statement on line 25 compares  $\text{in}(x_{\text{head}})$  to  $\text{in}(x''_{\text{head}})$ , we know that we will take the “then” branch, so we step to line 26. Since  $x_{\text{head}} = x_{\text{tail}}$  and  $\text{next}$  was set to  $\text{None}$ , we step to line 28 which returns  $\text{None}$ . The post-condition thus requires us to prove  $Q(\text{None})$ , which we already have.

**Case**  $\text{out}(x_{\text{head}}) \mapsto^{\square} \text{in}(x_{\text{head\_next}}) * x_{\text{head\_next}} \rightsquigarrow x_{\text{last}}$ . This means that the queue is not empty, and there is an element to be dequeued: the value in  $x_{\text{head\_next}}$ . The load resolves to  $\text{in}(x_{\text{head\_next}})$ , and the program variable  $\text{next}$  is set to this. Using lemmas 15 and 20 with 7.15 we get

$$\text{in}(x_{\text{head\_next}}) \mapsto^{\square} (\text{val}(x_{\text{head\_next}}), \text{out}(x_{\text{head\_next}})) \quad (7.16)$$

$$x_{\text{head\_next}} \dashrightarrow G.\gamma_{\text{Last}} \quad (7.17)$$

We close the invariant and step to the consistency check on line 25. We handle this similarly to the previous case and conclude that the consistency check succeeds, taking us to line 26 in the “then” branch. This line ensures that the  $\text{dequeue}$  will not make the tail node lag behind the head node. We can simply consider both cases of the check.

The case where the “if” succeeds takes us to the **CAS** on line 30, which attempts to swing the tail, and try dequeuing again. We handle the **CAS** with our swing-tail lemma (lemma 23), and the recursive call by the induction hypothesis.

If the “if” fails, then the tail node will not lag behind as a result of the  $\text{dequeue}$ , so we step to the else block on line 31 which attempts to  $\text{dequeue}$ . We first read the value out of  $x_{\text{head\_next}}$  on line 32. Next, we attempt to swing the head pointer on line 33. The rule for **CAS** demands a points-to predicate for  $\ell_{\text{head}}$ , so we open the invariant which gives us fresh nodes  $x'_{\text{head}}$ ,  $x'_{\text{tail}}$ , and  $x'_{\text{last}}$ , so that  $\ell_{\text{head}} \mapsto \text{in}(x'_{\text{head}})$ . The success of the **CAS** depends on whether  $\text{in}(x'_{\text{head}})$  equals  $\text{in}(x_{\text{head}})$ . If they aren't equal, the **CAS** fails and nothing is updated. We can thus close the invariant and we step to the recursive call on line 35 and apply the induction hypothesis. So for the remainder, we assume they are equal and the **CAS** succeeds. Since the **CAS** moved the head pointer to  $x_{\text{head\_next}}$ , the queue data structure got updated, so this is a linearisation point.

Using the abstract points-to and reachability propositions from the invariant together with lemmas 19 and 15, we may deduce that  $x'_{\text{head}}$  is a node. As we already know that  $x_{\text{head}}$  is a node (from 7.7), lemma 33 tells us that the nodes are equal:  $x_{\text{head}} = x'_{\text{head}}$ . From the invariant (specifically  $\text{isLL}(xs)$ ) we know that  $x'_{\text{head}}$  points to the first element of  $x_{\text{queue}}$ . But since  $x_{\text{head}} = x'_{\text{head}}$ , then this element must be our  $x_{\text{head\_next}}$ . In other words,  $x_{\text{queue}} = x_{\text{queue}}' ++ [x_{\text{head\_next}}]$ , for some  $x_{\text{queue}}'$ . This means that, when we apply the viewshift (giving up  $P$  and  $G.\gamma_{\text{Abst}} \models_{\bullet} x_{\text{sv}}$  as usual), only the second case of the resulting disjunct is possible:

$xs_v$  cannot be empty as  $xs_{\text{queue}}$  isn't. We therefore get that there are some  $xs'_v$  and  $v$  so that

$$xs_v = xs'_v ++ [v] \quad (7.18)$$

$$G.\gamma_{\text{Abst}} \Rightarrow_{\bullet} xs'_v \quad (7.19)$$

$$Q(\text{Some } v) \quad (7.20)$$

Since  $xs_v$  is reflected in  $xs_{\text{queue}}$  (according to the invariant), we may additionally conclude that  $xs'_v$  is reflected in  $xs'_{\text{queue}}$  and  $\text{Some } v = \text{val}(x_{\text{head\_next}})$ .

To close the invariant, we must update some of our resources. Since we now have  $\ell_{\text{head}} \mapsto \text{in}(x_{\text{head\_next}})$ , we must pick  $x_{\text{head\_next}}$  for the head node. But currently  $G.\gamma_{\text{Head}} \mapsto x_{\text{head}}$ . So we use lemma 21 to advance the pointer, and we get  $G.\gamma_{\text{Head}} \mapsto x_{\text{head\_next}}$ .

We are now required to show that  $x_{\text{head\_next}} \dashrightarrow G.\gamma_{\text{Tail}}$ . Since  $x_{\text{head}} \leadsto x_{\text{tail}}$ , and  $x_{\text{head}} \neq x_{\text{tail}}$ , then it must be the case that  $x_{\text{head\_next}} \leadsto x_{\text{tail}}$ . Lemma 19 with 7.13 now tells us that  $x_{\text{tail}}$  can reach the current tail node,  $x'_{\text{tail}}$ . By transitivity (lemma 14), we get  $x_{\text{head\_next}} \leadsto x'_{\text{tail}}$ , and hence by lemma 20, we get the desired  $x_{\text{head\_next}} \dashrightarrow G.\gamma_{\text{Tail}}$ . We now own all the resources required to close the invariant.

As the CAS succeed, we step to line 34 which simply returns  $\text{val}(x_{\text{head\_next}})$ . Thus, we must prove the post-condition  $Q(\text{val}(x_{\text{head\_next}}))$ . We can do this since we still own 7.20 and we deduced that  $\text{Some } v = \text{val}(x_{\text{head\_next}})$ .

**Case**  $\text{in}(x_{\text{head}}) \neq v_p$ . We step to line 24 which finds out what  $x_{\text{head}}$  points to. To access the relevant points-to predicate, open the invariant. Importantly, we get that there is some last node of the linked list,  $x_{\text{last}}$ , with  $G.\gamma_{\text{Last}} \mapsto x_{\text{last}}$ . By combining this with 7.10 and lemma 19 we deduce that  $x_{\text{head}} \leadsto x_{\text{last}}$  which means that either  $x_{\text{head}}$  is the last node, and hence  $\text{out}(x_{\text{head}}) \mapsto \text{None}$ , or there is some other node  $x_{\text{head\_next}}$  and  $\text{out}(x_{\text{head}}) \mapsto^{\square} x_{\text{head\_next}}$ . This shows that the load is safe. The actual value that the load resolves to is unimportant, as it won't be used in this case. We thus perform the load, close the invariant and step to line 25.

By 7.6, we know that  $!(\text{fst}(!Q))$  steps to  $!(\ell_{\text{head}})$ , so to resolve the prophecy, we show what  $!(\ell_{\text{head}})$  evaluates to. We open the invariant, which gives us that  $\ell_{\text{head}} \mapsto \text{in}(x'_{\text{head}})$  for some node  $x'_{\text{head}}$ . We close the invariant again, and resolve the prophecy:  $!(\text{fst}(!Q))$  evaluated to  $\text{in}(x'_{\text{head}})$ . That is,  $v_p = \text{in}(x'_{\text{head}})$ , and therefore  $\text{in}(x_{\text{head}}) \neq \text{in}(x'_{\text{head}})$ . We hence take the “else” branch and step to line 36. This consists of a recursive call to the loop function and we are done by the induction hypothesis. □

## 7.6 Discussion

It was shown in Vindum and Birkedal [2021] that a version of the Lock-Free M&S Queue with the consistency checks is contextually equivalent to a version without consistency checks. In the original presentation [Michael and Scott, 1996] the implementation language was assumed to not have a garbage collector. This meant that the ABA problem was an issue; if a node is freed by a dequeue, and afterwards a new node is allocated at the same location with the same value through an enqueue, it will look as though it is the exact same node. This can cause inconsistencies for threads that read the original node, went to sleep, and continued after the new allocation. To fix the issue, the authors added *modification counters* to their pointers, which means that the pointer to the newly allocated node will have a higher counter, and we can hence tell that it isn't the same node as the original. The essence of the consistency checks is to ensure that previously read nodes are the exact same nodes as from before, by ensuring that the counter is the same. We do this check *after* reading any “next” nodes to ensure that they are indeed the next of the node we originally read – the original node and its next node are consistent.

However, the language used in Vindum and Birkedal [2021] and this project, HeapLang, is a garbage-collected language. This means that we do not need to worry about freeing nodes, and the ABA problem doesn't occur. In other words, when we read the “next” of a node, we can be certain that it is the next of the node we originally read – no one could have freed it in a dequeue it and subsequently allocated a similar looking node in an enqueue. This then means that the consistency checks are no longer needed.



Indeed, in the Coq formalisations, we prove the Hocap-style specification for an implementation that doesn't have the consistency checks. As an additional benefit, removing the consistency check means that in dequeue, we know already when we read the “next” of the head node on line 24 whether or not the dequeue will conclude that the queue is empty, and hence whether or not the load is a linearisation point. The prophecy variable is thus not needed, which further simplifies the proof of the specification.

## Chapter 8

# On the Resource Algebras Used

### 8.1 Tokens

Tokens are defined using the exclusive resource algebra on the singleton set:  $\text{Ex}()$ . This resource algebra only has one valid element, which we here denote by  $T$ , and combining two elements will give the non-valid element  $\perp$ .

We can always create new tokens, which will just give us ownership of  $T$  with a fresh ghost name,  $\gamma$ .

**Lemma 26.**  $\vdash \models \exists \gamma. [\bar{T}]^\gamma$

*Proof.* This follows directly from  $\text{ghost-alloc} \blacktriangleright \text{ref} \blacktriangleleft$  and the fact that  $T$  is valid.  $\square$

Importantly, if we own a particular token, then we know that no one else can own the same token. Formally, we write this as:

**Lemma 27.**  $[\bar{T}]^\gamma * [\bar{T}]^\gamma \vdash \text{False}$

*Proof.* By  $\text{own-op} \blacktriangleright \text{ref} \blacktriangleleft$  we get ownership of  $[\bar{T} \cdot \bar{T}]^\gamma$ . By the exclusive RA, this is equivalent to  $[\bar{\perp}]^\gamma$ . We now own a non-valid element which proves **False**.  $\square$

### 8.2 Abstract State

We shall use the resource algebra  $\text{AUTH}((\text{FRAC} \times \text{AG}(\text{List Val}))^?)$  to define the abstract state predicate.  $\text{List Val}$  is the abstract state. It is wrapped in the agreement RA,  $\text{AG}$ , which ensures that if one owns two elements, then they agree on the abstract state. The fractional RA,  $\text{FRAC}$ , denotes how much of the fragmental view is owned; the fragmental view can be split up, which is handled by the clients. We collect  $\text{FRAC}$  and  $\text{AG}(\text{List Val})$  in the product RA, whose elements are then pairs of fractions and abstract states. The option RA  $?$ , makes the product RA unital which is required by the  $\text{AUTH}$  construction.  $\text{AUTH}$  is the authoritative resource algebra. It gives us the authoritative and fragmental views and, together with the fractional RA, governs that they can only be updated in unison.

We note that core of the fractional RA is always undefined, which means that the core of  $\text{AUTH}((\text{FRAC} \times \text{AG}(\text{List Val}))^?)$  is also undefined. Hence no elements of the RA can be persistent – not even the fragmental ones.

With this, we may now define  $\gamma \models_\bullet xs_v$  and  $\gamma \models_\circ xs_v$  formally.

**Definition 8.2.1** (Authoritative Abstract State). For an abstract state  $xs_v$  and a ghost name  $\gamma$ , we define:  $\gamma \models_\bullet xs_v \triangleq [\bullet(1, \text{ag } xs_v)]^\gamma$

**Definition 8.2.2** (Fragmental Abstract State). For an abstract state  $xs_v$  and a ghost name  $\gamma$ , we define:  $\gamma \models_\circ xs_v \triangleq [\circ(1, \text{ag } xs_v)]^\gamma$

**Lemma 1** (Abstract State Alloc).  $\forall xs_v. \vdash \models \exists \gamma. \gamma \models_\bullet xs_v * \gamma \models_\circ xs_v$

*Proof.* By own-op $\blacktriangleright$ ref $\blacktriangleleft$ , it suffices to show  $\models \exists \gamma. [\bullet(1, \text{ag } xs_v) \cdot \circ(1, \text{ag } xs_v)]^\gamma$ . This follows by ghost-alloc $\blacktriangleright$ ref $\blacktriangleleft$  if we can show that the element is valid. That is,  $\bullet(1, \text{ag } xs_v) \cdot \circ(1, \text{ag } xs_v) \in \mathcal{V}$ . However, this follows by the definitions of the involved resource algebras.  $\square$

In section 3.5 we postulated two important lemmas which showed that the authoritative and fragmental states agree, and how to update the abstract state. We now proceed to prove these lemmas.

**Lemma 2** (Abstract State Agree).  $\forall \gamma, xs_v, xs'_v.$

$$\gamma \models_\bullet xs'_v * \gamma \models_\circ xs_v \vdash xs_v = xs'_v$$

*Proof.* Since we own both  $\gamma \models_\bullet xs_v$  and  $\gamma \models_\circ xs'_v$ , and they have the same ghost name, we can use rules own-op $\blacktriangleright$ ref $\blacktriangleleft$  and own-valid $\blacktriangleright$ ref $\blacktriangleleft$  to conclude that the element  $\bullet(1, \text{ag } xs_v) \cdot \circ(1, \text{ag } xs'_v)$  is valid. By the authoritative RA, this means that  $\text{ag } xs'_v \preceq \text{ag } xs_v$ . By definition of the agreement RA, this means that  $xs'_v = xs_v$ , which is what we wanted.  $\square$

**Lemma 3** (Abstract State Update).  $\forall \gamma, xs_v, xs'_v, xs''_v.$

$$\gamma \models_\bullet xs'_v * \gamma \models_\circ xs_v \Rightarrow \gamma \models_\bullet xs''_v * \gamma \models_\circ xs''_v$$

*Proof.* This time, we use own-op $\blacktriangleright$ ref $\blacktriangleleft$  to conclude  $[\bullet(1, \text{ag } xs_v) \cdot \circ(1, \text{ag } xs'_v)]^\gamma$ . Further, by own-op $\blacktriangleright$ ref $\blacktriangleleft$ , it suffices to prove  $\models [\bullet(1, \text{ag } xs'_v) \cdot \circ(1, \text{ag } xs''_v)]^\gamma$ . We do this by applying the ghost-update $\blacktriangleright$ ref $\blacktriangleleft$  rule, which requires us to prove  $\bullet(1, \text{ag } xs_v) \cdot \circ(1, \text{ag } xs'_v) \rightsquigarrow \bullet(1, \text{ag } xs'_v) \cdot \circ(1, \text{ag } xs''_v)$ . Firstly by the product, fractional, and agreement RA, the element  $(1, \text{ag } xs'_v)$  is valid, so the product of the authoritative and fragmental parts of it is also valid. Next, note that we own the entire fraction of the fragmental element. Hence, there can be no other valid fragments. It hence follows that we can do the frame-preserving update.  $\square$

### 8.3 Abstract Points-to and Reachability

To define the abstract points-to predicate and the abstract reach predicate, we create the following resource algebra:  $\text{AUTH}(\mathcal{P}(\text{Node}))$ , where  $\text{Node} \triangleq (\text{Loc} \times \text{Val}) \times \text{Loc}$ . Here, the resource algebra  $\mathcal{P}(\text{Node})$  denotes the set of subsets of  $\text{Node}$ , with union as the operation. The empty set is the unit element, meaning that  $\mathcal{P}(\text{Node})$  is unital. We may now define abstract reach and abstract points-to as follows.

**Definition 8.3.1** (Abstract Reach).  $x \dashrightarrow \gamma \triangleq [\circ\{x\}]^\gamma$

**Definition 8.3.2** (Abstract Points-to).  $\gamma \mapsto x \triangleq \exists s. [\bullet s]^\gamma * \bigstar_{x_m \in s} x_m \rightsquigarrow x$

One should think of sets  $s \in \mathcal{P}(\text{Node})$  as specifying which nodes can abstractly reach a certain ghost variable. Due to how the authoritative resource algebra work, the assertion  $[\circ\{x\}]^\gamma$  essentially states that  $x$  is *one* of the nodes that can reach the node that  $\gamma$  points to. This is because, when combining a fragmental element  $s$  and authoritative element  $t$ , we get that the fragmental element is “smaller” than the authoritative,  $s \preceq t$ . In this case, “smaller” amounts to “subset”. Hence,  $[\circ\{x\}]^\gamma$  means that, whatever the authoritative set is, it will contain the node  $x$ .

The authoritative set is existentially quantified as it can change over time, but whatever it is, we assert that all the nodes it contains can *concretely* reach the node that the ghost name is currently pointing to. This choice of definitions enables us to prove the properties we desired of abstract reachability in section 7.2. We proceed to prove the four essential lemmas on abstract points-to and reachability that we postulated in said section.

**Lemma 18** (Abs-Reach-Alloc).  $x \rightsquigarrow x \Rightarrow (\exists \gamma. \gamma \mapsto x * x \dashrightarrow \gamma)$

*Proof.* We assume  $x \rightsquigarrow x$  and must show  $\models \exists \gamma'. \gamma' \mapsto x * x \dashrightarrow \gamma'$ . By definition or the authoritative RA, the element  $\bullet\{x\} \cdot \circ\{x\}$  is valid. Hence by the Ghost-alloc rule $\blacktriangleright$ ref $\blacktriangleleft$ , we may get  $\models \exists \gamma. [\bullet\{x\} \cdot \circ\{x\}]^\gamma$ . By Upd-mono $\blacktriangleright$ ref $\blacktriangleleft$ , we may strip away the update modality on the goal and the previous assertion. Thus, we must prove  $\exists \gamma'. \gamma' \mapsto x * x \dashrightarrow \gamma'$ , and we have that  $\exists \gamma. [\bullet\{x\} \cdot \circ\{x\}]^\gamma$ . We use  $\gamma$  as the witness in the goal meaning we must prove  $\gamma \mapsto x * x \dashrightarrow \gamma$ . We can split the ownership of the authoritative and fragmental parts up using Own-op $\blacktriangleright$ ref $\blacktriangleleft$ , giving us  $[\bullet\{x\}]^\gamma$  and  $[\circ\{x\}]^\gamma$ . The latter assertion is equivalent to  $x \dashrightarrow \gamma$ , which matches the second obligation in the goal. To prove the first obligation, we must give some set as witness, and show that all nodes in the set can reach  $x$ . We of course choose  $\{x\}$  as the witness, and must then prove that  $x \rightsquigarrow x$ , which we assumed in the beginning.  $\square$

**Lemma 19** (Abs-Reach-Concr).  $x_n \dashrightarrow \gamma_m * \gamma_m \rightarrow x_m \multimap x_n \leadsto x_m * \gamma_m \rightarrow x_m$

*Proof.* Assuming  $x_n \dashrightarrow \gamma_m$  and  $\gamma_m \rightarrow x_m$ , we must show  $x_n \leadsto x_m$  without “consuming”  $\gamma_m \rightarrow x_m$ . From  $\gamma_m \rightarrow x_m$  we can deduce that there is some set  $s$  so that  $\{\bullet s\}^{\gamma_m}$  and  $\bigstar_{x' \in s} x' \leadsto x_m$ . Since we own both  $\{\bullet s\}^{\gamma_m}$  and  $\{\circ \{x_n\}\}^{\gamma_m}$  (from  $x_n \dashrightarrow \gamma_m$ ), we may conclude that their product is valid, which in our instantiation of the authoritative RA equates to  $x_n \in s$ . We may now frame away the second part of the goal,  $\gamma_m \rightarrow x_m$ , using  $\{\bullet s\}^{\gamma_m}$  and  $\bigstar_{x' \in s} x' \leadsto x_m$ . Note that we get to keep  $\{\bullet s\}^{\gamma_m}$  as we only used it to prove persistent facts. Thus, from that assertion and by  $x_n \in s$ , we can deduce that  $x_n \leadsto x_m$ , which is what we had to prove.  $\square$

**Lemma 20** (Abs-Reach-Abs).  $x_n \leadsto x_m * \gamma_m \rightarrow x_m \Rightarrow x_n \dashrightarrow \gamma_m * \gamma_m \rightarrow x_m$

*Proof.* Assuming  $x_n \leadsto x_m$  and  $\gamma_m \rightarrow x_m$  we must conclude  $\models x_n \dashrightarrow \gamma_m$ . From  $\gamma_m \rightarrow x_m$  we know that there is some set  $s$  so that  $\{\bullet s\}^{\gamma_m}$  and  $\bigstar_{x' \in s} x' \leadsto x_m$ . There are now two cases to consider: either  $x_n \in s$  or  $x_n \notin s$ .

$x_n \in s$  By the definition of our authoritative RA, if a set  $y$  is a subset of  $s$ , then we may update our ghost resources to obtain ownership of the fragment  $y$ . In our case, since  $x_n \in s$ , we may update our resources to additionally get  $\{\circ \{x_n\}\}^{\gamma_m}$ , which is exactly what we wanted. Since we still have  $\{\bullet s\}^{\gamma_m}$ , we can also prove  $\gamma_m \rightarrow x_m$ .

$x_n \notin s$  In this case we may update  $\{\bullet s\}^{\gamma_m}$  so that the set also includes  $x_n$ . The reason we may do this, is because, according to the  $\mathcal{P}(\text{Node})$  RA, we may update a set  $X$  to  $Y$ , as long as  $X \subseteq Y$ . Thus, we can update our resource to get  $\{\bullet \{x_n\} \cup s\}^{\gamma_m}$ . As in the previous case, we can further get  $\{\circ \{x_n\}\}^{\gamma_m}$  out of this, which we use to frame away the goal  $x_n \dashrightarrow \gamma_m$ .

To prove  $\gamma_m \rightarrow x_m$ , we use the set  $\{x_n\} \cup s$ , and immediately frame away the authoritative part, which we owned. We are left with having to prove  $\bigstar_{x' \in \{x_n\} \cup s} x' \leadsto x_m$ . However, by  $\bigstar_{x' \in s} x' \leadsto x_m$  and our assumption that  $x_n \leadsto x_m$ , we can easily conclude this.  $\square$

**Lemma 21** (Abs-Reach-Advance).  $\gamma_m \rightarrow x_m * x_m \leadsto x_o \Rightarrow \gamma_m \rightarrow x_o * x_o \dashrightarrow \gamma_m$

*Proof.* Assuming  $\gamma_m \rightarrow x_m$  and  $x_m \leadsto x_o$  we must prove  $\models \gamma_m \rightarrow x_o * x_o \dashrightarrow \gamma_m$ . From  $\gamma_m \rightarrow x_m$ , we get some set  $s$  so that  $\{\bullet s\}^{\gamma_m}$  and  $\bigstar_{x' \in s} x' \leadsto x_m$ . As we did in the proof of lemma 20, we update  $\{\bullet s\}^{\gamma_m}$  so that the set additionally contains  $x_o$ . Thus, we get  $\{\bullet \{x_o\} \cup s\}^{\gamma_m}$ . From this, we may extract ownership of the fragmental part:  $\{\circ \{x_o\}\}^{\gamma_m}$ , which we use to prove the second part of the goal.

We are thus left with proving  $\gamma_m \rightarrow x_o$ . We use  $\{x_o\} \cup s$  as witness for the authoritative set, and immediately frame away the ownership assertion of the authoritative part. We are left with proving  $\bigstar_{x' \in \{x_o\} \cup s} x' \leadsto x_o$ . We already know that  $\bigstar_{x' \in s} x' \leadsto x_m$  and  $x_m \leadsto x_o$ . Thus, by transitivity of reach (lemma 14), we may conclude  $\bigstar_{x' \in \cup s} x' \leadsto x_o$ . Thus, we are done if we can prove that  $x_o \leadsto x_o$ , which by lemma 13 amounts to showing that  $x_o$  is a node. However, since  $x_m \leadsto x_o$ , then, by lemma 16, we know that this is the case.  $\square$

## Chapter 9

# Conclusion and Future Work

►conclude on the problem statement from the introduction◄

►Mention the possibility of simplifying queue invariant for lock-free by removing isLL (and adding x\_last -> None)◄

# Bibliography

- Lars Birkedal and Aleš Bizjak. Lecture notes on iris: Higher-order concurrent separation logic. <http://iris-project.org/tutorial-pdfs/iris-lecture-notes.pdf>, 2017.
- Maurice Herlihy and Jeannette M. Wing. Linearizability: A correctness condition for concurrent objects. *ACM Trans. Program. Lang. Syst.*, 12(3):463–492, 1990. doi: 10.1145/78969.78972. URL <https://doi.org/10.1145/78969.78972>.
- Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek Dreyer. Iris from the ground up: A modular foundation for higher-order concurrent separation logic. *J. Funct. Program.*, 28:e20, 2018. doi: 10.1017/S0956796818000151. URL <https://doi.org/10.1017/S0956796818000151>.
- Maged M. Michael and Michael L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In James E. Burns and Yoram Moses, editors, *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, pages 267–275. ACM, 1996. doi: 10.1145/248052.248106. URL <https://doi.org/10.1145/248052.248106>.
- Simon Friis Vindum and Lars Birkedal. Contextual refinement of the michael-scott queue (proof pearl). In Catalin Hritcu and Andrei Popescu, editors, *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, pages 76–90. ACM, 2021. doi: 10.1145/3437992.3439930. URL <https://doi.org/10.1145/3437992.3439930>.

# Appendix A

## Common Definitions and Lemmas

### A.1 Lists and Nodes

**Definition A.1.1** (First in Queue).  $\text{isFirst}(x, xs) \triangleq \exists xs_{rest}. xs = xs_{rest} ++ [x]$

**Definition A.1.2** (Last in Queue).  $\text{isLast}(x, xs) \triangleq \exists xs_{rest}. xs = x :: xs_{rest}$

**Definition A.1.3** (Second Last in Queue).  $\text{isSndLast}(x, xs) \triangleq \exists x_{last}, xs_{rest}. xs = x_{last} :: x :: xs_{rest}$

**Lemma 28** (Adding/Removing Non-Last).  $\forall x, y, xs, ys.$   
 $\text{isLast}(x, (xs ++ [y] ++ ys)) \iff \text{isLast}(x, (xs ++ [y]))$

**Lemma 29** (List Destruction Equality).  $\forall xs_1, x_1, xs_2, x_2.$   
 $xs_1 ++ [x_1] = xs_2 ++ [x_2] \implies xs_1 = xs_2 \wedge x_1 = x_2$

**Definition A.1.4** (Value Projection).

$$\begin{aligned} \text{projVal}([]) &\triangleq [] \\ \text{projVal}(x :: xs) &\triangleq \text{val}(x) :: \text{projVal}(xs) \end{aligned}$$

**Lemma 30** (Value Projection Split).  $\forall xs_1, xs_2.$   
 $\text{projVal}(xs_1 ++ xs_2) = \text{projVal}(xs_1) ++ \text{projVal}(xs_2)$

**Definition A.1.5** (Wrap Some).

$$\begin{aligned} \text{wrapSome}([]) &\triangleq [] \\ \text{wrapSome}(x :: xs) &\triangleq \text{Some } x :: \text{wrapSome}(xs) \end{aligned}$$

**Lemma 31** (Wrap Some Split).  $\forall xs_1, xs_2.$   
 $\text{wrapSome}(xs_1 ++ xs_2) = \text{wrapSome}(xs_1) ++ \text{wrapSome}(xs_2)$

**Definition A.1.6** (The “All” Predicate).

$$\begin{aligned} \text{All}([], \Phi) &\triangleq \text{True} \\ \text{All}(x :: xs, \Phi) &\triangleq \Phi(x) * \text{All}(xs) \end{aligned}$$

**Lemma 32** (All Split).  $\forall xs_1, xs_2, \Phi.$   
 $\text{All}(xs_1 ++ xs_2, \Phi) ** \text{All}(xs_1, \Phi) ++ \text{All}(xs_2, \Phi)$

**Lemma 33** (Node Equality).

$$\begin{aligned} &\forall x, y. \\ &\text{in}(x) = \text{in}(y) \multimap \\ &\text{in}(x) \mapsto^\square (\text{val}(x), \text{out}(x)) \multimap \\ &\text{in}(y) \mapsto^\square (\text{val}(y), \text{out}(y)) \multimap \\ &x = y \end{aligned}$$

## A.2 Results About the isLL Predicate

**Lemma 34** (Extract Chain from isLL).  $\forall xs. \text{isLL}(xs) \multimap \text{isLL}(xs) * \text{isLL\_chain}(xs)$

**Lemma 35** (isLL\_chain Nodes).  $\forall xs_1, x, xs_2.$   
 $\text{isLL\_chain}(xs_1 ++ [x] ++ xs_2) \multimap \text{in}(x) \mapsto^\square (\text{val}(x), \text{out}(x))$

**Lemma 36** (isLL\_chain Split).  $\forall xs, ys.$   
 $\text{isLL\_chain}(xs ++ ys) \multimap \text{isLL\_chain}(xs) * \text{isLL\_chain}(ys)$

**Lemma 37** (isLL Split).  $\forall xs, ys.$   
 $\text{isLL}(xs ++ ys) \multimap \text{isLL}(xs) * \text{isLL\_chain}(ys)$