
TITLE HERE

Mathias Pedersen, 201808137

Master's Thesis, Computer Science

March 2024

Advisor: Amin Timany



AARHUS
UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE

Abstract

► in English... ◄

Resumé

► in Danish... ◄

Acknowledgments



Mathias Pedersen
Aarhus, March 2024.

Contents

Abstract	iii
Resumé	v
Acknowledgments	vii
1 Introduction	1
2 Preliminaries	3
3 The Two-Lock Michael Scott Queue	5
3.1 Preliminaries	5
3.2 implementation	6
3.2.1 initialise	6
3.2.2 enqueue	6
3.2.3 dequeue	6
3.3 Sequential Specification	8
3.4 Proving the Sequential Specification	9
3.4.1 The isqueue predicate	9
3.4.2 Proof outline	10
3.5 Concurrent Specification	12
3.6 Proving the Concurrent Specification	13
3.6.1 The isqueue predicate	13
3.6.2 Proof outline	16
4 Conclusion	19
Bibliography	21
A The Technical Details	23

Chapter 1

Introduction

►motivate and explain the problem to be addressed◄

►example of a citation: [1]◄ ►get your bibtex entries from <https://dblp.org/>◄

Chapter 2

Preliminaries

►Description of HeapLang, Iris, Verified in Coq (Weakest precondition vs Hoare triples)◄

Chapter 3

The Two-Lock Michael Scott Queue

I present here an implementation of the Two-lock MS-Queue in HeapLang. This implementation differs slightly from the original, presented in [1], but most changes simply reflect the differences in the two languages.

3.1 Preliminaries

The underlying data structure making up the queue is a singly-linked list. The linked-list will always contain at least one element, called the *sentinel* node, marking the beginning of the queue. Note that the sentinel node is itself not part of the queue, but all nodes following it are. The queue keeps a head pointer (ℓ_{head}) which always points to the sentinel, and a tail pointer (ℓ_{tail}) which points to some node in the linked list.

In my implementation, a node can be thought of as a triple $(\ell_{i_in}, v_i, \ell_{i_out})$. The location ℓ_{i_in} points to the pair (v_i, ℓ_{i_out}) , where v_i is the value of the node, and ℓ_{i_out} either points to None which represents the null pointer, or to the next node in the linked list. When we say that a location ℓ points to a node $(\ell_{i_in}, v_i, \ell_{i_out})$, we mean that $\ell \mapsto \ell_{i_in}$. Hence, if we have two adjacent nodes $(\ell_{i_in}, v_i, \ell_{i_out})$, $(\ell_{i+1_in}, v_{i+1}, \ell_{i+1_out})$ in the linked list, then we have the following structure: $\ell_{i_in} \mapsto (v_i, \ell_{i_out})$, $\ell_{i_out} \mapsto \ell_{i+1_in}$, and $\ell_{i+1_in} \mapsto v_{i+1}, \ell_{i+1_out}$.

The reader may wonder why there is an extra, intermediary "in" pointer, between the pairs of the linked list, and why the "out" pointer couldn't point directly to the next pair. In the original implementation [1], nodes are allocated on the heap. To simulate this in HeapLang, when creating a new node, we create a pointer to a pair making up the node. Now, in the C-like language used in the original specification, an assignment operator is available which is not present in HeapLang. So in order to mimic this behaviour, we model variables as pointers. In this way, we can model a variable x as a location ℓ_x , and the value stored at ℓ_x is the current value of x . This means that the variable ℓ_{i_out} (called "next" in the original) becomes a location ℓ_{head} , and the value stored at the location is what head is currently assigned to. Since ℓ_{i_out} is supposed to be a variable containing a pointer, then the value saved at that location will also be a pointer.

3.2 implementation

The queue consists of 3 functions: initialize, enqueue, and dequeue which I now present in turn.

3.2.1 initialize

initialize will first create a single node – the sentinel – marking the start of the linked list. It then creates two locks, H_lock and T_lock , protecting the head and tail pointers, respectively. Finally, it creates the head and tail pointers, both pointing to the sentinel. The queue is then a pointer to a structure containing the head, the tail, and the two locks.

Figure 3.1 illustrates the structure of the queue after initialisation. Note that one of the pointers is coloured blue. This represents a *persistent* pointer; a pointer that will never be updated again. All "in" pointers ℓ_{i_in} , are persistent, meaning that they will always point to (v_i, ℓ_{i_out}) . We shall use the notation $\ell \mapsto \Box v$ (introduced in [2]) to mean that ℓ points persistently to v .

Note that in the original specification, a queue is a pointer to a 4-tuple $(\ell_{head}, \ell_{tail}, H_lock, T_lock)$. Since HeapLang doesn't support 4-tuples, we instead represent the queue as a pointer to a pair of pairs: $((\ell_{head}, \ell_{tail}), (H_lock, T_lock))$.

3.2.2 enqueue

To enqueue a value, we must create a new node, append it to the underlying linked-list, and swing the tail pointer to this new node. These three operations are depicted in figure 3.2.

enqueue takes as argument the value to be enqueued and creates a new node containing this value (corresponding to figure 3.2a). This creation doesn't interact with the underlying queue data-structure, hence why we don't acquire the T_lock first. After creating the new node, we must make the last node in the linked list point to it. Since this operation interacts with the queue, we first acquire the T_lock . Once we obtain the lock, we make the last node in the linked list point to our new node (figure 3.2b). Following this, we swing ℓ_{tail} to the new last node in the linked list (figure 3.2c).

Figure 3.2 also illustrates when pointers become persistent; once the previous last node is updated to point to the newly inserted node, that pointer will never be updated again, hence becoming persistent.

3.2.3 dequeue

It is of course only possible to dequeue an element from the queue if the queue contains at least one element. Hence, the first thing dequeue does is check if the queue is empty. We can detect an empty queue by checking if the sentinel is the last node in the linked list. Being the last node in the linked list corresponds to having the "out" node be None. If this is the case, then the queue is empty and the code returns None. Otherwise, there is a node just after the sentinel, which is the first node of the queue. To dequeue it, we first read the associated value, and next we swing the head to it, making it the new sentinel. Finally, we return the value we read.

Since all of these operations interact with the queue, we shall only perform them after having acquired H_lock .

Figure 3.3 illustrates running dequeue on a non-empty queue. Note that the only change is that the head pointer is swung to the next node in the linked list; the old sentinel is not deleted, it just become unreachable from the heap pointer. In this way, the linked list only ever grows.

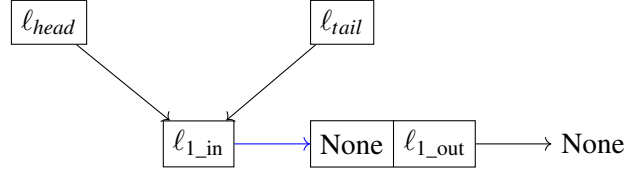
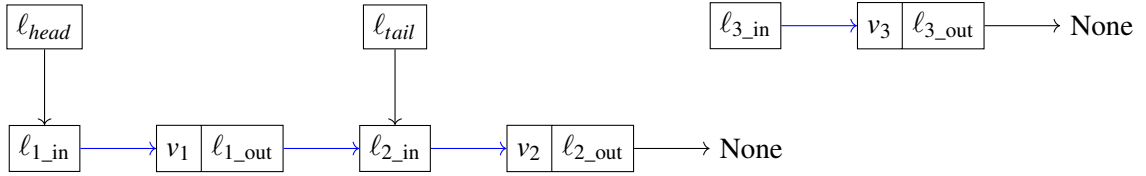
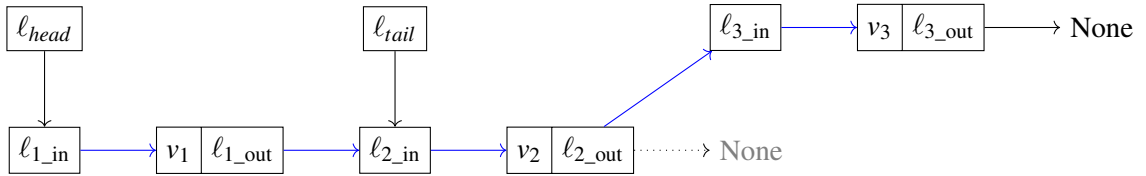


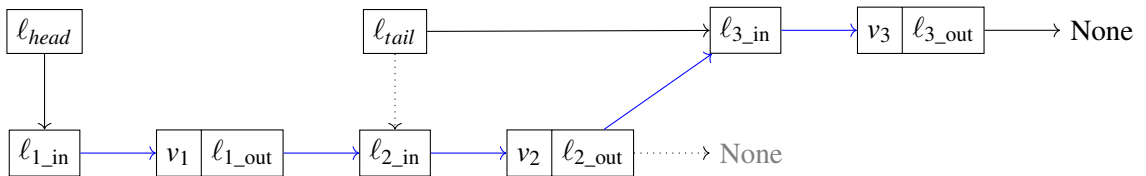
Figure 3.1: Queue after initialisation



(a) Queue after creating the new node $(\ell_{3_in}, v_3, \ell_{3_out})$ to be added to the queue.



(b) Queue after adding the new node to linked list.



(c) Queue after swinging tail pointer to the new node.

Figure 3.2: Enqueuing an element to a queue with one element.

let initialize :=

```

let node = ref ((None, ref (None))) in
let H_lock = newlock() in
let T_lock = newlock() in
ref ((ref (node), ref (node)), (H_lock, T_lock))

```

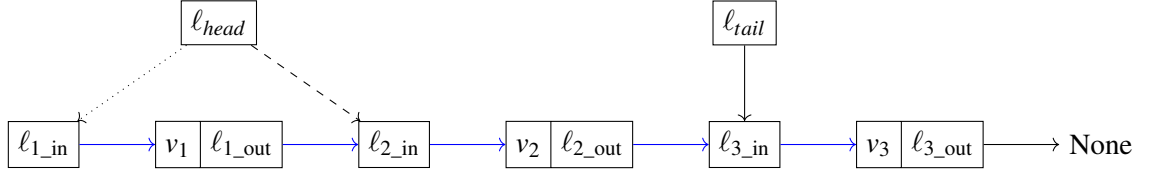


Figure 3.3: Dequeueing an element (v_2) from a queue with two elements (v_2, v_3). The dotted line represents the state before the dequeue, and the dashed line is the state after dequeuing.

```

let enqueue  $Q$  value :=
  let node = ref ((Some value, ref (None))) in
  acquire(snd(snd(! $Q$ )));
  snd(!(!snd(fst(! $Q$ ))))  $\leftarrow$  node;
  snd(fst(! $Q$ ))  $\leftarrow$  node;
  release(snd(snd(! $Q$ )))

```

```

let dequeue  $Q$  :=
  acquire(fst(snd(! $Q$ )));
  let node = !(fst(fst(! $Q$ ))) in
  let new_head = !(snd(!node)) in
  if new_head = None then
    release(fst(snd(! $Q$ )));
    None
  else
    let value = fst(!new_head) in
    fst(fst(! $Q$ ))  $\leftarrow$  new_head;
    release(fst(snd(! $Q$ )));
    value

```

3.3 Sequential Specification

Let us first prove a specification for the two-lock michael scott queue in the simple case where we don't allow for concurrency. In this case, we know that only a single thread will interact with the queue at any given point in a sequential manner. This means that we give a specification that tracks the exact contents of the queue. To this end, we shall define the abstract state of the queue, denoted xs_v as a list of HeapLang values. I.e. $xs_v : List\ Val$. We adopt the convention that enqueueing an element is done by adding it to the front of the list, and dequeuing removes the last element of the list (if such an element exists). The reason for this choice is purely technical.

Since the queue uses two locks, we will get two ghost names; one for each lock. For this specification, these are the only two ghost names we will need. However, for the later specifications, we will use more resource algebra, and will need more ghost names. Thus, to ease notation, we shall define the type "*Qghostnames*" whose purpose is to keep track of the ghost names used for a specific queue. Since we only have two ghost names for this specification, element of *Qghostnames* will simply be pairs. For an element $Q_\gamma \in Qghostnames$, the first element of the pair, written $Q_\gamma.\gamma_{Hlock}$, will contain the ghost name for the head lock, and the second element, $Q_\gamma.\gamma_{Tlock}$, the ghost name for the tail lock.

The sequential specification we wish to prove is the following:

$$\begin{aligned} & \exists \text{is_queue} : Val \rightarrow List\ Val \rightarrow Qghostnames \rightarrow Prop. \\ & \{True\} \text{ initialize() } \{v. \exists Q_\gamma, \text{is_queue } v \sqcap Q_\gamma\} \\ & \wedge \quad \forall q, v, xs_v, Q_\gamma. \{\text{is_queue } q\ xs_v\ Q_\gamma\} \text{ enqueue } q\ v \{w. \text{is_queue } q\ (v :: xs_v)\ Q_\gamma\} \\ & \wedge \quad \forall q, xs_v, Q_\gamma. \{\text{is_queue } q\ xs_v\ Q_\gamma\} \text{ dequeue } q \left\{ v. \begin{array}{l} (xs_v = [] * v = \text{None} * \text{is_queue } q\ xs_v\ Q_\gamma) \vee \\ (\exists x_v, xs'_v. xs_v = xs'_v ++ [x_v] * v = \text{Some } x_v * \text{is_queue } q\ xs'_v\ Q_\gamma) \end{array} \right\} \end{aligned}$$

The predicate $\text{is_queue } q\ xs_v\ Q_\gamma$ captures that the value q is a queue, whose content matches that of our abstract representation xs_v , and the queue uses the ghost names described by Q_γ . Note that the is_queue predicate is not required to be persistent, hence it cannot be duplicated and given to multiple threads. This is the sense in which this specification is sequential.

3.4 Proving the Sequential Specification

3.4.1 The is_queue Predicate

To prove the specification we must give a specific is_queue predicate. To help guide us in designing this, we give the following observations about the behaviour of the implementation.

1. Head always points to the first node in the queue.
2. Tail always points to either the last or second last node in the queue.
3. All but the last pointer in the queue (the pointer to None) never change.

Observation 2 captures the fact that, while enqueueing, a new node is first added to the linked list, and then later the tail is updated to point to the newly added node. Since only one thread can enqueue a node at a time (due to the lock), then the tail will only ever point to the last or second last due to the above. However, in a sequential setting, the tail will always appear to point to the last node, as no one can inspect the queue while the tail points to the second last.

Insight 3 means that we can mark all pointers in the queue (except the pointer to the null node) as persistent. This is technically not needed in the sequential case, but we will incorporate it now, as we will need it in the concurrent setting.

$$\begin{aligned}
\text{is_queue } q \, xs_v \, Q_\gamma &= \exists \ell_{\text{queue}}, \ell_{\text{head}}, \ell_{\text{tail}} \in \text{Loc}. \exists H_{\text{lock}}, T_{\text{lock}} \in \text{Val}. \\
q &= \ell_{\text{queue}} * \ell_{\text{queue}} \mapsto \square((\ell_{\text{head}}, \ell_{\text{tail}}), (H_{\text{lock}}, T_{\text{lock}})) * \\
&\exists xs_{\text{queue}} \in \text{List}(\text{Loc} \times \text{Val} \times \text{Loc}). \exists x_{\text{head}}, x_{\text{tail}} \in (\text{Loc} \times \text{Val} \times \text{Loc}). \\
\text{proj_val } xs_{\text{queue}} &= \text{wrap_some } xs_v * \\
\text{isLL}(xs_{\text{queue}} ++ [x_{\text{head}}]) &* \\
\ell_{\text{head}} &\mapsto (\text{in } x_{\text{head}}) * \\
\ell_{\text{tail}} &\mapsto (\text{in } x_{\text{tail}}) * \text{isLast } x_{\text{tail}} (xs_{\text{queue}} ++ [x_{\text{head}}]) * \\
\text{isLock } Q_\gamma \cdot \gamma_{H_{\text{lock}}} H_{\text{lock}} &\text{ True} * \\
\text{isLock } Q_\gamma \cdot \gamma_{T_{\text{lock}}} T_{\text{lock}} &\text{ True}.
\end{aligned}$$

This `is_queue` predicate states that the value q is a location, which always points to the structure containing the head, the tail, and the two locks. It also connects the abstract state xs_v with the concrete state (represented by xs_{queue}), by stating that if you strip away the locations connecting the nodes and remove the `Some` around the values, then you get the abstract state xs_v . Next, the predicate specifies the concrete state. There is some head node x_{head} , which the head points to. This head node and the nodes in xs_{queue} form the underlying linked list (specified using the `isLL` predicate below). There is also a tail node, which is the last node in the linked list, and the tail points to this node. Finally, we have the `isLock` predicate for our two locks. Since we are in a sequential setting, then the locks are superfluous, hence they simply protect `True`.

The `isLL` predicate essentially creates the structure seen in the examples of section 3.2. It is defined in two steps. Firstly, we create all the persistent pointers in the linked list using the `isLL_chain` predicate. Note that this in effect makes `isLL_chain xs` persistent for all xs .

Definition 3.4.1 (Linked List Chain Predicate)

$$\begin{aligned}
\text{isLL_chain } [] &\equiv \text{True} \\
\text{isLL_chain } [x] &\equiv \text{in } x \mapsto \square(\text{val } x, \text{out } x) \\
\text{isLL_chain } x :: x' :: xs &\equiv \text{in } x \mapsto \square(\text{val } x, \text{out } x) * \text{out } x' \mapsto \square \text{in } x * \text{isLL_chain } x' :: xs
\end{aligned}$$

Then, to define `isLL`, we add that the last node in the linked list points to `None`.

Definition 3.4.2 (Linked List Predicate)

$$\begin{aligned}
\text{isLL } [] &\equiv \text{True} \\
\text{isLL } x :: xs &\equiv \text{out } x \mapsto \text{None} * \text{isLL_chain } x :: xs
\end{aligned}$$

3.4.2 Proof outline

initialise

Proving the `initialise` spec amounts to stepping through the code, giving us the required resources, and then using these to create an instance of `is_queue` with the obtained resources. To begin with, we step through the lines creating the first node,

giving us locations ℓ_{1_in} , ℓ_{1_out} with $\ell_{1_out} \mapsto \text{None}$ and $\ell_{1_in} \mapsto (\text{None}, \ell_{1_out})$. We can then update the later points-to predicate to become persistent, giving us $\ell_{1_in} \mapsto \text{persistent}(() \text{None}, \ell_{1_out})$. We then step to the creation of the two locks, where we shall use the newlock specification asserting that the locks should protect True. This gives us two ghost names, γ_{lock} , γ_{lock} , which we will collect in a $Qnames$ pair, Q_γ . Next, we step through the allocations of the head, tail, and queue, which gives us locations ℓ_{head} , ℓ_{tail} , ℓ_{queue} , such that both ℓ_{head} and ℓ_{tail} point to node 1, and such that ℓ_{queue} points to the structure containing the head, tail, and two locks. This last points to predicate we update to become persistent. With this, we now have all the resources needed to prove the post-condition: $\exists Q_\gamma. \text{is_queue } \ell_{queue} Q_\gamma$. Proving this follows by a sequence of framing away the resources we obtained and instantiating existentials with the values we got above. Most noteworthy, we pick the empty list for xs_{queue} , and node 1 for xs_{head} and xs_{tail} .

Enqueue

►add line numbers to code, and refer to them in proof◄ For enqueue, we get in our pre-condition $\text{is_queue } q \ xs_v \ Q_\gamma$, and we wish to that, if we run enqueue $q \ v$, then we will get the $\text{is_queue } q \ (v :: xs_v) \ Q_\gamma$. The proposition $\text{is_queue } q \ xs_v \ Q_\gamma$ gives us all the resources we will need to step through the code. Firstly, we create a new node, node x_{new} , with $\text{val } x_{new} = v$. We then have to acquire the lock, which will just give us True.

The next line adds node x_{new} to the linked list, by first finding the tail, from the queue pointer ℓ_{queue} , and then finding the node that the tail points to, denoted x_{tail} , and finally writing updating the out location of x_{tail} to point to x_{new} . The resources needed to do this are all described in $\text{is_queue } q \ xs_v \ Q_\gamma$. Firstly, it tells us that ℓ_{queue} points to the structure containing ℓ_{tail} . Secondly, it tells us that ℓ_{tail} points to x_{tail} , which is the last node in the linked list ($xs_{queue} ++ [x_{head}]$). Thirdly, since we know that x_{tail} is the last node in the linked list, then by the isLL predicate, we know that x_{tail} points to None and that it has the node-like structure described by isLL_chain . This is all we need to step through the line, adding x_{new} to the linked list. After performing the write, we then get that x_{tail} points to x_{new} , instead of None. We make this points-to predicate persistent.

The next line swings the tail to x_{new} . As describe above, we already know that ℓ_{tail} points to x_{tail} , so we have the required resources to perform the write. Afterwards, we get that ℓ_{tail} points to x_{new} .

Finally, we release the lock using the release specification (and we simply give back True), and the only thing left is to prove the postcondition: $\text{is_queue } q \ (v :: xs_v) \ Q_\gamma$. For the existentials, we shall pick the ones we got from the precondition, with the exception for xs_{queue} and x_{tail} . For xs_{queue} , we shall use the same xs_{queue} we got from the precondition, but with x_{new} cons'ed to it, and for x_{tail} , we chose the new tail node: x_{new} . With these choices, proving $\text{is_queue } q \ (v :: xs_v) \ Q_\gamma$ is fairly straightforward.

Dequeue

For dequeue q , our precondition is $\text{is_queue } q \ xs_v \ Q_\gamma$, and we our post condition states that either the queue is empty, or there is a tail element which is returned by the function, and removed from the queue.

Stepping through the function, we first do the superfluous acquire. Next, we get the head node x_{head} through the queue pointer ℓ_{queue} . As described above for Enqueue, we get the resource to do this through $is_queue \ q \ x_{sv} \ Q_\gamma$. The is_queue predicate also tells us that x_{head} is a node in the linked list (described by the $isLL$ predicate), hence we can step through the code in the next line, which finds the node that x_{head} is pointing to. Now, depending on whether or not the queue is empty, x_{head} either points to `None`, or some node x_{head_next} . Thus, we shall perform a case analysis on xs_{queue} .

xs_{queue} is empty: In this case, we will have that $isLL[x_{head}]$, which tells us that x_{head} points to `None`. Hence, the "then" branch of the "if" will be taken. This branch simply releases the lock and returns `None`. In this case, we prove the first disjunction in the post-condition. Since xs_v is reflected in xs_{queue} , then we will be able to conclude that xs_v is empty, and since we haven't modified the queue, we can create $is_queue \ q \ x_{sv} \ Q_\gamma$ using the same resources we got from the pre-condition.

xs_{queue} is not empty: In this case, we can conclude that there must be some node x_{head_next} , which is the first node in xs_{queue} . I.e. $xs_{queue} = xs'_{queue} ++ [x_{head_next}]$. We can thus use the $isLL$ predicate to conclude that x_{head} must point to x_{head_next} . Hence the else branch will be taken. Since x_{head_next} is part of the linked list, then $isLL$ tells us it has the node-like structure, allowing us to extract its value in the first line of the else branch.

In the next line, we make the head pointer, ℓ_{head} point to x_{head_next} , and we have the resource to do this through $is_queue \ q \ x_{sv} \ Q_\gamma$.

Finally, we release the lock and return the value we got from x_{head_next} . We must now prove the post-condition, and this time we prove the second disjunct. Since xs_v is reflected in xs_{queue} , then it must also be the case that xs_v is non-empty, and it has a first element, x_v , which is related to the first element of xs_{queue} , i.e. x_{head_next} . This allows us to conclude that the returned value (`val x_{head_next}`) is exactly x_v , but wrapped in a `Some`, as we had to prove. Finally, we must prove $is_queue \ q \ xs'_v \ Q_\gamma$, where xs'_v is xs_v but with x_v removed. For the existentials, we pick the same values we got from the precondition, with the exception of xs_{queue} and x_{head} . For xs_{queue} we pick the same xs_{queue} we got from the precondition, but with the first element, x_{head_next} removed. By doing this, xs_{queue} will be reflexed in xs'_v . For x_{head} , we pick the new head, which we have obtained that ℓ_{head} points to: x_{head_next} . With these choices, we can prove the predicate.

3.5 Concurrent Specification

For the concurrent specification, we will need is_queue to be duplicable. To achieve this, we shall initially give up on tracking the abstract state of the queue, and instead add a predicate Φ , which we will ensure holds for all elements of the queue. In this way, when dequeuing, we at least know that if we get some value, then Φ holds of this

value. The specification we wish to prove is as follows.

$\exists \text{is_queue} : (Val \rightarrow Prop) \rightarrow Val \rightarrow Qnames \rightarrow Prop.$

$\forall \Phi : Val \rightarrow Prop.$

$\forall v, Q_\gamma. \text{is_queue } \Phi \ v \ Q_\gamma \implies \Box \text{is_queue } \Phi \ v \ Q_\gamma$

$\wedge \{ \text{True} \} \text{ initialize}() \{ v. \exists Q_\gamma. \text{is_queue } \Phi \ v \ Q_\gamma \}$

$\wedge \forall q, v, Q_\gamma. \{ \text{is_queue } \Phi \ q \ Q_\gamma * \Phi \ v \} \text{ enqueue } q \ v \{ v. \text{True} \}$

$\wedge \forall q, Q_\gamma. \{ \text{is_queue } \Phi \ q \ Q_\gamma \} \text{ dequeue } q \{ v. v = \text{None} \vee (\exists x_v, v = \text{Some } x_v * \Phi \ x_v) \}$

3.6 Proving the Concurrent Specification

3.6.1 The is_queue Predicate

As we did for the sequential specification, we note here some useful observations

1. Nodes in the linked list are never deleted. Hence, the linked list only ever grows.
2. The tail can lag one node behind Head.
3. At any given time, the queue is in one of four states:
 - (a) No threads are interacting with the queue (**Static**)
 - (b) A thread is enqueueing (**Enqueue**)
 - (c) A thread is dequeuing (**Dequeue**)
 - (d) A thread is enqueueing and a thread is dequeuing (**Both**)

Observation 2 might seem a little surprising, and indeed it stands in contrast to property 5 in [1], which states that the tail never lags behind head. I also didn't realise this possibility until a proof attempt using a model that "forgot" old nodes lead to an unprovable case (see section 3.6.2). The situation can occur when the queue is empty, and a thread performs an incomplete enqueue; it attaches the new node to the end, but before it can swing the tail to this new node, another thread performs a dequeue, which dequeues this new node, swinging the head to it. Now the tail is lagging a node behind the head.

It is not possible for the tail to point more than one node behind the head, as in order for this to happen, more nodes must be enqueued, but this can't happen before the current enqueue finishes, which will update the tail and bring it up to speed with the head.

Fortunately, this isn't an issue for safety, but a consequence of this possibility is that when modelling the queue, we must remember at least one "old" node (i.e. a dequeued node), as the tail might be pointing to this node. For the sake of simplicity in the model, the choice is made to remember an arbitrary amount of old nodes, which is represented by the list xs_{old} .

Observation 3 is a simple consequence of the implementation using two locks.

Since we want is_queue to be persistent, then we cannot directly state the points-to predicates as we did in the sequential case. However, we will still need all the same resources to be able to prove the specification. The solution is to have an invariant

which describes the concrete state of the queue. In the proofs, when we need access to some resource, we shall then access it by opening the invariant. We now present the invariant and explain it afterwards.

Definition 3.6.1 (Two-Lock M&S-Queue Invariant)

$$\begin{aligned}
& \text{queue_invariant } \Phi \ell_{\text{head}} \ell_{\text{tail}} Q_\gamma = \\
& \exists x_{s_v}. \quad \text{(the abstract state)} \\
& \exists xs, xs_{\text{queue}}, xs_{\text{old}}, x_{\text{head}}, x_{\text{tail}}. \quad \text{(the concrete state)} \\
& xs = xs_{\text{queue}} + +[x_{\text{head}}] + +xs_{\text{old}} * \\
& \text{isLL } xs * \\
& \text{proj_val } xs_{\text{queue}} = \text{wrap_some } x_{s_v} * \\
& \text{All } x_{s_v} \Phi * \\
& (\\
& \quad \ell_{\text{head}} \mapsto (\text{in } x_{\text{head}}) * \ell_{\text{tail}} \mapsto (\text{in } x_{\text{tail}}) * \text{isLast } x_{\text{tail}} xs * \quad \text{(Static)} \\
& \quad \text{ToknEQ}_\gamma * \text{ToknDQ}_\gamma * \text{TokUpdatedQ}_\gamma \\
& \quad \vee \\
& \quad \ell_{\text{head}} \mapsto (\text{in } x_{\text{head}}) * \ell_{\text{tail}} \mapsto 1/2(\text{in } x_{\text{tail}}) * \quad \text{(Enqueue)} \\
& \quad (\text{isLast } x_{\text{tail}} xs * \text{TokBefore } Q_\gamma \vee \text{isSndLast } x_{\text{tail}} xs * \text{TokAfter } Q_\gamma) * \\
& \quad \text{TokEQ}_\gamma * \text{ToknDQ}_\gamma \\
& \quad \vee \\
& \quad \ell_{\text{head}} \mapsto 1/2(\text{in } x_{\text{head}}) * \ell_{\text{tail}} \mapsto (\text{in } x_{\text{tail}}) * \text{isLast } x_{\text{tail}} xs * \quad \text{(Dequeue)} \\
& \quad \text{ToknEQ}_\gamma * \text{TokDQ}_\gamma * \text{TokUpdatedQ}_\gamma \\
& \quad \vee \\
& \quad \ell_{\text{head}} \mapsto 1/2(\text{in } x_{\text{head}}) * \ell_{\text{tail}} \mapsto 1/2(\text{in } x_{\text{tail}}) * \quad \text{(Both)} \\
& \quad (\text{isLast } x_{\text{tail}} xs * \text{TokBefore } Q_\gamma \vee \text{isSndLast } x_{\text{tail}} xs * \text{TokAfter } Q_\gamma) * \\
& \quad \text{TokEQ}_\gamma * \text{TokDQ}_\gamma \\
&)
\end{aligned}$$

In contrast to the sequential specification, the abstract state is now existentially quantified, hence the exact contents of the queue are not tracked. But the concrete state of the queue is still reflected in the abstract state through a value projection and wrapping of values. Another difference is that we now also keep track of an arbitrary number of "old" nodes; nodes that are behind the head node, x_{head} . As discussed above, this inclusion is due to observation 2.

As before, we also assert that the concrete state forms a linked list, as described by the `isLL` predicate.

The proposition `All $x_{s_v} \Phi$` states that all values in x_{s_v} (i.e. the values currently in the queue) satisfy the predicate Φ . This will allow us to conclude that dequeued values satisfy Φ .

The final part of the invariant describes the four possible states of the queue, as described in 3. Since the resources used by the queue are inside an invariant, and enqueueing/dequeueing threads need to access the resources of the queue multiple times, then

we will have to open and close the invariant multiple times. Each time we open the invariant, the existentially quantified variables will not be the same as those from early accesses of the invariant (as they are existentially quantified). Thus, the threads must be able to "match up" variables from previous accesses to later accesses. The way we shall achieve this is by allowing threads to keep a *fraction* of the points-to predicate that it is using. For instance, an enqueueing thread will have to access the points-to predicate concerning ℓ_{tail} multiple times, and in between accesses of the invariant, it can get to keep half of the points-to predicate. Thus, when it opens the invariant later, it will have $\ell_{tail} \mapsto 1/2x_{tail}$ from an earlier access, and it will obtain the existence of some new x'_{tail} , such that $\ell_{tail} \mapsto 1/2x'_{tail}$. Combining the two points-to predicates allows us to conclude that $x_{tail} = x'_{tail}$. In this way, we can match up variables from earlier access to variables in later accesses.

In the **Static** state where no thread is interacting with the queue, the queue owns all of the points-to predicates concerning the head and tail.

In the **Enqueue** state, the enqueueing thread owns half of the tail pointer, and we distinguish between two cases, as discussed in 2: either the enqueueing thread has yet to add the new node to the linked list, or the new node has been added, but the tail pointer is still pointing to the previous tail node.

In the **Dequeue** state, the dequeueing thread owns half of the head pointer, and the tail is as in the **Static** state.

Finally, the **Both** state is essentially a combination of the **Enqueue** and **Dequeue** states.

To track which state the queue is in, we use *tokens*. Tokens are defined using the exclusive resource algebra on the singleton set: $\text{EX}()$. This resource algebra only has one valid element, and combining two elements will give the non-valid element \perp . Thus, if we own a particular token, then, upon opening the invariant, we can rule out certain states simply because they mention the token we own.

We will use several tokens, each of which is the valid element of their own instance of $\text{EX}()$. Different instances are distinguish between using ghost names. Hence, each token will be represented by a ghost name. As we did for the sequential specification, we group these ghost names into a tuple Q_γ , and write, for instance TokEQ_γ to refer to the valid element of a particular instance. We proceed to explain the meaning of each of the tokens used in the invariant. **►make macros for the tokens◄**

- ToknEQ_γ represents that no threads are enqueueing.
- TokEQ_γ represents that a thread is enqueueing.
- ToknDQ_γ represents that no threads are dequeueing.
- TokDQ_γ represents that a thread is dequeueing.
- TokBeforeQ_γ represents that an enqueueing thread has not yet added the new node to the linked list.
- TokAfterQ_γ represents that an enqueueing thread has added the new node to the linked list, but not yet swung the tail.
- $\text{TokUpdatedQ}_\gamma$ is defined as $\text{TokBeforeQ}_\gamma * \text{TokAfterQ}_\gamma$, and represents that the queue is up to date.

Note: The concurrent specification for the two-lock Michael Scott Queue *can* be proven using the queue invariant 3.6.1, and the proof outline below will also be using this. However, a simpler (but arguably less intuitive) queue invariant was discovered. This simpler invariant is equivalent to 3.6.1 and has the benefit of being easier to work with in the mechanised proofs. Thus, in the mechanised proofs, the simpler variant is used. The simpler variant can be found in the appendix **►add appendix◄**.

3.6.2 Proof outline

►prove persistency of isqueue◄ The proofs structure for the specifications are largely similar to the sequential counterparts. The major difference is that we don't have access to the resource all the time; we must get it from the invariant, and we also have to keep track of which state we are in.

initialise

►do◄

Enqueue

►do◄

Dequeue

►do◄

Discussing the need for xs_{old}

►Update to reflect changes in Coq◄ As mentioned in the insights, it is possible for the tail to lag one node behind the head. This insight lead to including the old nodes of the queue in the queue invariant. This addition manifests in the end of the proof of dequeue. When we open the invariant to swing the head to the new node, we get that the entire queue is xs . After performing the write, we can then close the invariant with the same xs that we opened the queue to (just written differently to signify that x_{head} is now "old"). Because of this, we can supply the same predicate concerning the *tail* (the or) that we opened the queue with, since these only mention xs , which remains the same.

Had we not used an xs_{old} and essentially just "forgotten" old nodes in the queue, we couldn't have done this. Say that we defined xs as $xs = x_{head} :: xs_{rest}$ instead. Then, once we have to close the invariant, we cannot supply the xs , which we got when we opened the invariant. Our only choice (due to the fact that *head* must point to $x_{n_{head}}$) is to close the invariant with $xs' = xs_{rest} = x_{n_{head}} :: xs''_{n_{rest}}$. However, clearly $xs' \neq xs$, so we cannot supply the same predicate concerning the *tail* (the or) that we got when opening the invariant, since this predicate talks about xs , not xs' . Now, if we opened the invariant in the Dequeue case, then we could assert that $lastxs = lastxs'$, and hence still be able close the invariant. However, if we opened the invariant in the Both case, then we would need to assert that $2lastxs = 2lastxs'$. This is however not provable,

since it might be the case that $xs''_{n_{rest}}$ is empty, and hence $2lastxs'$ is *None*, whereas $2lastxs = x_{n_{head}}$.

Chapter 4

Conclusion

►conclude on the problem statement from the introduction◄

Bibliography

- [1] Maged M. Michael and Michael L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In James E. Burns and Yoram Moses, editors, *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, pages 267–275. ACM, 1996.
- [2] Simon Friis Vindum and Lars Birkedal. Contextual refinement of the michael-scott queue (proof pearl). In Catalin Hritcu and Andrei Popescu, editors, *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, pages 76–90. ACM, 2021.

Appendix A

The Technical Details

