TITLE HERE
Mathias Pedersen, 201808137

Master's Thesis, Computer Science
March 2024
Advisor: Amin Timany

AARHUS
UNIVERSITY
DEPARTMENT OF COMPUTER SCIENCE

Abstract

►in English...◄

Resumé

▶in Danish...◀

Acknowledgments

►…◄

Mathias Pedersen
Aarhus, March 2024.

Contents

Chapter 1

Introduction

►motivate and explain the problem to be addressed◄

►example of a citation: [1]◄ ►get your bibtex entries from `https://dblp.org/`◄

Chapter 2

Preliminaries

►Description of HeapLang, Iris, Verified in Coq (Weakest precondition vs Hoare triples)◄

Chapter 3

The Two-Lock Michael Scott Queue

I present here he an implementation of the Two-lock MS-Queue in HeapLang. This implementation differs slightly from the original, presented in [1], but most changes simply reflect the differences in the two languages.

3.1 Preliminaries

The underlying data structure making up the queue is a singly-linked list. The linked-list will always contain at least one element, called the *sentinel* node, marking the beginning of the queue. Note that the sentinel node is itself not part of the queue, but all nodes following it are. The queue keeps a head pointer (ℓ_{head}) which always points to the sentinel, and a tail pointer (ℓ_{tail}) which points to some node in the linked list.

In my implementation, a node can be thought of as a triple $(\ell_{i\_in}, v_i, \ell_{i\_out})$. The location ℓ_{i_in} points to the pair (v_i, ℓ_{i_out}), where v_i is the value of the node, and ℓ_{i_out} either points to None which represents the null pointer, or to the next node in the linked list. When we say that a location ℓ points to a node $(\ell_{i\_in}, v_i, \ell_{i\_out})$, we mean that $\ell \mapsto \ell_{i\_in}$. Hence, if we have two adjacent nodes $(\ell_{i\_in}, v_i, \ell_{i\_out})$, $(\ell_{i+1\_in}, v_{i+1}, \ell_{i+1\_out})$ in the linked list, then we have the following structure: $\ell_{i\_in} \mapsto (v_i, \ell_{i\_out})$, $\ell_{i\_out} \mapsto \ell_{i+1\_in}$, and $\ell_{i+1\_in} \mapsto v_{i+1}, \ell_{i+1\_out}$.

The reader may wonder why there is an extra, intermediary "in" pointer, between the pairs of the linked list, and why the "out" pointer couldn't point directly to the next pair. In the original implementation [1], nodes are allocated on the heap. To simulate this in HeapLang, when creating a new node, we create a pointer to a pair making up the node. Now, in the C-like language used in the original specification, an assignment operator is available which is not present in HeapLang. So in order to mimic this behaviour, we model variables as pointers. In this way, we can model a variable x as a location ℓ_x, and the value stored at ℓ_x is the current value of x. This means that the variable ℓ_{i_out} (called "next" in the original) becomes a location ℓ_{head}, and the value stored at the location is what head is currently assigned to. Since ℓ_{i_out} is supposed to be a variable containing a pointer, then the value saved at that location will also be a pointer.

3.2 implementation

The queue consists of 3 functions: initialize, enqueue, and dequeue which I now present in turn.

3.2.1 initialize

initialize will first create a single node – the sentinel – marking the start of the linked list. It then creates two locks, *H_lock* and *T_lock*, protecting the head and tail pointers, respectively. Finally, it creates the head and tail pointers, both pointing to the sentinel. The queue is then a pointer to a structure containing the head, the tail, and the two locks.

Figure 3.1 illustrates the structure of the queue after initialisation. Note that one of the pointers is coloured blue. This represents a *persistent* pointer; a pointer that will never be updated again. All "in" pointers ℓ_{i_in}, are persistent, meaning that they will always point to $(v_i, \ell_{i\_\text{out}})$. We shall use the notation $\ell \mapsto \Box v$ (introduced in [2]) to mean that ℓ points persistently to v.

Note that in the original specification, a queue is a pointer to a 4-tuple $(\ell_{head}, \ell_{tail}, H\_lock, T\_lock)$. Since HeapLang doesn't support 4-tuples, we instead represent the queue as a pointer to a pair of pairs: $((\ell_{head}, \ell_{tail}), (H\_lock, T\_lock))$.

3.2.2 enqueue

To enqueue a value, we must create a new node, append it to the underlying linked-list, and swing the tail pointer to this new node. These three operations are depicted in figure 3.2.

enqueue takes as argument the value to be enqueued and creates a new node containing this value (corresponding to figure 3.2a). This creation doesn't interact with the underlying queue data-structure, hence why we don't acquire the *T_lock* first. After creating the new node, we must make the last node in the linked list point to it. Since this operation interacts with the queue, we first acquire the *T_lock*. Once we obtain the lock, we make the last node in the linked list point to our new node (figure 3.2b). Following this, we swing ℓ_{tail} to the new last node in the linked list (figure 3.2c).

Figure 3.2 also illustrates when pointers become persistent; once the previous last node is updated to point to the newly inserted node, that pointer will never be updated again, hence becoming persistent.

3.2.3 dequeue

It is of course only possible to dequeue an element from the queue if the queue contains at least one element. Hence, the first thing dequeue does is check if the queue is empty. We can detect an empty queue by checking if the sentinel is the last node in the linked list. Being the last node in the linked list corresponds to having the "out" node be None. If this is the case, then the queue is empty and the code returns None. Otherwise, there is a node just after the sentinel, which is the first node of the queue. To dequeue it, we first read the associated value, and next we swing the head to it, making it the new sentinel. Finally, we return the value we read.

Since all of these operations interact with the queue, we shall only perform them after having acquired H_lock.

Figure 3.3 illustrates running dequeue on a non-empty queue. Note that the only change is that the head pointer is swung to the next node in the linked list; the old sentinel is not deleted, it just become unreachable from the heap pointer. In this way, the linked list only ever grows.
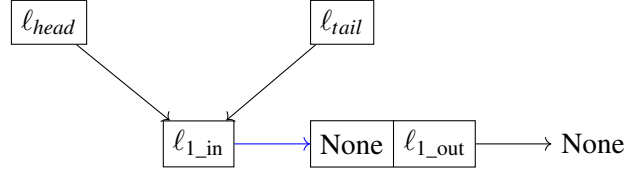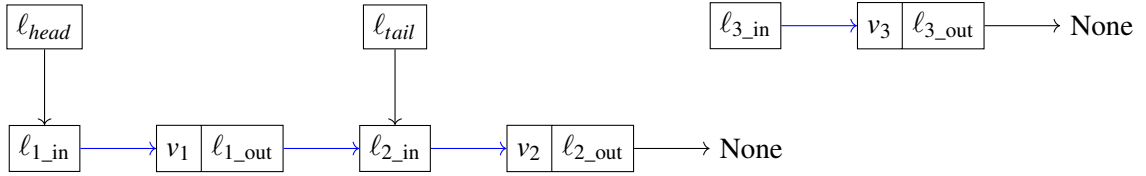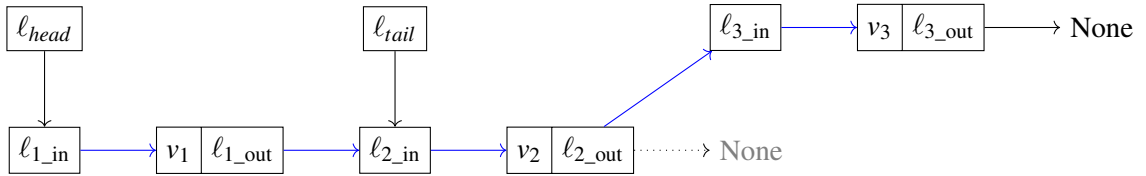

Figure 3.1: Queue after initialisation

(a) Queue after creating the new node $(\ell_{3\_in}, v_3, \ell_{3\_out})$ to be added to the queue.

(b) Queue after adding the new node to linked list.
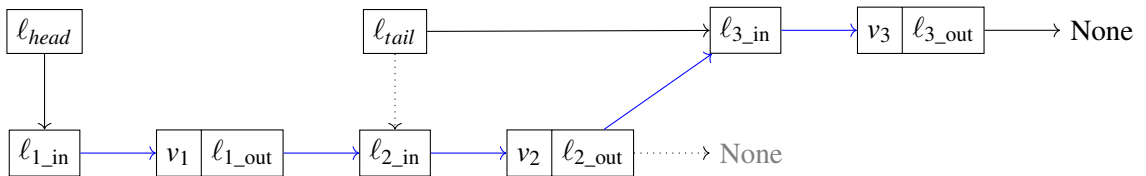
(c) Queue after swinging tail pointer to the new node.

Figure 3.2: Enqueuing an element to a queue with one element.

$$
\begin{aligned}
&\mathsf{let}\ \text{initialize} := \\
&\qquad \mathsf{let}\ node = \mathsf{ref}\,((\text{None}, \mathsf{ref}\,(\text{None})))\ \mathsf{in} \\
&\qquad \mathsf{let}\ H\_lock = newlock()\ \mathsf{in} \\
&\qquad \mathsf{let}\ T\_lock = newlock()\ \mathsf{in} \\
&\qquad \mathsf{ref}\,((\mathsf{ref}\,(node), \mathsf{ref}\,(node)), (H\_lock, T\_lock))
\end{aligned}
$$

7

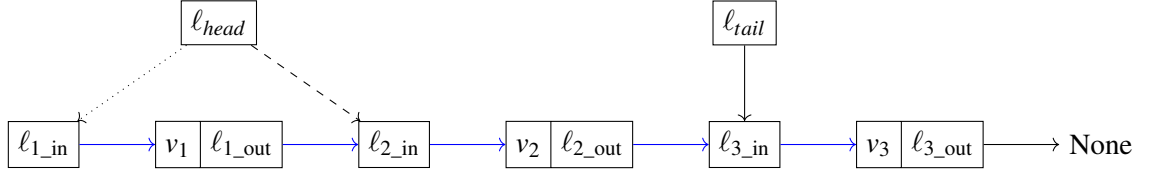Figure 3.3: Dequeueing an element (v_2) from a queue with two elements (v_2, v_3). The dotted line represents the state before the dequeue, and the dashed line is the state after dequeuing.

$$
\begin{aligned}
&\text{let } \text{enqueue } Q \; value := \\
&\qquad \text{let } node = \text{ref } ((\text{Some } value, \text{ref } (\text{None}))) \text{ in} \\
&\qquad \text{acquire}(\text{snd}(\text{snd}(!Q))); \\
&\qquad \text{snd}(!(!(\text{snd}(\text{fst}(!Q))))) \leftarrow node; \\
&\qquad \text{snd}(\text{fst}(!Q)) \leftarrow node; \\
&\qquad \text{release}(\text{snd}(\text{snd}(!Q)))
\end{aligned}
$$

$$
\begin{aligned}
&\text{let } \text{dequeue } Q := \\
&\qquad \text{acquire}(\text{fst}(\text{snd}(!Q))); \\
&\qquad \text{let } node = !(\text{fst}(\text{fst}(!Q))) \text{ in} \\
&\qquad \text{let } new\_head = !(\text{snd}(!node)) \text{ in} \\
&\qquad \text{if } new\_head = \text{None } \text{then} \\
&\qquad\qquad \text{release}(\text{fst}(\text{snd}(!Q))); \\
&\qquad\qquad \text{None} \\
&\qquad \text{else} \\
&\qquad\qquad \text{let } value = \text{fst}(!new\_head) \text{ in} \\
&\qquad\qquad \text{fst}(\text{fst}(!Q)) \leftarrow new\_head; \\
&\qquad\qquad \text{release}(\text{fst}(\text{snd}(!Q))); \\
&\qquad\qquad value
\end{aligned}
$$

3.3 Sequential Specification

Let us first prove a specification for the two-lock michael scott queue in the simple case where we don't allows for concurrency. In this case, we know that only a single thread will interact with the queue at any given point in a sequential manner. This means that we give a specification that tracks the exact contents of the queue. To this end, we shall define the abstract state of the queue, denoted xs_v as a list of HeapLang values. I.e. $xs_v : List\ Val$. We adopt the convention that enqueueing an element is done by adding it to the front of the list, and dequeueing removes the last element of the list (if such an element exists). The reason for this choice is purely technical.

Since the queue uses two locks, we will get two ghost names; one for each lock. For this specification, these are the only two ghost names we will need. However, for the later specifications, we will use more resource algebra, and will need more ghost names. Thus, to ease notation, we shall define the type "*Qgnames*" whose purpose is to keep track of the ghost names used for a specific queue. Since we only have two ghost names for this specification, element of *Qgnames* will simply be pairs. For an element $Q_\gamma \in$ *Qgnames*, the first element of the pair, written $Q_\gamma.\gamma_{Hlock}$, will contain the ghost name for the head lock, and the second element, $Q_\gamma.\gamma_{Tlock}$, the ghost name for the head lock.

The sequential specification we wish to prove is the following:

$\exists$ is_queue : *Val* $\rightarrow$ *List Val* $\rightarrow$ *Qgnames* $\rightarrow$ Prop.

$\qquad \{\text{True}\}$ initialize() $\{v.\exists Q_\gamma, \text{is\_queue } v \text{ } [] \text{ } Q_\gamma\}$

$\wedge \quad \forall q, v, xs_v, Q_\gamma.\{\text{is\_queue } q \text{ } xs_v \text{ } Q_\gamma\}$ enqueue $q \text{ } v \text{ } \{w. \text{is\_queue } q \text{ } (v :: xs_v) \text{ } Q_\gamma\}$

$\wedge \quad \forall q, xs_v, Q_\gamma.\{\text{is\_queue } q \text{ } xs_v \text{ } Q_\gamma\}$ dequeue $q \left\{ v. \begin{array}{l} (xs_v = [] * v = \text{None} * \text{is\_queue } q \text{ } xs_v \text{ } Q_\gamma) \vee \\ (\exists x_v, xs'_v \text{ . } xs_v = xs'_v \text{ ++ } [x_v] * v = \text{Some } x_v * \text{is\_queue } q \text{ } xs'_v \text{ } Q_\gamma) \end{array} \right\}$

The predicate is_queue $q \text{ } xs_v \text{ } Q_\gamma$ captures that the value q is a queue, whose content matches that of our abstract representation xs_v, and the queue uses the ghost names described by Q_γ. Note that the is_queue predicate is not required to be persistent, hence it cannot be duplicated and given to multiple threads. This is the sense in which this specification is sequential.

3.4 Proving the Sequential Specification

3.4.1 The is_queue Predicate

To prove the specification we must give a specific is_queue predicate. To help guide us in designing this, we give the following observations about the behaviour of the implementation.

1. Head always points to the first node in the queue.

2. Tail always points to either the last or second last node in the queue.

3. All but the last pointer in the queue (the pointer to None) never change.

Observation 2 captures the fact that, while enqueueing, a new node is first added to the linked list, and then later the tail is updated to point to the newly added node. Since only one thread can enqueue a node at a time (due to the lock), then the tail will only ever point to the last or second last due to the above. However, in a sequential setting, the tail will always appear to point to the last node, as no one can inspect the queue while the tail points to the second last.

Insight 3 means that we can mark all pointers in the queue (except the pointer to the null node) as persistent. This is technically not needed in the sequential case, but we will incorporate it now, as we will need it in the concurrent setting.

9

is_queue $q\ xs_v\ Q_\gamma = \exists \ell_{queue}, \ell_{head}, \ell_{tail} \in Loc. \exists H_{lock}, T_{lock} \in Val.$

$\qquad\qquad q = \ell_{queue} * \ell_{queue} \mapsto \Box((\ell_{head}, \ell_{tail}), (H_{lock}, T_{lock})) *$

$\qquad\qquad \exists xs_{queue} \in List(Loc \times Val \times Loc). \exists x_{head}, x_{tail} \in (Loc \times Val \times Loc).$

$\qquad\qquad proj\_val\ xs_{queue} = wrap\_some\ xs_v *$

$\qquad\qquad isLL(xs_{queue} + +[x_{head}]) *$

$\qquad\qquad \ell_{head} \mapsto (\text{in } x_{head}) *$

$\qquad\qquad \ell_{tail} \mapsto (\text{in } x_{tail}) * isLast\ x_{tail}\ (xs_{queue} + +[x_{head}]) *$

$\qquad\qquad isLock\ Q_\gamma.\gamma_{Hlock}\ H_{lock}\ \mathsf{True} *$

$\qquad\qquad isLock\ Q_\gamma.\gamma_{Tlock}\ T_{lock}\ \mathsf{True}.$

This is_queue predicate states that the value q is a location, which always points to the structure containing the head, the tail, and the two locks. It also connects the abstract state xs_v with the concrete state (represented by xs_{queue}), by stating that if you strip away the locations connecting the nodes and remove the Some around the values, then you get the abstract state xs_v. Next, the predicate specifies the concrete state. There is some head node x_{head}, which the head points to. This head node and the nodes in xs_{queue} form the underlying linked list (specified using the isLL predicate below). There is also a tail node, which is the last node in the linked list, and the tail points to this node. Finally, we have the isLock predicate for our two locks. Since we are in a sequential setting, then the locks are superfluous, hence they simply protect True.

The isLL predicate essentially creates the structure seen in the examples of section 3.2. It is defined in two steps. Firstly, we create all the persistent pointers in the linked list using the isLL_chain predicate. Note that this in effect makes isLL_chain xs persistent for all xs.

Definition 3.4.1 (Linked List Chain Predicate)

$$\text{isLL\_chain } [] \equiv \textit{True}$$
$$\text{isLL\_chain } [x] \equiv \text{in } x \mapsto \Box(\text{val } x, \text{out } x)$$
$$\text{isLL\_chain } x :: x' :: xs \equiv \text{in } x \mapsto \Box(\text{val } x, \text{out } x) * \text{out } x' \mapsto \Box \text{in } x * \text{isLL\_chain } x' :: xs$$

Then, to define isLL, we add that the last node in the linked list points to None.

Definition 3.4.2 (Linked List Predicate)

$$\text{isLL } [] \equiv \textit{True}$$
$$\text{isLL } x :: xs \equiv \text{out } x \mapsto \text{None} * \text{isLL\_chain } x :: xs$$

3.4.2 Proof outline

initialise

Proving the initialise spec amounts to stepping through the code, giving us the required resources, and then using these to create an instance of is_queue with the obtained resources. To begin with, we step through the lines creating the first node,

giving us locations ℓ_{1_in}, ℓ_{1_out} with $\ell_{1\_out} \mapsto$ None and $\ell_{1\_in} \mapsto$ (None, ℓ_{1_out}). We can then update the later points-to predicate to become persistent, giving us $\ell_{1\_in} \mapsto$ persistent(() None, ℓ_{1_out}). We then step to the creation of the two locks, where we shall use the newlock specification asserting that the locks should protect True. This gives us two ghost names, γ_{Tlock}, γ_{Tlock}, which we will collect in a *Qgnames* pair, Q_γ. Next, we step through the allocations of the head, tail, and queue, which gives us locations ℓ_{head}, ℓ_{tail}, ℓ_{queue}, such that both ℓ_{head} and ℓ_{tail} point to node 1, and such that ℓ_{queue} points to the structure containing the head, tail, and two locks. This last points to predicate we update to become persistent. With this, we now have all the resources needed to prove the post-condition: $\exists Q_\gamma.$ is_queue $\ell_{queue} \, Q_\gamma$. Proving this follows by a sequence of framing away the resources we obtained and instantiating existentials with the values we got above. Most noteworthy, we pick the empty list for xs_{queue}, and node 1 for xs_{head} and xs_{tail}.

Enqueue

►**add line numbers to code, and refer to them in proof**◄ For enqueue, we get in our pre-condition is_queue $q \, xs_v \, Q_\gamma$, and we wish to that, if we run enqueue $q \, v$, then we will get the is_queue $q \, (v :: xs_v) \, Q_\gamma$. The proposition is_queue $q \, xs_v \, Q_\gamma$ gives us all the resources we will need to step through the code. Firstly, we create a new node, node x_{new}, with val $x_{new} = v$. We then have to acquire the lock, which will just give us True.

The next line adds node x_new to the linked list, by first finding the tail, from the queue pointer ℓ_{queue}, and then finding the node that the tail points to, denoted x_{tail}, and finally writing updating the out location of x_{tail} to point to x_new. The resources needed to do this are all described in is_queue $q \, xs_v \, Q_\gamma$. Firstly, it tells us that ℓ_{queue} points to the structure containing ℓ_{tail}. Secondly, it tells us that ℓ_{tail} points to x_{tail}, which is the last node in the linked list ($xs_{queue} + +[x_{head}]$). Thirdly, since we know that x_{tail} is the last node in the linked list, then by the isLL predicate, we know that x_{tail} points to None and that it has the node-like structure described by isLL_chain . This is all we need to step through the line, adding x_{new} to the linked list. After performing the write, we then get that x_{tail} points to x_{new}, instead of None. We make this points-to predicate persistent.

The next line swings the tail to x_{new}. As describe above, we already know that ℓ_{tail} points to x_{tail}, so we have the required resources to perform the write. Afterwards, we get that ℓ_{tail} points to x_{new}.

Finally, we release the lock using the release specification (and we simply give back True), and the only thing left is to prove the postcondition: is_queue $q \, (v :: xs_v) \, Q_\gamma$. For the existentials, we shall pick the ones we got from the precondition, with the exception for xs_{queue} and x_{tail}. For xs_{queue}, we shall use the same xs_{queue} we got from the precondition, but with xs_new cons'ed to it, and for x_{tail}, we chose the new tail node: x_{new}. With these choices, proving is_queue $q \, (v :: xs_v) \, Q_\gamma$ is fairly straightforward.

Dequeue

For dequeue q, our precondition is is_queue $q \, xs_v \, Q_\gamma$, and we our post condition states that either the queue is empty, or there is a tail element which is returned by the function, and removed from the queue.

Stepping through the function, we first do the superfluous acquire. Next, we get the head node x_{head} through the queue pointer ℓ_{queue}. As described above for Enqueue, we get the resource to do this through is_queue $q\ xs_v\ Q_\gamma$. The is_queue predicate also tells us that x_{head} is a node in the linked list (described by the isLL predicate), hence we can step through the code in the next line, which finds the node that x_{head} is pointing to. Now, depending on whether or not the queue is empty, x_{head} either points to None, or some node x_{head_next}. Thus, we shall perform a case analysis on xs_{queue}.

xs_{queue} **is empty**: In this case, we will have that isLL$[x_{head}]$, which tells us that $x_h ead$ points to None. Hence, the "then" branch of the "if" will be taken. This branch simply releases the lock and returns None. In this case, we prove the first disjunction in the post-condition. Since xs_v is reflected in xs_{queue}, then we will be able to conclude that xs_v is empty, and since we haven't modified the queue, we can create is_queue $q\ xs_v\ Q_\gamma$ using the same resources we got from the pre-condition.

xs_{queue} **is not empty**: In this case, we can conclude that there must be some node x_{head_next}, which is the first node in xs_{queue}. I.e. $xs_{queue} = xs'_{queue} + +[x_{head\_next}]$. We can thus use the isLL predicate to conclude that x_{head} must point to x_{head_next}. Hence the else branch will be taken. Since x_{head_next} is part of the linked list, then isLL tells us it has the node-like structure, allowing us to extract its value in the first line of the else branch.
In the next line, we make the head pointer, ℓ_{head} point to x_{head_next}, and we have the resource to do this through is_queue $q\ xs_v\ Q_\gamma$.
Finally, we release the lock and return the value we got from x_{head_next}. We must now prove the post-condition, and this time we prove the second disjunct. Since xs_v is reflected in xs_{queue}, then it must also be the case that xs_v is non-empty, and it has a first element, x_v, which is related to the first element of xs_{queue}, i.e. x_{head_next}. This allows us to conclude that the returned value (val x_{head_next}) is exactly x_v, but wrapped in a Some, as we had to prove. Finally, we must prove is_queue $q\ xs'_v\ Q_\gamma$, where xs'_v is xs_v but with x_v removed. For the existentials, we pick the same values we got from the precondition, wit the exception of xs_{queue} and x_{head}. For xs_{queue} we pick the same xs_{queue} we got from the precondition, but with the first element, x_{head_next} removed. By doing this, xs_{queue} will be reflexed in xs'_v. For $x_h ead$, we pick the new head, which we have obtained that ℓ_{head} points to: x_{head_next}. With these choices, we can prove the predicate.

3.5   Concurrent Specification

For the concurrent specification, we will need is_queue to be duplicable. To achieve this, we shall initially give up on tracking the abstract state of the queue, and instead add a predicate Φ, which we will ensure holds for all elements of the queue. In this way, when dequeueing, we at least know that if we get some value, then Φ holds of this

value. The specification we wish to prove is as follows.

$\exists \text{is\_queue} : (Val \rightarrow \text{Prop}) \rightarrow Val \rightarrow Qgnames \rightarrow \text{Prop}.$

$\forall \Phi : Val \rightarrow \text{Prop}.$

$\quad\quad \forall v, Q_\gamma. \text{is\_queue } \Phi \ v \ Q_\gamma \implies \Box \text{is\_queue } \Phi \ v \ Q_\gamma$

$\wedge \quad \{\text{True}\} \text{ initialize}() \ \{v.\exists Q_\gamma, \text{is\_queue } \Phi \ v \ Q_\gamma\}$

$\wedge \quad \forall q, v, Q_\gamma. \{\text{is\_queue } \Phi \ q \ Q_\gamma * \Phi \ v\} \text{ enqueue } q \ v \ \{v.\text{True}\}$

$\wedge \quad \forall q, Q_\gamma. \{\text{is\_queue } \Phi \ q \ Q_\gamma\} \text{ dequeue } q \ \{v.v = \text{None} \vee (\exists x_v, v = \text{Some} \, x_v * \Phi \ x_v)\}$

3.6 Proving the Concurrent Specification

3.6.1 The is_queue Predicate

As we did for the sequential specification, we note here some useful observations

1. Nodes in the linked list are never deleted. Hence, the linked list only ever grows.

2. The tail can lag one node behind Head.

3. At any given time, the queue is in one of four states:

   (a) No threads are interacting with the queue (**Static**)

   (b) A thread is enqueueing (**Enqueue**)

   (c) A thread is dequeuing (**Dequeue**)

   (d) A thread is enqueueing and a thread is dequeuing (**Both**)

Observation 2 might seem a little surprising, and indeed it stands in contrast to property 5 in [1], which states that the tail never lags behind head. I also didn't realise this possibility until a proof attempt using a model that "forgot" old nodes lead to an unprovable case (see section 3.6.2). The situation can occur when the queue is empty, and a thread performs an incomplete enqueue; it attaches the new node to the end, but before it can swing the tail to this new node, another thread performs a dequeue, which dequeues this new node, swinging the head to it. Now the tail is lagging a node behind the head.

It is not possible for the tail to point more than one node behind the head, as in order for this to happen, more nodes must be enqueued, but this can't happen before the current enqueue finishes, which will update the tail and bring it up to speed with the head.

Fortunately, this isn't an issue for safety, but a consequence of this possibility is that when modelling the queue, we must remember at least one "old" node (i.e. a dequeued node), as the tail might be pointing to this node. For the sake of simplicity in the model, the choice is made to remember an arbitrary amount of old nodes, which is represented by the list xs_{old}.

Observation 3 is a simple consequence of the implementation using two locks.

Since we want is_queue to be persistent, then we cannot directly state the points-to predicates as we did in the sequential case. However, we will still need all the same resources to be able to prove the specification. The solution is to have an invariant

which describes the concrete state of the queue. In the proofs, when we need access to some resource, we shall then access it by opening the invariant. We now present the invariant and explain it afterwards.

Definition 3.6.1 (Two-Lock M&S-Queue Invariant) ►*Write it in*◄

►**explain the invariant: Concrete xs, and abstract xs_v, with xs_queue being reflected in xs_v, tokens (and RA), the four states, quotient points to, isSndLast**◄
The specification for the two-lock Michael Scott Queue *can* be proven using the queue invariant 3.6.1, and the proof outline below will also be using this. However, a simpler (but arguably less intuitive) queue invariant was discovered. This simpler invariant is equivalent to 3.6.1 and has the benefit of being easier to work with in the mechanised proofs. Thus, in the mechanised proofs, the simpler variant is used. The simpler variant can be found in the appendix ►**add appendix**◄.

3.6.2 Proof outline

►**prove persistency of isqueue**◄ The proofs structure for the specifications are largely similar to the sequential counterparts. The major difference is that we don't have access to the resource all the time; we must get it from the invariant, and we also have to keep track of which state we are in.

initialise

►**do**◄

Enqueue

►**do**◄

Dequeue

►**do**◄

Discussing the need for xs_{old}

►**Update to reflect changes in Coq**◄ As mentioned in the insights, it is possible for the tail to lag one node behind the head. This insight lead to including the old nodes of the queue in the queue invariant. This addition manifests in the end of the proof of dequeue. When we open the invariant to swing the head to the new node, we get that the entire queue is *xs*. After performing the write, we can then close the invariant with the same *xs* that we opened the queue to (just written differently to signify that x_{head} is now "old"). Because of this, we can supply the same predicate concerning the *tail* (the or) that we opened the queue with, since these only mention *xs*, which remains the same.

Had we not used an xs_{old} and essentially just "forgotten" old nodes in the queue, we couldn't have done this. Say that we defined *xs* as $xs = x_{head} :: xs_{rest}$ instead. Then, once we have to close the invariant, we cannot supply the *xs*, which we got when we

14

opened the invariant. Our only choice (due to the fact that *head* must point to x_{n_head}) is to close the invariant with $xs' = xs_{rest} = x_{n_head} :: xs''_{n_r est}$. However, clearly $xs' \neq xs$, so we cannot supply the same predicate concerning the *tail* (the or) that we got when opening the invariant, since this predicate talks about *xs*, not xs'. Now, if we opened the invariant in the Dequeue case, then we could assert that $lastxs = lastxs'$, and hence still be able close the invariant. However, if we opened the invariant in the Both case, then we would need to assert that $2lastxs = 2lastxs'$. This is however not provable, since it might be the case that $xs''_{n_r est}$ is empty, and hence $2lastxs'$ is *None*, whereas $2lastxs = x_{n_head}$.

Chapter 4

Conclusion

►conclude on the problem statement from the introduction◄

Bibliography

[1] Maged M. Michael and Michael L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In James E. Burns and Yoram Moses, editors, *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, pages 267–275. ACM, 1996.

[2] Simon Friis Vindum and Lars Birkedal. Contextual refinement of the michael-scott queue (proof pearl). In Catalin Hritcu and Andrei Popescu, editors, *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, pages 76–90. ACM, 2021.

Appendix A

The Technical Details

▶...◀